

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33

# Contents

- 1 FP Recollections 7**
  - 1.1 Terms and Reduction . . . . . 7
  - 1.2 Types . . . . . 7
  - 1.3 Functions in Haskell . . . . . 9
  - 1.4 Data in Haskell . . . . . 11
  - 1.5 Type Classes in Haskell . . . . . 12
  - 1.6 Monads in Haskell . . . . . 13
  
- 2 Parametric Polymorphism 16**
  - 2.1 Parametric polymorphism in the language . . . . . 16
    - 2.1.1 Emulating type abstractions and applications (**Proxy**) . . . . . 17
    - 2.1.2 First-class polymorphism . . . . . 18
    - 2.1.3 Higher-order/kinded polymorphism . . . . . 20
    - 2.1.4 Generalized algebraic data types (GADTs) . . . . . 20
    - 2.1.5 Type-level structures, data promotion . . . . . 21
  - 2.2 Implementation of parametric polymorphism . . . . . 25
    - 2.2.1 Monomorphization . . . . . 25
    - 2.2.2 Type erasure . . . . . 26
    - 2.2.3 Hybrid approach . . . . . 26
    - 2.2.4 Using a value witness table for types . . . . . 27
  - 2.3 Representation polymorphism . . . . . 28
    - 2.3.1 Varieties of runtime representations in Haskell . . . . . 28
    - 2.3.2 Classification of values by runtime representation . . . . . 29
    - 2.3.3 Representation polymorphism . . . . . 30
  
- 3 Ad-hoc Polymorphism 31**
  - 3.1 Type classes in the language . . . . . 32
    - 3.1.1 Dictionaries . . . . . 32
    - 3.1.2 Implicit arguments . . . . . 34
    - 3.1.3 Instance inference . . . . . 34
    - 3.1.4 Constructing a type from a value . . . . . 36
    - 3.1.5 Implicits and coherence . . . . . 37

---

<sup>1</sup>Author Andrey Stoyan (andrey.stoyan.csam@gmail.com).  
<sup>2</sup>Thanks to Ilya Kolegov for the first careful read and many comments. Thanks to my students whose interest in the subject made this work possible.  
<sup>3</sup>Version 1 мая 2026 г., 18:10.

34	3.1.6	Rules and specialization . . . . .	38
35	3.1.7	Digression: defunctionalization . . . . .	39
36	3.1.8	Emulating higher-kinded polymorphism . . . . .	40
37	3.2	Families . . . . .	41
38	3.2.1	Data families . . . . .	41
39	3.2.2	Synonym families . . . . .	42
40	3.2.3	Injective families . . . . .	43
41	3.2.4	First-class families . . . . .	43
42	3.3	The <code>Constraint</code> kind . . . . .	44
43	3.4	Uses of ad-hoc polymorphism . . . . .	45
44	3.4.1	Serialization . . . . .	45
45	3.4.2	Existential types . . . . .	46
46	3.4.3	Name resolution . . . . .	47
47	3.4.4	Non-syntactic type equalities, System FC . . . . .	48
48	3.4.5	Coercions and roles . . . . .	49
49	3.4.6	Type reflection . . . . .	51
50	3.4.7	Data reflection . . . . .	52
51	3.4.8	Open structures . . . . .	53
52	3.4.9	Exceptions and the open hierarchy . . . . .	53
53	3.4.10	Lightweight partial stack traces . . . . .	55
54	3.4.11	Customizable type errors . . . . .	55
55	<b>4</b>	<b>Data Types</b> . . . . .	<b>57</b>
56	4.1	Variance . . . . .	57
57	4.2	Isomorphism . . . . .	59
58	4.2.1	Cardinality: sums, products, exponentials . . . . .	60
59	4.2.2	Algebraic representation of a type . . . . .	61
60	4.3	Recursive types . . . . .	63
61	4.3.1	Just a list . . . . .	63
62	4.3.2	Fixpoint of a functor . . . . .	64
63	4.3.3	Recursion schemes . . . . .	65
64	4.4	Everything through folds . . . . .	67
65	4.4.1	Deforestation & list fusion . . . . .	69
66	4.4.2	Visitor pattern . . . . .	71
67	4.5	Everything through unfolds . . . . .	72
68	4.5.1	Abstract data types . . . . .	73
69	4.5.2	Stream fusion . . . . .	74
70	4.6	Ubiquitous duality . . . . .	75
71	4.6.1	Push vs pull streaming . . . . .	75
72	4.6.2	Data vs codata . . . . .	76
73	4.7	Appendix: category of algebras . . . . .	78

74	<b>5</b>	<b>Interpreters</b>	<b>81</b>
75	5.1	Interpreters as the foundation . . . . .	81
76	5.1.1	The tower of interpreters . . . . .	81
77	5.1.2	Interpreters everywhere . . . . .	82
78	5.1.3	Interpreters and programming language semantics . . . . .	83
79	5.1.4	Embedded domain-specific languages (eDSL) . . . . .	84
80	5.1.5	Example: the Accelerate library . . . . .	86
81	5.2	Value types . . . . .	87
82	5.2.1	Untyped tagless interpreters . . . . .	87
83	5.2.2	Typed tagged interpreters . . . . .	87
84	5.2.3	Typed tagless interpreters . . . . .	88
85	5.3	Bindings and first-class functions . . . . .	88
86	5.3.1	Name semantics . . . . .	89
87	5.3.2	Substitutions . . . . .	89
88	5.3.3	Environment . . . . .	91
89	5.3.4	Closures . . . . .	91
90	5.3.5	Typed context . . . . .	92
91	5.3.6	Meta-circular interpretation . . . . .	93
92	5.3.7	Higher-order abstract syntax . . . . .	94
93	5.3.8	Serialization of functions . . . . .	94
94	5.4	Tagless final interpreters . . . . .	95
95	5.4.1	Different interpretations for shallow embedding . . . . .	96
96	5.4.2	Going all the way . . . . .	96
97	5.4.3	Recovering compositionality of semantics . . . . .	97
98	5.4.4	Typed tagless final interpreter . . . . .	98
99	5.4.5	Meeting old friends: <b>Applicative</b> , <b>Monad</b> . . . . .	99
100	5.5	Expression problem . . . . .	101
101	5.5.1	Coproduct of functors . . . . .	102
102	5.5.2	Product of algebras . . . . .	103
103	<b>6</b>	<b>Continuations</b>	<b>105</b>
104	6.1	The concept of continuations . . . . .	105
105	6.1.1	Reduction semantics . . . . .	107
106	6.1.2	Continuation semantics . . . . .	108
107	6.1.3	First-class continuations . . . . .	110
108	6.2	Continuations by hand . . . . .	112
109	6.2.1	Defunctionalization and accumulators . . . . .	113
110	6.2.2	<b>Monad Cont</b> . . . . .	114
111	6.2.3	call/cc . . . . .	116
112	6.3	Delimited continuations . . . . .	116
113	6.3.1	Implementing the operators . . . . .	118

114	6.3.2	In <b>Monad Cont</b> . . . . .	119
115	6.4	Applications of continuations . . . . .	120
116	6.4.1	Everything through continuations . . . . .	120
117	6.4.2	The mother of all monads . . . . .	122
118	6.4.3	Generators and coroutines . . . . .	123
119	6.5	Efficient work with continuations . . . . .	124
120	6.5.1	Contiguous stack . . . . .	124
121	6.5.2	Segmented stack . . . . .	125
122	6.5.3	Finite state machine (FSM) . . . . .	125
123	<b>7</b>	<b>Effects and Modular Interpreters</b>	<b>126</b>
124	7.1	The notion of an effect . . . . .	126
125	7.2	Effect handlers . . . . .	128
126	7.2.1	Handlers via delimited continuations . . . . .	129
127	7.2.2	Efficient implementation of handlers . . . . .	129
128	7.2.3	Embedded handlers as explicit client-server communication . . . . .	129
129	7.2.4	Extensible messages and forwarding . . . . .	131
130	7.2.5	Free monads . . . . .	132
131	7.2.6	Handlers via free monads . . . . .	133
132	7.2.7	Applications of handlers . . . . .	134
133	7.2.8	Monad transformers . . . . .	135
134	7.2.9	Algebraicity and higher-order effects . . . . .	135
135	<b>8</b>	<b>Effect Systems TODO</b>	<b>137</b>
136	<b>9</b>	<b>Datatype-generic programming TODO</b>	<b>138</b>
137	9.1	Specialisation . . . . .	138
138	9.2	Ad-hoc solutions . . . . .	138
139	9.3	Approaches to datatype-generic programming in Haskell . . . . .	138
140	9.3.1	Template Haskell . . . . .	138
141	9.3.2	Deriving strategies . . . . .	138
142	9.3.3	GHC.Generics . . . . .	139
143	9.3.4	Sum Of Products . . . . .	139
144	9.3.5	Uniplate . . . . .	139
145	<b>10</b>	<b>Persistence and Optics TODO</b>	<b>140</b>
146	10.1	Persistent data structures . . . . .	140
147	10.2	The idea of optics . . . . .	141
148	10.3	Using optics . . . . .	143
149	10.4	Data optics . . . . .	143
150	10.4.1	Lenses . . . . .	143
151	10.4.2	Prisms . . . . .	144

152	10.4.3 Composition of lenses and prisms . . . . .	145
153	10.5 The path to profunctor optics . . . . .	146
154	10.6 Generating optics . . . . .	146

## 155 Introduction

156 Many complex concepts in language design and programming can be understood as special  
157 cases of some simple fundamental principles that are usually treated as common folklore and do  
158 not require extra explanation. However, the difficulty is that this knowledge is scattered across  
159 books, papers, and “cult” blog posts, and it takes a lot of time and effort to reconstruct the whole  
160 picture.

161 The goal of this course is to gather such folklore knowledge in one place and organize it into  
162 a coherent system. The course will explicitly rely on classical works studying the principles of  
163 language design and will help in studying them. Reading the referenced papers is an important part  
164 of independent work in the course.

165 By functional programming, as highlighted in the course title, we mean a careful attitude towards  
166 the notion of effect, which in FP — unlike in other schools of thought — is not considered an  
167 axiomatic given but a subject for study, conscious construction, and careful handling. This approach  
168 turns out to be very useful for learning languages and for building powerful language constructs;  
169 it is also the foundation of a productive programming style. Moreover, functional languages are  
170 relatively simple, which means new ideas and approaches often originate there and spread further.

171 Haskell is chosen as the main language of the course because, on the one hand, it embodies  
172 many concepts that are often brought to a certain logical completion and is powerful enough to  
173 encode others. On the other hand, it remains a practical industrial programming language.

174 Because of the breadth of the topic, this course is not always deep. Thus, implementation  
175 details in GHC or the category-theoretic foundations of some concepts may be given in general  
176 terms, without going into specifics. At the same time, in the realm of language design through the  
177 lens of functional programming, the course tries to be as detailed as possible.

178 Thus, this course may be useful for those who are interested in language design and elegant  
179 generalizations of programming concepts, want to improve their API design skills, or plan to work  
180 practically with functional languages.

181 A prerequisite for the course is knowledge of the basics of functional programming: algebraic  
182 data types, pattern matching, folds, parametric polymorphism, type classes, and basic monads. It is  
183 also useful to be able to read type judgments, and to be familiar with polymorphic lambda calculus  
184 and Church encoding.


185 **1 FP Recollections**

186 In this section we will recall the main concepts of functional programming and the Haskell  
187 language.

188 **1.1 Terms and Reduction**


189 In FP, programs are expressions. Program execution is the reduction of those expressions to  
190 simpler ones. Expressions can be represented either as a linear sequence of symbols or as a tree,  
191 which makes them easier to understand without requiring knowledge of auxiliary associativity rules  
192 and so on.

193 The simplest functional language is the lambda calculus. Expressions in it are called lambda  
194 terms, which consist of nodes of three kinds ( $V$  is the set of valid identifiers,  $\Lambda$  is the set of lambda  
195 terms):

196 **Variables**  $x \in \Lambda$ , if  $x \in V$  

197 **Abstraction**  $(\lambda x. M) \in \Lambda$ , if  $x \in V, M \in \Lambda$

198 **Application**  $(M N) \in \Lambda$ , if  $M \in \Lambda, N \in \Lambda$



199 In an arbitrary expression, some subterm can be replaced with a formal parameter that must be  
200 declared higher in the tree by a special lambda node. Instead of the formal parameter, we can later  
201 substitute various concrete arguments using the application node @ and reuse the expression for  
202 different purposes (for example, see Fig. 1). Reduction is defined by the following rewriting rule: an  
203 application of a lambda function to an argument is found, and the argument is substituted for all  
204 free occurrences of the variable bound by the lambda in the function body (Fig. 2).

205 **1.2 Types**

206 Software is hard. Therefore, bugs constantly and inevitably appear in programs. They can be  
207 found, among other ways, statically, that is, without running the program. One form of static  
208 analysis is type checking.

209 The idea of type analysis is that we attempt to assign a syntactic label to every node of the  
210 program tree according to certain rules. If every node can be labeled that way, then we say the




Figure 1: The expression represented with a lambda node is transformed into a one-argument function.




Figure 2: Reduction rewrites the tree by substituting a concrete argument for the formal parameter.




Figure 3: The tree corresponding to the expression  $(\lambda y. x \div (y + 2)) 3$ .

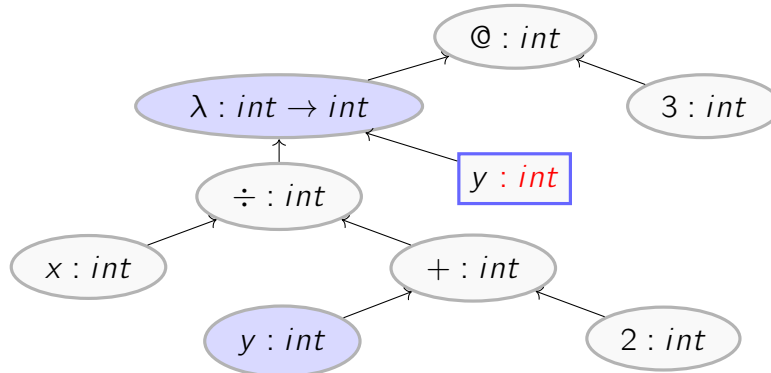


Figure 4: The tree for  $(\lambda y. x \div (y + 2)) 3$  after assigning type labels.

211 program type-checks and is “good”. For example, Fig. 3 shows an expression, and Fig. 4 shows the  
 212 same tree with labels assigned consistently with a type system.

213 **A type system** defines the syntax of type labels and the rules by which they can be assigned.  
 214 Syntax is usually described in a classical BNF-like notation, and rules are expressed as typing  
 215 judgments. For example, the rules for simply-typed lambda calculus look like this:

$$\frac{(x : \sigma) \in \Gamma}{\Gamma \vdash x : \sigma} \text{ ctx} \quad \frac{\Gamma \vdash M : \sigma \rightarrow \tau \quad \Gamma \vdash N : \sigma}{\Gamma \vdash M N : \tau} \text{ elim} \rightarrow \quad \frac{\{x : \sigma\} \cup \Gamma \vdash M : \tau}{\Gamma \vdash \lambda x^\sigma. M : \sigma \rightarrow \tau} \text{ intro} \rightarrow$$

216 Type labels are purely syntactic, but they can be interpreted. The most popular interpretation  
 217 is to treat a type label as a set. Thus, the label  $int \rightarrow int$  can correspond to the set of functions  
 218 between the sets of integers  $\mathbb{Z} \rightarrow \mathbb{Z}$ .

### 219 1.3 Functions in Haskell

220 At its core, Haskell is an extended typed lambda calculus enriched with primitive types, the  
 221 ability to declare new names, data structures, and type classes.

222 Examples of lambda abstractions in the GHCi REPL:

```
1 ghci> (\x -> x + 1) 4
2 5
```

223 You can learn the type of a function in the interpreter (in reality numbers are polymorphic, but  
224 we will discuss that later):

```
1 ghci> :t \x -> x + 1
2 \x -> x + 1 :: Int -> Int
```

225 Functions can be given names. Names can be annotated with types, and it is recommended to  
226 do so explicitly for top-level declarations in source files:

```
1 f :: Int -> Int
2 f x = x + 1
```

227 If a type name begins with a lowercase letter, it is not a concrete type but a type variable that can  
228 take different values depending on the call site. This capability is called **parametric polymorphism**.  
229 Thus, a function that simply returns its argument does not constrain the argument type. At the  
230 same time, the result type must match the argument type:

```
1 id :: a -> a
2 id x = x

3 ghci> :t id 5
4 id 5 :: Int
```

231 Functions can take other functions as arguments (such functions are called **higher-order functions**).  
232 A function name can consist of special symbols, in which case it is treated as an operator and can  
233 be used in infix form:

```
1 ($) :: (a -> b) -> a -> b
2 f $ x = f x
```

234 Example of a recursive function using guards to distinguish the base case:

```
1 factorial :: Int -> Int
2 factorial n
3   | n < 1 = 1
4   | otherwise = n * factorial (n - 1)
```

235 **Exercise 1** What does the command `ghci> :t uncurry (flip const)` return?

236 **Exercise 2** What does `ghci> :t first . first` return given

```
1 first :: (a -> a') -> (a, b) -> (a', b)
```

237 **Exercise 3** Implement factorial using the accumulator pattern.

## 238 1.4 Data in Haskell

239 Haskell has built-in support for declaring new data types from other types and constructing  
240 values of those types.

241 Define a data type describing animals:

```
1 data Animal
2   = Cat String Int
3   | Dog String
```

242 We defined the data type `Animal` and two ways to create values of this type: one for cats and one  
243 for dogs. `Cat` and `Dog` are **data constructors**. They are functions whose implementation is provided  
244 by the language. They allocate memory for instances of the type and place components positionally.  
245 Cats are described by a name and the remaining number of lives, while dogs are described only by  
246 a name.

```
1 Cat :: String -> Int -> Animal
2 Dog :: String -> Animal
```

247 To use the information stored in a data structure, it must be deconstructed using pattern  
248 matching. We match the value against a pattern. If the pattern resembles the way the value was  
249 constructed, it is selected among other alternatives and the variables declared in it bind to the  
250 corresponding positional contents of the data structure:

```
1 show :: Animal -> String
2 show animal = case animal of
3   Cat name nLives -> "This is cat " ++ name ++ " " ++ show nLives
4   Dog name -> "This is dog " ++ name
```

251 Haskell offers special syntax for declaring record fields with named labels.

```
1 data Penguin = Penguin { getName :: String, getAge :: Int }
2 penguin = Penguin { getName = "Andrey", getAge = 500 }
```

252 Haskell generates accessor functions for field access:

```
1 ghci> :t getName :: Penguin -> String
```

253 Often functions in programming are partial — for some argument values they return a result,  
254 and for others they do not. Let us model this with a dedicated data type. If there is a real result,  
255 we return it. If there is none, we return a specially reserved constant value of that type.

```
1 data MaybeD = NothingD | JustD Double
2 sqrt :: Double -> MaybeD
3 sqrt x = if x < 0 then NothingD else JustD (calcSqrt x)
```

256 One can notice that this would require defining a separate type `MaybeT` for each type `T`. Therefore  
257 Haskell allows type-level abstraction, just as we can abstract over values in terms.

```

1 data Maybe a = Nothing | Just a
2 sqrt :: Double -> Maybe Double
3 sqrt x = if x < 0 then Nothing else Just (calcSqrt x)

```

258 Note that `Maybe` is not a concrete type itself; it requires a type parameter to produce a concrete  
259 type. `Maybe` is called a **type constructor**.

260 With abstraction at the type level also comes application of types to types. What if a type  
261 constructor is given fewer parameters than expected? What if it is given more? Correctness of type  
262 applications is enforced by the **kind system**<sup>4</sup>. Kinds are the simplest “types for types”, syntactic  
263 labels that control whether the types written by the programmer are valid. Ordinary types have kind  
264 `*`. Type constructors have arrow kinds. For example, `Maybe :: * -> *`. Applying a type constructor  
265 to a type of the appropriate kind consumes one arrow:

```

1 ghci> :k Int
2 Int :: *
3 ghci> :k Maybe
4 Maybe :: * -> *
5 ghci> :k Maybe Int
6 Maybe Int :: *

```

266 In addition to brand-new data types, Haskell allows type synonyms. These are names that can  
267 be used instead of other types when the original type expression is too long to write everywhere:

```

1 type T a = VeryLongType Int (a -> AnotherLongType a)

```

268 If a data type has only one constructor and only one field, there is no need to allocate runtime  
269 memory containing a constructor tag and field references. In that case, the value of the new type  
270 can be represented by the wrapped value itself, leaving the new type present only at compile time  
271 and reducing runtime overhead. To declare such wrapper types, use `newtype` instead of `data`:

```

1 newtype CourseId = CourseId Int64
2 newtype ModuleId = ModuleId Int64

```

272 **Exercise 4** *Define the kind of the type constructor*

```

1 data Free f a = Pure a | Free (f (Free f a))

```

## 273 1.5 Type Classes in Haskell

274 Parametric polymorphism allows the same code to work for different input types. Type classes  
275 allow one identifier to refer to different implementations for different data types (which is analogous  
276 to overloading in other languages). Type classes are, as they say, the mechanism of **ad-hoc**  
277 **polymorphism**. For example, we can declare the symbol `==`, whose implementation depends on  
278 the chosen type `a`:

---

<sup>4</sup>In Russian literature, kinds are sometimes called “genera of types”, but we will not use that terminology.

```

1 class Eq a where
2   (==) :: a -> a -> Bool

```

279 For each type, one can declare its own implementation of Eq:

```

1 instance Eq CourseId where
2   CourseId x == CourseId y = x == y

3 instance Eq a => Eq [a] where
4   [] == [] = True
5   x:xs == y:ys = x == y && xs == ys

```

280 Now, depending on the concrete type `a` at the call site, the appropriate implementation for that  
281 type will be selected:

```

1 ghci> CourseId 1 == CourseId 2
2 False
3 ghci> [CourseId 1, CourseId 2] == [CourseId 1, CourseId 2]
4 True

```

282 Earlier, parametrically polymorphic functions could only do one of two things with their arguments:  
283 return them or pass them to other polymorphic functions. To do anything else, additional information  
284 about the type is needed, because otherwise there is no guarantee that the required operations are  
285 defined for a value of that type. For example, the function `suc n = n + 1` will not work for strings,  
286 because addition is not defined for them. Therefore it would be incorrect to assign it the polymorphic  
287 type `suc :: a -> a`.

288 Type classes, in addition to overloading, are also a mechanism for restricting polymorphism. We  
289 can explicitly specify that a function requires not any type as input, but any type for which the  
290 operations we need are defined. Thus, for the type of `suc` it is enough to constrain it with the  
291 presence of addition (the operation denoted by `+` is declared in the `Num` type class):

```

1 suc :: Num a => a -> a

```

292 **Exercise 5** *Implement a function that checks whether all elements of a list are equal.*

293 **Exercise 6** *Implement the semigroup instance for functions.*

294 **Exercise 7** *Implement equality checking for functions.*

## 295 1.6 Monads in Haskell

296 The type class `Functor` is declared for type constructors and allows replacing every element of  
297 one type in a container with elements of another type while preserving the container structure.

```

1 class Functor (f :: * -> *) where
2   fmap :: (a -> b) -> f a -> f b

3 instance Functor [] where
4   fmap :: (a -> b) -> [a] -> [b]
5   fmap _ [] = []
6   fmap f (x:xs) = f x : fmap f xs

```

298 In Haskell, every function simply produces a result of some type. However, in programming  
299 we often need functions that not only produce a result but also do something else. For example,  
300 modify some state or write to the console. In other words, they perform side effects. In any case, in  
301 Haskell we can only return a result from a function, so such side effects are encoded as additional  
302 structure wrapping the pure result. Thus, where a pure function would return a value of type *a*,  
303 after introducing effects it instead returns a *computation* of type *f a*.

- 304 • If a function can fail, then *f* = *Maybe*.
- 305 • If a function reads a global environment of type *e*, then *f* = *e -> \_*.
- 306 • If a function reads and updates state *s*, then *f* = *s -> (s, \_)*.

307 The Haskell standard library provides several type classes for working with values of the form *f*  
308 a. They allow you to abstract away from the structure *f* and work with values of type *a* inside as  
309 if there were no additional structure.

310 The first such type class allows writing expressions over computations *f a*.

```

1 class Functor f => Applicative (f :: * -> *) where
2   pure :: a -> f a
3   liftA2 :: (a -> b -> c) -> f a -> f b -> f c

4 instance Applicative Maybe where
5   pure :: a -> Maybe a
6   pure = Just

7   liftA2 :: (a -> b -> c) -> Maybe a -> Maybe b -> Maybe c
8   liftA2 _ Nothing _ = Nothing
9   liftA2 _ _ Nothing = Nothing
10  liftA2 f (Just x) (Just y) = Just (f x y)

```

311 The second type class allows sequential composition of computations in an imperative style:

```

1 class Applicative m => Monad (m :: * -> *) where
2   (>=) :: m a -> (a -> m b) -> m b

3 newtype State s a = State { runState :: s -> (s, a) }

```

```

4 instance Monad (State s) where
5   (>>=) :: State s a -> (a -> State s b) -> State s b
6   m >>= k = State \s ->
7     let (s', x) = runState m s in
8     runState (k x) s'

```

312 Now, if we define the basic state operations, we can write code in an imperative style with  
313 effects.

```

1 get :: State s s
2 get = State \s -> (s, s)

3 put :: s -> State s ()
4 put newS = State \oldS -> (newS, ())

5 example :: State Int Int
6 example =
7   get >>= \x ->
8   put 42 >>= \() ->
9   get >>= \y ->
10  pure (x + y)

11 ghci> runState example 1
12 43

```

314 For such monadic chains there is special syntactic sugar:

```

1 example :: State Int Int
2 example = do
3   x <- get
4   put 42
5   y <- get
6   pure (x + y)

```

315 **Exercise 8** Implement `liftA3` using `liftA2`.

316 **Exercise 9** Implement `>>=` using `join` and vice versa.

317 **Exercise 10** Read two numbers from the console, divide one by the other with integer division, and  
318 print the result; if the remainder is non-zero, print it too.

## 319 2 Parametric Polymorphism

320 No non-trivial property of programs can be checked algorithmically<sup>5</sup>. To remain decidable (in  
321 terms of type checking and/or inference), many type systems sacrifice completeness and, besides  
322 incorrect programs, reject many correct ones. At the same time, type systems also try to provide  
323 various capabilities that make it possible to type as many correct programs as possible. One of  
324 these is parametric polymorphism.

325 By **parametric polymorphism** we mean the ability of code to work uniformly with arbitrary  
326 data types Strachey [2000], Cardelli and Wegner [1985], which in many cases allows avoiding code  
327 duplication.

328 In this chapter we will look at how polymorphism is described in its simplest form — in typed  
329  $\lambda$ -calculus. We will study different forms of parametric polymorphism and the associated safe  
330 programming techniques. We will analyze possible ways to implement parametric polymorphism  
331 efficiently. And finally, we will consider runtime representation polymorphism, "polymorphism over  
332 polymorphism."

### 333 2.1 Parametric polymorphism in the language

334  $\lambda$ -abstraction allows generalizing expressions over values; each abstraction adds an arrow to the  
335 type of the expression, while application removes an arrow:

$$\frac{x : \tau, \Gamma \vdash M : \sigma}{\Gamma \vdash \lambda x : \tau. M : \tau \rightarrow \sigma} \text{Lam} \quad \frac{\Gamma \vdash M : \tau \rightarrow \sigma \quad \Gamma \vdash N : \tau}{\Gamma \vdash M N : \sigma} \text{App}$$

336 At the same time,  $\Lambda$ -abstraction allows generalizing expressions over types by adding a quantifier to  
337 the type (a  $\Pi$ -abstraction), and term application to a type (**universal application**) allows choosing  
338 which concrete type this term is intended to use [Pierce, 2002, chapter 23]:

$$\frac{\Gamma \vdash M : \tau}{\Gamma \vdash \Lambda \alpha. M : \forall \alpha. \tau} \text{TLam} \quad \frac{\Gamma \vdash M : \forall \alpha. \tau}{\Gamma \vdash M \sigma : [\alpha \rightarrow \sigma] \tau} \text{TApp}$$

339 Thus, for example, the function *id* actually takes *two arguments: a type and a value*:

$$\begin{aligned} id &: \forall \alpha. \alpha \rightarrow \alpha \\ id &= \Lambda \alpha. \lambda x : \alpha. x \\ id \text{ nat} &: \text{nat} \rightarrow \text{nat} \\ id \text{ nat } 42 &: \text{nat} \end{aligned}$$

340 In Haskell, type abstractions and applications are attached implicitly by type inference. However,  
341 there are language extensions that allow writing them explicitly: TypeAbstractions, TypeApplications.  
342 This can help, for example, when the term does not provide enough information to infer the type.  
343 Thus, you can explicitly specialize *id* to the desired type:

<sup>5</sup>[https://en.wikipedia.org/wiki/Rice%27s\\_theorem](https://en.wikipedia.org/wiki/Rice%27s_theorem)

```

1 id :: forall a . a -> a
2 ghci> :t id @Int
3 id @Int :: Int -> Int

```

344 Quantifiers are also implicitly attached at the beginning of a type, following a naming convention:  
345 concrete types start with an uppercase letter, while polymorphic ones start with a lowercase letter.  
346 Similarly, the user may explicitly write `forall`s using the `ExplicitForAll` extension. This may be  
347 useful either to manually fix the order of type abstractions, or to refer to the abstracted type inside  
348 the body of a function (the `ScopedTypeVariables` extension).

349 Polymorphic data types are expressed by another construct. If previously we controlled the type  
350 at the term level by universal application, now we want to control it at the type level. For this we  
351 introduce  $\lambda$  abstraction in types, type application, and, accordingly,  $\beta$ -reduction. The kind system  
352 (for now) is just the simplest "type system for types" and ensures well-formedness of types and  
353 strong normalization<sup>6</sup>. For example, we can write the type of a pair, abstracted over the concrete  
354 types of its components, so that the user can pick the desired instantiation:

$$\begin{aligned}
Pair &: * \rightarrow * \rightarrow * \\
Pair &= \lambda\tau^* \sigma^*. \forall\gamma. (\tau \rightarrow \sigma \rightarrow \gamma) \rightarrow \gamma \\
pair &: \forall\alpha \beta. \alpha \rightarrow \beta \rightarrow Pair \alpha \beta \\
pair &= \Lambda\alpha^* \beta^*. \lambda x^\alpha y^\beta. (\Lambda\gamma^*. \lambda f^{\alpha \rightarrow \beta \rightarrow \gamma}. f x y) \\
fst &: \forall\alpha \beta. Pair \alpha \beta \rightarrow \alpha \\
fst &= \Lambda\alpha^* \beta^*. \lambda p^{Pair \alpha \beta}. p \alpha (\mathbf{K} \alpha \beta)
\end{aligned}$$

355 In Haskell, the computational semantics of polymorphic types can be traced in type synonyms:

```

1 type Pair a b = forall c . (a -> b -> c) -> c
2 intPair :: Pair Int Int -- forall c . (Int -> Int -> c) -> c

```

356 Regular type constructors are nominal. For example, `(Int, Int)` or `Maybe Int` do not further  
357 compute.

358 Haskell does not allow writing functions at the type level in place using an explicit type lambda<sup>7</sup>  
359 due to the difficulty of this construct for type inference. However, named functions at the type  
360 level do exist, and we will examine them later 3.2. Scala has a non-trivial trick<sup>8,9</sup>, which makes this  
361 possible. Scala 3, however, has added this capability directly to the language<sup>10</sup>.

### 362 2.1.1 Emulating type abstractions and applications (Proxy)

363 In Haskell, extensions that allow manually writing type applications and abstractions appeared  
364 relatively recently<sup>11</sup>. Before that, the following technique was used.

<sup>6</sup>**Strong normalization** means any order of reductions leads to a normal form.

<sup>7</sup><https://stackoverflow.com/questions/4069840/lambda-for-type-expressions-in-haskell>

<sup>8</sup>(stackoverflow) Scala type lambdas.

<sup>9</sup><https://stackoverflow.com/questions/9443004/what-does-the-operator-mean-in-scala>

<sup>10</sup><https://docs.scala-lang.org/scala3/reference/new-types/type-lambdas.html>

<sup>11</sup>TypeApplications, TypeAbstractions.

365 The standard library defines the type `Proxy` with one parameter. This is a **phantom type**  
366 **parameter** — values of the corresponding type are not stored in the data structure, it only allows  
367 placing additional information at the type level<sup>12</sup>. Consequently, the uninformative constant `Proxy`  
368 can be annotated with the desired type and passed to a function to specialize the type parameter  
369 to the required type. Or you can accept a `Proxy` and use `ScopedTypeVariables` for type signatures  
370 in patterns<sup>13</sup>.

```
1 data Proxy a = Proxy

2 id :: Proxy a -> a -> a
3 ghci> :t id (Proxy :: Proxy Int)
4 id (Proxy :: Proxy Int) :: Int -> Int

5 id (Proxy :: Proxy [a]) x = (x :: [a])
```

371 Sometimes the proxy type is left polymorphic so as not to depend on `Data.Proxy`. Instead of  
372 a concrete value, one may pass a specialized value `undefined`, and the receiver, not knowing the  
373 type, cannot force it (however, any occurrence of `undefined` in a term is too suspicious, so this is  
374 not a very good practice).

```
1 id :: proxy a -> a -> a
2 id (_ :: proxy a) x = (x :: a)

3 ghci> :t id (undefined :: Proxy Int)
4 id (undefined :: Proxy Int) :: Int -> Int
```

## 375 2.1.2 First-class polymorphism

376 There is the possibility to write functions that accept other polymorphic functions as arguments.  
377 The types of such functions are called **higher-rank types**, they can be used with the extension  
378 `RankNTypes`. Thus, the type parameter of function `g` is determined by the function `f`, not by the  
379 caller of `f`:

```
1 f :: (forall a . a -> a) -> (Int, Char)
2 f g = (g @Int 42, g @Char 'a') -- universal application for clarity
3 ghci> f (\x -> x)
```

380 The problem with higher-rank types is that their inference is undecidable, meaning Haskell's  
381 global type inference does not work in this case. But if higher-rank types are annotated manually,  
382 the rest of the inference works as before. For example, Church numerals have higher rank<sup>14,15</sup>:

---

<sup>12</sup>[https://wiki.haskell.org/Phantom\\_type](https://wiki.haskell.org/Phantom_type)

<sup>13</sup>The type parameter actually has polymorphic kinds: `data Proxy (a :: k) = Proxy`, so that this technique works with types of arbitrary kinds (see also 2.1.5).

<sup>14</sup><https://okmij.org/ftp/tagless-final/course/Boehm-Berarducci.html>

<sup>15</sup>Anyone interested in programming languages is invited to spend more than a month of their life on Oleg Kiselyov's site: <https://okmij.org/ftp/README.html>.

```

1  suc :: (forall a . (a -> a) -> a -> a) -> (a -> a) -> a -> a
2  suc n s z = s (n s z)

```

383 **Exercise 11** What rank does the type `Int -> (forall a . a -> a)` have?

384 Many problems associated with higher-rank types can be avoided by wrapping them. For example,  
 385 for Church numerals we can define a wrapper `newtype Church`. Now code that works with the  
 386 wrapper can be typed with first-rank types; only the constructor has a higher-rank type:

```

1  newtype Church = Church (forall a . (a -> a) -> a -> a)
2  (+) :: Church -> Church -> Church -- rank 1

```

387 Code like that can also be written in Java (Kotlin):

```

1  interface Church { fun <a> fold(s: (a) -> a, z: a): a }
2  fun plus(n: Church, m: Church): Church = object : Church {
3      override fun <a> fold(s: (a) -> a, z: a): a = n.fold(s, m.fold(s, z))
4  }

```

388 By default, type parameters can only be instantiated to concrete types. The extension `ImpredicativeTypes`  
 389 allows specializing type parameters to polymorphic types (including `forall`s inside them) — **impredicative**  
 390 **application**.

```

1  runST :: (forall s. ST s a) -> a
2  ($) :: forall a b . (a -> b) -> a -> b
3  foo = runST $ ... -- only typechecks with ImpredicativeTypes

```

391 Higher-rank types can be used as type-based escape analysis, that is, to prevent the user  
 392 from passing some value outside a certain scope. For example, Haskell provides the efficient `ST`  
 393 monad, which within a restricted scope allows working with mutable memory cells Launchbury and  
 394 Peyton Jones [1995][Maguire, a, 7.2, ST trick]:

```

1  newtype ST s a = ST (IO a)
2  runST :: (forall s. ST s a) -> a

3  sumTo :: Int -> Int
4  sumTo n = runST do
5      ref <- newSTRef 0
6      forM [0..n] \i -> modifySTRef ref (+ i)
7      readSTRef ref

```

395 Notice that if you try to return a reference to a mutable cell from `runST`, the resulting type will  
 396 not pass the well-formedness check because it would contain the phantom parameter `s` that is not  
 397 bound anywhere:

```
1 newSTRef :: a -> ST s (Ref s a)
2 ghci> runST (newSTRef 0) :: Ref s Int -- error
```

398 In practice, to distinguish such locally bound type variables, one uses the concept of levels<sup>16</sup> Jones  
399 [2019].

400 Higher-rank types together with impredicative application form **first-class polymorphism**, when  
401 polymorphic types can be used almost as freely as any other. The classical Hindley-Milner global  
402 inference algorithm cannot handle this (and in general the problem is undecidable), so there is a  
403 variety of solutions that make different tradeoffs. You can make inference local, relying only on  
404 nearby AST nodes and auxiliary type annotations Pierce and Turner [2000], Christiansen [2013],  
405 Dunfield and Krishnaswami [2019]. Or you can try to help global inference with additional preprocessing  
406 (Quick Look<sup>17</sup> Serrano et al. [2020], recently implemented in Haskell) or with additional controlling  
407 constructs (FreezeML Emrich et al. [2020]).

### 408 2.1.3 Higher-order/kinded polymorphism

409 Haskell also allows abstraction over types of arbitrary kinds, not only `Type`, both in `data`  
410 declarations (**higher-order/kinded types (HKT)**<sup>18</sup>) and in polymorphic functions. We will encounter  
411 many examples later. For example, `Fix` has kind `(Type -> Type) -> Type`, and the catamorphism  
412 is abstracted over the arrow kind:

```
1 newtype Fix f = Fix (f (Fix f))
2 cata :: forall (f :: Type -> Type) a . Functor f => (f a -> a) -> Fix f -> a
```

413 Later we will study a technique that allows encoding higher-order kinds in a language that does  
414 not support them (see 3.1.8).

### 415 2.1.4 Generalized algebraic data types (GADTs)

416 Generalized algebraic data types (GADTs) allow attaching more information to data at the type  
417 level. As a model example, consider the syntax of a tiny programming language. Our goal is to  
418 prevent the construction of ill-typed syntax trees in Haskell.

```
1 data Expr = Const Int | IsZero Expr | If Expr Expr Expr
```

419 As we know, data constructors in Haskell are ordinary functions, except that their implementation  
420 is generated by the compiler (allocation, field layout, ...). Functions have types. For example,  
421 `IsZero :: Expr -> Expr`.

422 Haskell has syntax for defining `data` by explicitly listing constructor types<sup>19</sup>. It is completely  
423 analogous to the form we saw earlier, only much more convenient for complex data structures. The  
424 previously considered term type `Expr` looks like this:

---

<sup>16</sup><https://okmij.org/ftp/ML/generalization.html>

<sup>17</sup>(youtube) A Quick Look at Impredicativity (Simon Peyton Jones)

<sup>18</sup><https://serokell.io/blog/kinds-and-hkts-in-haskell>

<sup>19</sup>[https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/gadt\\_syntax.html#gadt-style](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/gadt_syntax.html#gadt-style)

```

1 data Expr where
2   Const :: Int -> Expr
3   IsZero :: Expr -> Expr
4   If :: Expr -> Expr -> Expr -> Expr

```

425 For polymorphic data structures, the list example uses the following syntax. The name `elem` is  
426 needed only for documentation and cannot be used otherwise; it only marks the presence of a type  
427 parameter and allows giving it a kind manually<sup>20</sup>.

```

1 data List (elem :: Type) where
2   Nil :: List a
3   Cons :: a -> List a -> List a

```

428 Let us add a phantom type parameter `ty` to `Expr`, indicating the Haskell type into which the  
429 expression should be interpreted, and use a GADT to assign concrete `ty` values to constructor result  
430 types. Thus we state that a program built with `Const` evaluates to an integer, `IsZero` evaluates  
431 to a boolean, and a conditional evaluates to the type of its branches.

```

1 data Expr ty where
2   Const :: Int -> Expr Int
3   IsZero :: Expr Int -> Expr Bool
4   If :: forall ty . Expr Bool -> Expr ty -> Expr ty -> Expr ty
5 eval :: Expr ty -> ty

```

432 Now we can write a safe typed interpreter. Note that when pattern matching on constructors,  
433 the information about the type parameter is refined:<sup>21</sup>

```

1 eval :: Expr ty -> ty
2 eval = \case
3   Const x -> x           -- ty           sim Int
4   IsZero t -> eval t == 0 -- ty           sim Bool
5   If c t e -> if eval c then eval t else eval e

```

434 Next we will examine how GADTs in Haskell are expressed through more basic language mechanisms 3.4.4.

### 435 2.1.5 Type-level structures, data promotion

436 To gain greater control over program correctness, let us learn to encode arbitrary data structures  
437 at the type level. As a model task, we will define a vector with length information.

438 First, define Peano-style natural numbers at the type level:

---

<sup>20</sup>The kind may be omitted. Alternatively, you can omit parameter names and simply write  
`data List :: Type -> Type where ...`

<sup>21</sup>This uses the convenient `LambdaCase` extension, which avoids introducing extra names.

```

1 data Zero
2 data Suc n

```

439 **Exercise 12** *How many inhabitants does the type `Suc (Suc Zero)` have?*

440 Now we can define a vector type that carries length information:

```

1 data Vec (size :: Type) (elem :: Type) where
2   VNil :: Vec Zero a
3   VCons :: a -> Vec n a -> Vec (Suc n) a

4 example :: Vec (Suc (Suc Zero)) Int
5 example = VCons 1 (VCons 2 VNil)

```

441 For such a type, for example, one can write a safe `vzip` function that only works on vectors of  
442 the same length:

```

1 vzip :: Vec n a -> Vec n b -> Vec n (a, b)
2 vzip VNil VNil = VNil -- n sim Zero
3 vzip (VCons x xs) (VCons y ys) = VCons (x, y) (vzip xs ys) -- n sim Suc n'

```

443 Note that in the other branches of `vzip`, equalities must arise that begin with different constructors,  
444 for example `Zero sim Suc n`. Since it is impossible to construct such function arguments, Haskell  
445 allows the corresponding branches to be omitted.

446 **Exercise 13** *Write a function that appends an element to the end of a vector. Proceed step by  
447 step, filling in type holes and tracking the equalities that arise.*

448 Currently the kind language that should control types is too poor. Indeed, the kind `Suc` is  
449 `Suc :: Type -> Type`, so nothing prevents writing, for example, `Suc (Maybe Int)`. At the same  
450 time it is too restrictive because it does not support polymorphism, which led to a lot of duplication  
451 like `Typeable (ty :: Type)`, `Typeable1 (ty :: Type -> Type)`, ...

452 Modern Haskell has the `TypeData` extension, which allows declaring new types and kinds similarly  
453 to how `data` allows declaring new types.

```

1 type data Nat = Zero | Suc Nat

```

454 Now the vector can be given a more precise kind:

```

1 data Vec (size :: Nat) (elem :: Type) where
2   VNil :: Vec Zero a
3   VCons :: a -> Vec n a -> Vec (Suc n) a

```

455 **Exercise 14** *What does `ghci> :k Vec print`?*

456 Another way to achieve the same result is using DataKinds Yorgey et al. [2012]. This extension  
 457 automatically promotes all `data` declarations one level higher. Specifically: any type constructor  
 458 also becomes a kind, and any data constructor becomes a type constructor. Thus, in the example  
 459 with natural numbers, we can declare them normally and use them at the type level:

```
1 data Nat = Zero | Suc Nat
2 ghci> :k Suc :: Nat -> Nat -- here it is clear that Suc is used as a type
```

460 Because types and terms live in different namespaces in Haskell, one can name type constructors  
 461 and data constructors identically. However, if such a data type is promoted, ambiguity arises: do  
 462 we mean the type or the promoted constructor? Haskell allows explicitly indicating that we mean  
 463 the promoted constructor with a leading apostrophe.

```
1 data T = T Nat
2 ghci> :k T
3 T :: Type          -- about the type constructor
4 ghci> :k 'T
5 'T :: Nat -> T    -- about the promoted data constructor
```

464 Not every `data` declaration is suitable for promotion, while `type data` declarations allow explicitly  
 465 requesting a type-level structure and getting sensible errors if the declaration is written incorrectly.  
 466 In the case of promoting a polymorphic type, we obtain polymorphic kinds (PolyKinds):

```
1 data [a] = [] | (:) a [a]
2 ghci> :k '(:)
3 '(:) :: forall k . k -> [k] -> [k]
```

Term	Type	Kind
<code>Zero</code>	<code>Nat</code>	<code>Type</code>
<code>[Zero, Suc Zero]</code>	<code>[Nat]</code>	<code>Type</code>
<code>[]</code>	<code>forall a. [a]</code>	<code>Type</code>
<code>(:)</code>	<code>forall a. a -&gt; [a] -&gt; [a]</code>	<code>Type</code>
	<code>'Suc 'Zero</code>	<code>Nat</code>
	<code>'['Zero, 'Suc 'Zero]</code>	<code>[Nat]</code>
	<code>'[Int, Double]</code>	<code>[Type]</code>
	<code>'[]</code>	<code>forall k. [k]</code>
	<code>'(:)</code>	<code>forall k. k -&gt; [k] -&gt; [k]</code>

Figure 5: An example of promotions in Haskell.

467 Examples of promotions for various constructions can be seen in Table 5.  
 468 As an example, let us define a heterogeneous list indexed by the types of its elements:

```

1 data HList (tys :: [Type]) where
2   HNil :: HList '[]
3   HCons :: ty -> HList tys -> HList (ty ': tys)

4 example :: HList '[Int, Bool, Double]
5 example = HCons 42 $ HCons True $ HCons 12.5 HNil

```

469 Data structures can also be kind-polymorphic. Consider the type `Tagged`, which allows attaching  
 470 an additional type tag to the value type. The kind of the tag may be arbitrary, so for example we  
 471 can use type-level constants from `TypeLits` (another example of polykind usage we saw earlier in  
 472 2.1.1):

```

1 newtype Tagged (tag :: k) (a :: Type) = Tagged a
2 ghci> :t Tagged
3 Tagged :: forall k (tag :: k) a. a -> Tagged tag a

4 example :: Tagged ("dbId" :: Symbol) Int
5 example = Tagged 42

```

473 Modern Haskell has ultimately arrived at the point where the type system does not distinguish  
 474 between types and kinds (Fig. 6). In particular, `Type :: Type`. This is required to extend Haskell  
 475 in the direction of dependent-type programming by adding non-syntactic equivalences for kinds  
 476 (`TypeInType`). *System FC* was presented in the work Weirich et al. [2013]<sup>2223</sup>.

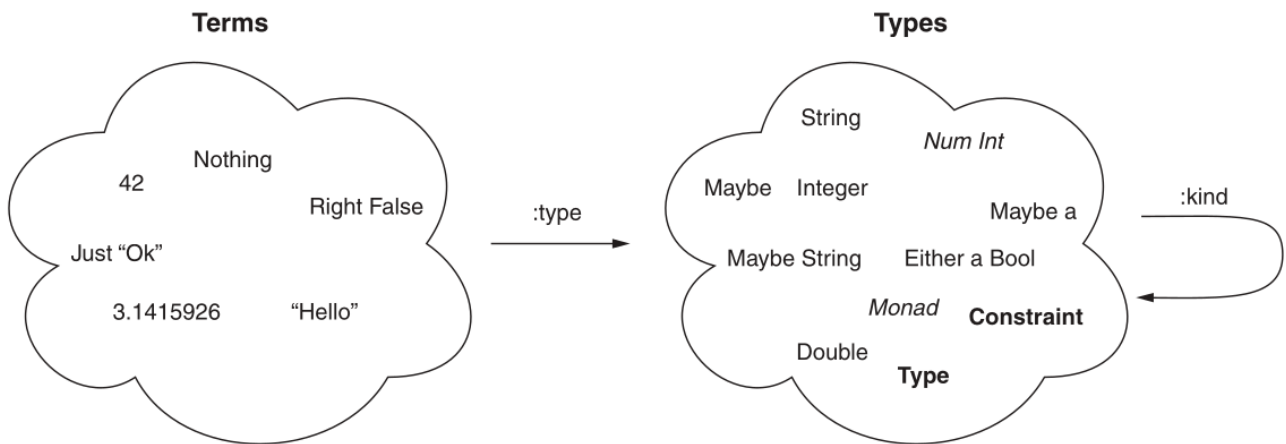


Figure 6: Types and kinds are the same thing (Bragilevsky).

<sup>22</sup>(youtube) Mini-course in Russian on the development of Haskell toward dependent typing.  
<sup>23</sup>(youtube) Mini-course in Russian — Haskell type inference system.

## 477 2.2 Implementation of parametric polymorphism

478 A **calling convention**<sup>24</sup> is a set of agreements between how a function is compiled and how  
479 it must be called. For example, a function takes two word-sized arguments and returns one word-  
480 sized result. Then the generated low-level code may, for example, expect both arguments to be  
481 passed through a specific pair of registers and place the result in a third register. In that case, the  
482 caller must provide the arguments in the correct registers and expect the result in a predetermined  
483 register.

484 In general, the calling convention of a function depends on the types of its arguments and  
485 result. You need to know at least their size to decide whether to place them in registers or on the  
486 stack. You need to know whether it is a pointer (**reference type**) or a value itself (**value type**) to  
487 understand how to handle it. For data structures, you need to know field offsets.

488 Thus, implementing parametric polymorphism in a language is not a trivial task. Different  
489 languages use different approaches, each with its own strengths and weaknesses.

### 490 2.2.1 Monomorphization

491 **Monomorphization** is the most straightforward approach: compile polymorphic functions and  
492 structures for each set of type arguments. Thus, if there are 100 different sets of type arguments  
493 with which a function is called, then its code will be compiled 100 times and occupy 100 times more  
494 space in the binary. This is what C++ and Rust do, for example.

495 In fact it is even worse. If the project is multi-module and consists of many compilation units  
496 (pieces that are compiled separately), the same specialization of the function for type arguments  
497 will be compiled again in every compilation unit where it is needed. Then the linker will remove  
498 duplicates, which is also not the fastest or most efficient process.

- 499 + Generated code is maximally efficient for each type;
- 500 + is-style runtime type checks are easy to implement at compile time (in other approaches this  
501 is difficult);
- 502 – Compilation time becomes very large;
- 503 – The resulting binary size increases significantly, which may be critical for some applications;
- 504 – It may work inefficiently due to instruction cache pollution;
- 505 – Polymorphic methods cannot exist in interfaces, because at the call site we do not know which  
506 concrete implementation should be specialized (similarly, higher-rank polymorphism does not  
507 work);
- 508 – Polymorphic functions cannot be dynamically linked in general (they have no code before  
509 specialization);
- 510 – In general, variance cannot be supported because the code is compiled for a concrete type  
511 and cannot in general work for an arbitrary subtype or supertype (if reference and value types  
512 may live in the same subtyping hierarchy).

---

<sup>24</sup>[https://en.wikipedia.org/wiki/Calling\\_convention](https://en.wikipedia.org/wiki/Calling_convention)

513 Some languages (for example, C++ and Zig) do not treat instantiation as a hidden implementation  
514 detail of the language, but provide it as a tool for users. This provides the following possibilities:

- 515 ● If values are allowed in types, instantiation can be used as a compile-time computation  
516 mechanism.
- 517 ● If error checking is deferred to instantiation time, we get a kind of static duck typing. This  
518 allows not describing complex signatures of polymorphic functions. However, then functions  
519 for testing must be manually instantiated against all possible types, otherwise it is impossible  
520 to determine statically whether they compile for those types.

## 521 2.2.2 Type erasure

522 One can do the opposite: unify the values that arrive at polymorphic functions and are stored  
523 in polymorphic data structures, instead of compiling code for each type.

524 Suppose each value is allocated on the heap and passed by pointer. Then we can reuse the same  
525 code for different type arguments — it will simply expect pointers.

- 526 + Each function is compiled exactly once — fast;
- 527 + It is possible to dynamically load new polymorphic functions and types and use them together;
- 528 + Flexibility — variance, polymorphic methods in interfaces, higher-rank types, etc. just work;
- 529 − Heap allocation and pointer dereferencing can significantly slow down code;
- 530 − Since type information is erased, one cannot do anything with a type argument without having  
531 its inhabitants (for example, query reflection information or perform an `is` check).

532 This approach is followed by the JVM, Haskell, and generally other functional languages due to  
533 its flexibility and compile-time speed.

534 The special problem is working with primitives and other value types, because each value has to  
535 be boxed first and then used in a polymorphic context. Therefore languages fight this problem as  
536 best they can. Some languages reduce the range of primitive values to reserve a bit that indicates  
537 whether it is a pointer or a value. Code consults this bit to operate (similar to 2.2.4). This is  
538 done, for example, by OCaml and Koka. Aggressive inlining together with other optimizations also  
539 helps Blanchet [1999]. Java, for example, is also trying to move carefully toward the possibility of  
540 monomorphization<sup>2526</sup>.

## 541 2.2.3 Hybrid approach

542 C# implements a hybrid approach<sup>27</sup>. It distinguishes values stored on the heap — reference types  
543 — from values stored on the stack — value types. For the former it generates one specialization  
544 that works with pointers. For each set of value types it lazily generates specializations at runtime.

545 Thus, traces of generics exist both in the CIL intermediate representation and in the runtime.

---

<sup>25</sup>Type Specialization of Java Generics - What If Casts Have Teeth?

<sup>26</sup><https://cr.openjdk.org/~jrose/values/parametric-vm.html>

<sup>27</sup>Generics in the runtime (C# programming guide).

```

struct value_witness_table {
    size_t size, align;
    void (*copy_init)(opaque *dst, const opaque *src, type *T);
    void (*copy_assign)(opaque *dst, const opaque *src, type *T);
    void (*move_init)(opaque *dst, opaque *src, type *T);
    void (*move_assign)(opaque *dst, opaque *src, type *T);
    void (*destroy)(opaque *val, type *T);
};

```

Figure 7: Swift value witness table.

Example:

```

func f<T>(_ t: T) -> T {
    let copy = t
    return copy
}

```

Implementation:

```

void f(opaque *result, opaque *t, type *T) {
    opaque *copy = alloca(T->vwt->size);
    T->vwt->copy_init(copy, t, T);
    T->vwt->move_init(result, copy, T);
    T->vwt->destroy(t, T);
}

```

Figure 8: Code generated by the Swift compiler for a polymorphic function.

- 546 + Value types are stored and passed as-is without boxing;
- 547 + Reflection over generics is available;
- 548 + Compilation time is small;
- 549 - Runtime instantiation slows execution;
- 550 - Variance works only for reference types (which is odd — there is “correct” subtyping and
- 551 “incorrect” subtyping).

## 552 2.2.4 Using a value witness table for types

553 Swift<sup>28</sup> passes with each type parameter a value witness table (Fig. 7). This is a table with all the  
554 necessary information about the type: size and alignment, what to do when copying and moving  
555 the object (for example, incrementing a reference count). Thus, the generated code constantly  
556 consults this table and makes virtual calls to functions from it (Fig. 8).

- 557 + Compilation time is small;
- 558 + The performance behavior is predictable (it does not lead to surprising runtime pauses);
- 559 + Efficient work with value types;
- 560 + High flexibility;
- 561 + Type information is not erased;

<sup>28</sup>(youtube) 2017 LLVM Developers' Meeting: "Implementing Swift Generics"

562 — A significant constant overhead from virtual calls through the table, and performance strongly  
563 depends on compiler optimizations.

564 This is a kind of implementation of parametric polymorphism through a special ...

## 565 2.3 Representation polymorphism

566 As we already discussed above 2.2.2, parametric polymorphism in Haskell is implemented as  
567 follows: all values are stored on the heap and passed to polymorphic functions by pointer. However, if  
568 performance is important for computational code, that approach is not suitable because of the heavy  
569 load on the memory management subsystem and many indirections. Therefore Haskell also allows  
570 writing code using unboxed values. And if the calling convention is not essential, one can abstract  
571 over it and write one code path for both boxed and unboxed values Eisenberg and Peyton Jones  
572 [2017].

### 573 2.3.1 Varieties of runtime representations in Haskell

	Boxed	Unboxed
Lifted	<i>Int</i> <i>Bool</i>	
Unlifted	<i>ByteArray#</i>	<i>Int#</i> <i>Char#</i>

Figure 9: Kinds of values in Haskell with examples Eisenberg and Peyton Jones [2017].

574 Figure 9 shows the classification of values in Haskell with examples of types. **Unboxed types** are  
575 held and passed by value. **Boxed types**, conversely, are passed by pointer and stored on the heap.  
576 The ordinary `Int` is just a declaration of the following form, where `I#` is an ordinary constructor  
577 with an unusual name that contains an unboxed value.

```
1 data Int = I# Int#
```

578 **Lifted types** may contain  $\perp$  as a value. In other words, they may represent suspended computations  
579 (the compiler and runtime provide this specially). **Unlifted types**, in contrast, cannot be suspended.  
580 Operations producing values of unlifted types are always strict. The property lifted/unlifted is called  
581 **levity**. To extend the exposition to eager languages, one can replace levity with boxity and everything  
582 remains valid.

583 A hash sign in type and function names is a convention that indicates that unlifted values are  
584 involved nearby<sup>29</sup>.

---

<sup>29</sup>You need to enable the MagicHash extension to use `#` in identifiers.

585 Haskell also has unboxed tuples, which do not exist as runtime objects. For example, the following  
586 function appears to return a pair of values, but in fact the compiler may place them, for example,  
587 in a pair of registers. Accordingly, pattern matching on such tuples simply allows referring to each  
588 of those values.

```
1 divMod# :: Int -> Int -> (# Int, Int #)
2 case divMod# n k of (# quot, rem #) -> ...
```

589 Consequently, there is no distinction between differently nested unboxed tuples:

```
1 (# A, (# B, C #) #) ≡ (# (# A, B #), C #) ≡ (# A, B, C #)
```

## 590 2.3.2 Classification of values by runtime representation

591 Values of different types may be arranged differently at runtime. That is, we need some system  
592 for classifying types. But such a system already exists in Haskell — kinds. We describe the domain  
593 as data structures and then promote it to the required level using DataKinds 2.1.5.

594 The standard Haskell library provides the following data types:

```
1 TYPE :: RuntimeRep -> Type
2 data Levity = Lifted | Unlifted
3 data RuntimeRep = BoxedRep Levity
4                 | IntRep | DoubleRep
5                 | TupleRep [RuntimeRep]
6                 | SumRep [RuntimeRep]
7                 | ...
8 type LiftedRep = BoxedRep Lifted
9 type Type = TYPE LiftedRep
```

595 `TYPE` is a magic type defined in the compiler. It is parameterized by the runtime representation  
596 of values. Now the familiar `Type` is simply a special case with boxed lifted values.

- 597 ● `Int :: TYPE (BoxedRep Lifted)` or `:: Type`
- 598 ● `IntRep` and `DoubleRep` correspond to the runtime representation of numeric constants (depending  
599 on the processor architecture, integers and floating-point numbers may need to be placed in  
600 different special registers)  
601 `Int# :: TYPE IntRep`
- 602 ● `Maybe Int :: Type`
- 603 ● `Maybe :: Type -> Type`

604 • `TupleRep` and `SumRep` are unboxed algebraic representations, parameterized by the representations  
605 of their stored values

```
606 (# Int, Bool #) :: TYPE (TupleRep '[LiftedRep, LiftedRep])
```

607 • For simplicity, nested tuple types are not unified

```
1 (# Int#, (# Int, Double# #) #)
```

```
2 :: TYPE (TupleRep '[IntRep, TupleRep '[LiftedRep, DoubleRep]])
```

### 608 2.3.3 Representation polymorphism

609 Having exposed the runtime representation in the kind structure, we can now parameterize over  
610 it. For example, the kind of the function arrow looks as follows<sup>30</sup>:

```
1 ghci> :k (->)
```

```
2 (->) :: forall {q :: RuntimeRep} {r :: RuntimeRep}. TYPE q -> TYPE r -> Type
```

611 **Exercise 15** *Think about why the function has a boxed type. Can it be different? Could that be*  
612 *useful?*

613 Unfortunately, Haskell imposes a rather strict restriction: binders cannot have a type that is  
614 polymorphic over runtime representation. One can easily guess why — it is impossible to generate  
615 code for a function that works with a parameter of an arbitrary runtime representation. This can  
616 only be solved by monomorphization 2.2.1, but Haskell avoids that approach<sup>31</sup>. The community is  
617 also trying to find other solutions<sup>32</sup> (something like 2.2.4).

618 For example, initially the application operator was generalized only over the result type. This does  
619 not cause problems, because the caller can infer the representation and generate the appropriate  
620 code:

```
1 ($) :: forall r a (b :: TYPE r). (a -> b) -> a -> b
```

```
2 f $ x = f x
```

621 However, it was noticed that for the application operator one can obtain another implementation  
622 that does not use a levity-polymorphic binder<sup>33</sup>:

```
1 ($) :: forall ra rb (a :: TYPE ra) (b :: TYPE rb). (a -> b) -> a -> b
```

```
2 ($) f = f
```

623 Thus, in Haskell representation polymorphism is somewhat degenerate and helps only in a small  
624 number of cases, but not unimportant ones. If one allows monomorphization over `RuntimeRep`  
625 parameters, one gets a system similar to the hybrid implementation of parametric polymorphism 2.2.3,  
626 only with greater programmer control over monomorphization.

---

<sup>30</sup>Disable simplifications: `ghci> :set -fprint-explicit-foralls -fprint-explicit-runtime-reps`

<sup>31</sup><https://gitlab.haskell.org/ghc/ghc/-/issues/14917>

<sup>32</sup><https://mail.haskell.org/pipermail/haskell-cafe/2023-January/135770.html>

<sup>33</sup>[https://gitlab.haskell.org/ghc/ghc/-/merge\\_requests/10131](https://gitlab.haskell.org/ghc/ghc/-/merge_requests/10131)

## 627 3 Ad-hoc Polymorphism

628 Once, Joe Fasel in a conversation with Philip Wadler expressed the idea that function overloading  
629 should be reflected in types. Wadler misunderstood him Hudak et al. [2007]. But what he understood  
630 turned out to be type classes Wadler and Blott [1989].

631 Christopher Strachey introduced a classification of polymorphism into two categories Strachey  
632 [2000]. Parametric — the same code works with data of different types. **Ad-hoc polymorphism** —  
633 code is selected depending on the type. For example, the same multiplication symbol acts differently  
634 on integers and on floating-point numbers.

635 Overloading in languages refers to the ability to give the same name to several functions with  
636 different sets of input parameters. At the call site, the compiler statically determines by the argument  
637 types which one should actually be called.

```
1 string toString(x: int) { ... }
2 string toString(fmt: String, d: double) { ... }
```

638 Type classes require first declaring a named entity (the type class itself) that includes a set of  
639 function declarations that can be overloaded for different types.

```
1 class Show a where
2   show :: a -> String

3 instance Show Int where
4   show :: Int -> String
5   show = ...
```

640 It is important to note that a type class declaration contains a formal type parameter, and by its  
641 occurrences in the function type, the overloading is selected. There can be many such parameters  
642 (MultiParamTypeClasses), and they can have arrow kinds. For example, in the case of the type  
643 class **Applicative**, the choice of implementation for the operation **pure** happens based on the  
644 type constructor of the result, i.e. not even on a full type.

```
1 class Functor f => Applicative (f :: Type -> Type) where
2   pure :: a -> f a
3   ...

4 instance Applicative Maybe where
5   pure :: a -> Maybe a
6   ...
```

645 Also, unlike overloading, type classes are compatible with parametric polymorphism. In the type  
646 of a polymorphic function, one cannot specify that a certain overloading must exist for a type.  
647 Type classes, however, allow constraining the set of possible type arguments to those for which an  
648 instance of the required type class is implemented:

```
1 showPrefixed :: Show a => a -> String -> String
```

649 Comparing type classes with overriding in OOP languages, the resolution of a virtual function  
650 call uses a table stored in the object of the first parameter (the call receiver). Type classes, on the  
651 other hand, rely exclusively on the type, so, for example, it is possible to define constants in type  
652 classes<sup>34</sup>:

```
1 class Enum a => Bounded a where  
2   minBound :: a  
3   maxBound :: a
```

653 At the same time, type classes are not types, but rather predicates on types. A type satisfies  
654 such a predicate, or property, if there is a corresponding instance for it. Therefore, in particular,  
655 the familiar OOP way of creating a heterogeneous collection of elements with a common interface  
656 does not work directly with type classes. For example, this type would not be valid: [Show]. We  
657 will return to this problem in 3.4.2.

## 658 3.1 Type classes in the language

659 Despite its remarkable power, the idea of implementing type classes is extremely simple. It was  
660 already fully presented in the original paper Wadler and Blott [1989]. Subsequent work refined the  
661 type inference mechanism by reducing it to the classical Hindley-Milner type system Hall et al.  
662 [1996]. Other work mostly proposes a huge variety of extensions and applications Jones et al.  
663 [1997].

### 664 3.1.1 Dictionaries

665 Let us consider the idea of implementing type classes using the example of polymorphic sorting.  
666 Writing a sort for a list of elements of a specific type is trivial:

```
1 sort :: [Int] -> [Int]  
2 sort = \case [] -> []; x:xs -> insert x (sort xs)  
3   where  
4     insert x xs = let (l, r) = List.partition (< x) xs in l ++ x : r
```

667 In the implementation, the only type-specific information we use is the ordering on its inhabitants.  
668 Thus, when moving to polymorphic sorting, we need to accept a dictionary with predicates defining  
669 the required ordering for the given type.

```
1 data OrdDict a = OrdDict { less :: a -> a -> Bool }  
2 sort :: OrdDict a -> [a] -> [a]
```

---

<sup>34</sup>Modern OOP languages, however, strive to support static functions in interfaces, which makes them closer to type classes and allows doing similar things. For example, Swift.

```

3  sort d@OrdDict{ less } = \case [] -> []; x:xs -> insert x (sort d xs)
4  where
5  insert x xs = let (l, r) = List.partition (['less' x) xs in l ++ x : r

```

670 Now, to use sorting on a list of numbers, we need to construct the appropriate record and call  
671 the function with it on a list of specific types:

```

1  intOrd :: OrdDict Int
2  intOrd = OrdDict { less = (<) }

3  ghci> sort intOrd [3, 2, 1]

```

672 There can be a situation where an instance for one type depends on an instance for another.  
673 For example, ordering on lists can be derived automatically, knowing the ordering on elements. In  
674 the case of dictionaries, we model this as a function between dictionaries:

```

1  listDict :: OrdDict a -> OrdDict [a]
2  listDict d = OrdDict { less = ... ['less d'] ... }

```

675 Now we can sort a list of lists by constructing the appropriate dictionary:

```

1  ghci> sort (listDict intOrd) [[3, 2], [2, 1], [0]]

```

676 A comparison of explicit dictionary passing and type classes can be seen in the following table:

1. Definition of a dictionary of functions	1. Definition of a type class
<pre> 1  data MyOrd a = MyOrd 2  { less :: a -&gt; a -&gt; Bool } </pre>	<pre> 1  class MyOrd a where 2  less :: a -&gt; a -&gt; Bool </pre>
2. Dictionary instance for a specific type	2. Declaring a type as a member of a type class
<ul style="list-style-type: none"> <li>• Named value</li> </ul>	<ul style="list-style-type: none"> <li>• Has no name</li> </ul>
<pre> 1  intMyOrd :: MyOrd Int 2  intMyOrd = MyOrd { less = (&lt;) } </pre>	<pre> 1  instance MyOrd Int where 2  less = (&lt;) </pre>
3. Explicit function parameter	3. Implicit function parameter
<pre> 1  sort :: MyOrd a -&gt; [a] -&gt; [a] </pre>	<pre> 1  sort :: MyOrd a =&gt; [a] -&gt; [a] </pre>
4. Passed by the user	4. Passed by the compiler
<pre> 1  test = sort ['intMyOrd'] [3, 2, 1] </pre>	<pre> 1  test = sort [3, 2, 1] </pre>

678 Thus, a dictionary is a **witness** or proof that the type satisfies the constraint.

679 **Exercise 16** *What dictionary would correspond to the higher-kinded type class `Functor`?*

### 680 3.1.2 Implicit arguments

681 One can think of it as follows: to the left of `=>`, implicit function arguments are passed, inferred  
682 by the compiler from the context. That is, for example, one should not be surprised by occurrences  
683 of `=>` in the type of an argument — it is simply a function with an implicit argument. Thus, the  
684 following code will not compile because at the use site of variable `y` there is no value of type `Show b`:

```
1 f :: (Show b => b) -> b
2 f x = [x] -- error
```

685 We can accept this value in function `f`, and then it will automatically propagate to `x`:

```
1 f :: Show b => (Show b => b) -> b
2 f x = [x]
```

686 The extension `ImplicitParams` provides the ability to make some function arguments implicit.  
687 In fact, this is an implementation of dynamic binding in a static language Lewis et al. [2000]  
688 (see 5.3.1). Implicit arguments are taken from the scope by name and substituted automatically:

```
1 sortBy :: (a -> a -> Bool) -> [a] -> [a]
2 sort :: (?cmp :: a -> a -> Bool) => [a] -> [a]
3 sort = sortBy ?cmp
```

689 Haskell also provides the ability to store dictionaries in data structures:

```
1 data ShowDict a where
2   ShowDict :: Show a => ShowDict a
3 f :: ShowDict b -> (Show b => b) -> b
4 f d x = case d of ShowDict -> [x] -- the Show b instance is available in scope
```

690 **Exercise 17** *Is exactly this semantics possible in a strict language? Why?*

### 691 3.1.3 Instance inference

692 To call a constrained-polymorphic function, GHC performs instance inference, or in other words,  
693 automatically constructs witnesses. Instance inference is tightly integrated with Haskell's overall  
694 type inference system Peyton Jones [2019].

695 In reality, instance inference is nothing other than *the type inhabitation problem*. Indeed, after  
696 translation to Core (the intermediate representation in GHC), type classes are represented as  
697 dictionaries of functions. We have in context specific dictionaries and functions that allow deriving  
698 new dictionaries from existing ones. We need to find a term that constructs a dictionary of the  
699 required type.

700 Suppose, for example, that inside function `f :: Show a => ..` there is a call to a constrained-  
701 polymorphic function `g :: Show [a] -> ...`. That is, we have a dictionary `d1 :: ShowDict a`,

702 as well as a function `d2 :: ShowDict a -> ShowDict [a]`, coming from imports<sup>35</sup>. We need to  
703 construct a term of type `ShowDict [a]`. Obviously, it is simply the application of one to the other:  
704 `d2 d1`.

705 Instance inference proceeds recursively. To infer `ShowDict [a]`, the premise `ShowDict a` is first  
706 inferred. That is, we get recursion on the structure of the type. In other words, instance inference  
707 can be exploited as a type-level computational primitive. For example, we can lower information  
708 from types into terms (analogous to `GHC.TypeLits`):

```
1  type data Nat = Zero | Suc Nat
2  class KnownNat (n :: Nat) where
3    natVal :: Int
4  instance KnownNat Zero where
5    natVal = 0
6  instance KnownNat n => KnownNat (Suc n) where
7    natVal = 1 + natVal @n
8  ghci> natVal @(Suc (Suc Zero))
9  -- will be inferred as natVal {knownSuc (knownSuc knownZero)}
```

709 In the general case, the type inhabitation process, as one might expect by computational analogy,  
710 is undecidable. Therefore GHC imposes many restrictions on the form of instances that guarantee  
711 totality of inference. These restrictions are described in detail in Sulzmann et al. [2007a]. GHC also  
712 provides various extensions that relax these restrictions, shifting part of the responsibility to the  
713 programmer<sup>36</sup>. For example, with `UndecidableInstances` one can easily write type-level list reversal  
714 at compile time, as well as any other function:

```
1  class Reverse (acc :: [Type]) (tys :: [Type]) where
2    showReverse :: String
3  instance ShowT acc => Reverse acc '[] where
4    showReverse = showTypes @acc
5  instance Reverse (ty : acc) tys => Reverse acc (ty : tys) where
6    showReverse = showReverse @(ty : acc) @tys
7  ghci> showReverse @'[] @'[Char, Int, Double]
```

715 One can notice that the process of type class inference is very similar to the evaluation of logic  
716 programs, for example in Prolog, only without backtracking (trying different solution variants in

---

<sup>35</sup>Instances can be imported with an empty import: `import Module ()`.

<sup>36</sup>[https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/instances.html](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/instances.html)

717 search of a suitable one). As, indeed, type inference in Haskell Peyton Jones [2019] in general: the  
718 equivalences collected across the program can be viewed as a logic program, and solving this system  
719 of type equations as executing that program.

720 There is an interesting synergy between type classes and type inference (Fig. 10)<sup>37</sup>. Based on  
721 terms, types are inferred. Then, based on types, type class instances are inferred. That is, we write  
722 some interesting intellectual code, and in parallel the compiler writes out the uninteresting code.




Figure 10: Type classes + instance inference = code generation.

723 Instance inference relies only on the form of the “head” of the declaration — to the right of  
724 `=>`, while constraints on the left are applied after the fact. This can be used to write more general  
725 instances. For example, this is how the **constraint trick**<sup>38</sup> works, allowing one to resolve ad-hoc  
726 polymorphic functions in a parametrically-polymorphic context.

### 727 3.1.4 Constructing a type from a value

728 After we have learned to lower values from types, it is natural to learn the reverse — lifting values  
729 into types. We use the technique described in Kiselyov and Shan [2004]. There is a corresponding  
730 library `Data.Reflection` (see 3.4.7).

731 In reality, we of course cannot honestly obtain a syntactically correct type of the needed size,  
732 simply because types exist strictly before the execution stage. However, as we know, type class  
733 dictionaries have a runtime representation (the case of polymorphic recursion is precisely an example  
734 where this cannot be entirely avoided). Therefore, we use continuation passing style, which will be  
735 examined in detail later in chapter 6: instead of returning a result, we accept a continuation that  
736 can work with any type with `KnownNat` (we use type abstractions and applications, see 2.1)<sup>39</sup>:

```
1 reify :: Int -> (forall n. KnownNat n => a) -> a
2 reify n k
3   | n <= 0 = k @Zero
4   | otherwise = reify (n - 1) \@n' -> k @(Suc n')
```

<sup>37</sup>(youtube) Hackett: a metaprogrammable Haskell.

<sup>38</sup><https://chrisdone.com/posts/haskell-constraint-trick/>

<sup>39</sup>In not the most recent versions of GHC, you will need to use the `Proxy` technique from 2.1.1.

737 The continuation passed to the recursive call captures the dictionary for type `n` and constructs  
738 a dictionary for `Suc n`.

739 Finally, we can write the following surprising identity function that first lifts a value into a type  
740 and then lowers the type back into terms:

```
1 wonderId :: Int -> Int
2 wonderId n = reify n (\@t -> natVal @t)
```

### 741 3.1.5 Implicits and coherence

742 Type classes need not be a special language mechanism; instead, one can provide at the language  
743 level implicit parameters and inhabitation sufficient to implement type classes.

744 Thus, in Scala there is the mechanism of implicits Křikava et al. [2019]<sup>40</sup>. Function parameters  
745 can be marked with the keyword `implicit`, and Scala will try to infer them automatically using  
746 the `implicit` declarations available in scope. Declarations of variables, functions, and object  
747 constructors can also be marked `implicit`, and they will be used during inhabitation. Now we  
748 can model a dictionary of functions, for example using interfaces (which in Scala are called `trait`)  
749 and OOP singletons, to obtain type classes Oliveira et al. [2010]:

```
1 // A bundle of functions.
2 trait Show[T] {
3     def show(x: T): String
4 }
5
6 // Wrapper for convenient calling.
7 def show[T](x: T)(implicit ev: Show[T]): String = ev.show(x)
8
9 // Singleton object, a value for the bundle of functions.
10 implicit object intShow extends Show[Int] {
11     def show(x: Int): String = x.toString
12 }
13
14 def showAll[T](xs: List[T])(implicit ev: Show[T]): String =
15     xs.map(show(_)).join(", ")
```

750 As we discussed earlier (2.1), working with type parameters is similar to working with ordinary  
751 ones. Also, type inference can be viewed as the process of recovering omitted type applications.  
752 Note that the implicit inference mechanism does approximately the same thing. Therefore, when  
753 trying to make functions with implicits first-class functions, difficulties similar to those of first-class  
754 polymorphism arise (2.1.2).

755 In dependently typed languages, implicit parameters<sup>41</sup> are especially needed because, for example,  
756 types are exactly the same parameters of a function as all others. Therefore, type inference is

---

<sup>40</sup>The design of implicit parameters in Scala 3 has changed (youtube) Scala Implicits Revisited, Martin Odersky.

<sup>41</sup><https://agda.readthedocs.io/en/v2.7.0.1/language/implicit-arguments.html>




Figure 11: Instance coherence — the diagram commutes.

757 essentially the inference of implicit function arguments. Moreover, dependent functions, along with  
 758 arguments, often accept proofs of some properties of those arguments, which one also wants to  
 759 infer from context automatically where possible. Such an inference mechanism can be reused for  
 760 emulating type classes<sup>42</sup> Devriese and Piessens [2011]. It also works in the reverse direction — one  
 761 can emulate dependent typing mechanisms using type classes McBride [2002].

762 As we will see later, implicit parameters on their own are also needed as a static approximation  
 763 of dynamic free variables for implementing effect systems (see 8). However, having type classes as a  
 764 separate language feature is still useful, despite them seemingly being the same implicit parameters  
 765 (albeit with a recursive inhabitation mechanism). One can support an important property when all  
 766 restrictions are respected, i.e. in the absence of **orphan instances**<sup>43</sup>. **Instance coherence** — for  
 767 a given type, all instances of a given type class obtained by different means are indistinguishable  
 768 (Fig. 11). Accordingly, the origin of any particular instance does not matter. In other words, one  
 769 need not think about this, which removes a significant amount of cognitive load and simplifies  
 770 refactoring<sup>44</sup>. In contrast, other approaches require careful attention to the calling context, because  
 771 an unexpected implementation may come from it.

### 772 3.1.6 Rules and specialization

773 GHC allows specifying optimizing rewrite rules for the compiler directly in code using a special  
 774 pragma<sup>45</sup> Jones et al. [2001]. For example:

```

1 {-# RULES
2   "map/map" forall f g xs. map f (map g xs) = map (f . g) xs
3   "map/append" forall f xs ys. map f (xs ++ ys) = map f xs ++ map f ys
4   #-}
```

<sup>42</sup><https://agda.readthedocs.io/en/v2.7.0.1/language/instance-arguments.html>

<sup>43</sup><https://stackoverflow.com/questions/3079537/orphaned-instances-in-haskell>

<sup>44</sup>Edward Kmett - Type Classes vs. the World.

<sup>45</sup>[https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/rewrite\\_rules.html](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/rewrite_rules.html)

775 The first law is nothing other than the functor law. Ideally, we formulate laws at the design  
776 stage Maguire [b], verify their correctness using property-based testing<sup>46</sup>, and then use them for  
777 optimizations.

778 One can rewrite a polymorphic version of a function to a specialized one when the types match.  
779 For this, one needs to implement the specialized version (semantic equivalence is entirely the  
780 programmer's responsibility) and set up the corresponding rewrite rule:

```
1 genericLookup :: Ord a => Table a b -> a -> b
2 intLookup    ::          Table Int b -> Int -> b

3 {-# RULES "genericLookup/Int" genericLookup = intLookup #-}
```

781 The main effect of this optimization is the guaranteed conversion of dynamic type class function  
782 calls into static ones (because the type is known, and therefore so is the corresponding dictionary).

### 783 3.1.7 Digression: defunctionalization

784 **Defunctionalization** — a technique for eliminating higher-order functions from a program<sup>47</sup> Xia.  
785 It was first proposed in Reynolds [1972, 1998].

786 The idea is to replace every creation of a lambda function with a call to a constructor of some  
787 algebraic data type. And every call site of a function is replaced with a call to a special first-order  
788 function `apply` that interprets this algebraic type.

789 Consider the example with the higher-order function `map` and two call sites that create lambda  
790 functions:

```
1 map :: (Int -> Int) -> [Int] -> [Int]
2 map f = \case [] -> []; x:xs -> f x : map f xs

3 example1 xs = map (\x -> x + 1) xs
4 example2 y xs = map (\x -> x * y) xs
```

791 For each lambda literal we create a constructor that stores the closure. Application is replaced with  
792 a call to `apply`. Thus we obtained a first-order version of `map`.

```
1 data Fun = F1 | F2 Int
2 apply :: Fun -> Int -> Int
3 apply df x = case df of F1 -> x + 1; F2 y -> x * y

4 map :: Fun -> [Int] -> [Int]
5 map df = \case [] -> []; x:xs -> apply df x : map df xs

6 example1 xs = map F1 xs
7 example2 y xs = map (F2 y) xs
```

<sup>46</sup>(youtube) John Hughes - Keynote: How to specify it!

<sup>47</sup><https://en.wikipedia.org/wiki/Defunctionalization>

### 793 3.1.8 Emulating higher-kinded polymorphism

794 Not all languages have higher-rank polymorphism, but sometimes it is useful. Its most common  
795 application is emulating type classes of arrow kinds like `Monad`.

796 Note that a type constructor of kind `Type -> Type` is a function on types that takes one type  
797 and returns another. Let us apply defunctionalization to avoid the need to parameterize one type  
798 constructor by another Xia, Yallop and White [2014].

799 We associate to the type constructor `List` a type “symbol” `ListSym` (for illustration we use  
800 Kotlin):

```
1 class ListSym
```

801 We create a type corresponding to the application of a symbol to a type, storing the original value  
802 with erased type:

```
1 class Apply<Sym, T>(val value: Any)
```

803 We establish an isomorphism between the original type obtained by type application of the constructor  
804 and the new symbol application:<sup>48</sup>

```
1 fun <T> List<T>.to(): Apply<ListSym, T> = Apply(this)  
2 fun <T> Apply<ListSym, T>.from(): List<T> = this.value as List<T>
```

805 Now we can declare a monad interface and provide an implementation for lists using a singleton  
806 object:

```
1 interface Monad<M> {  
2     fun <T> pure(x: T): Apply<M, T>  
3     infix fun <T, R> Apply<M, T>.bind(k: (T) -> Apply<M, R>): Apply<M, R>  
4 }
```

```
5 object ListMonad : Monad<ListSym> {  
6     override fun <T> pure(x: T): Apply<ListSym, T> = listOf(x).to()  
7     override fun <T, R> Apply<ListSym, T>.  
8         bind(k: (T) -> Apply<ListSym, R>): Apply<ListSym, R> =  
9         this.from().flatMap { k(it).from() }.to()  
10 }
```

807 And finally we can write functions over arbitrary monads:<sup>49</sup>

```
1 fun <M> Monad<M>.go(x: Apply<M, Int>): Apply<M, Int> =  
2     x bind { it -> pure(it + 1) } bind { it -> pure(it + 2) }  
  
3 fun test(xs: List<Int>): List<Int> = ListMonad.go(xs.to()).from()
```

---

<sup>48</sup>To the left of the dot in the declaration, an additional function argument is specified with call syntax matching a method call on an object. From the body of the function, it can be referenced using `this`.

<sup>49</sup>Arguments to the left of the dot can automatically jump into subsequent calls. In fact, they are implicit parameters.

808 It is worth noting that the resulting code is rather unwieldy. Most likely, using this technique  
809 does not pay for itself and one should choose a different programming style.

## 810 3.2 Families

811 The idea of ad-hoc polymorphism is to obtain different code depending on the type. Families  
812 started as an extension of this idea into the domain of data and types. Thus, associated type  
813 synonyms provide different types for different indices (type parameters) Chakravarty et al. [2005a].  
814 Associated `data` provide different representations for different indices Chakravarty et al. [2005b].  
815 Eventually these ideas were generalized to open families Schrijvers et al. [2008], and then closed  
816 families were introduced Eisenberg et al. [2014].

817 One can think of families<sup>5051</sup> as type constructors that define a set of types. A specific type from  
818 the set can be selected by passing a type parameter called an **index**. Compare this with ordinary  
819 polymorphic type constructors, which behave the same regardless of type parameters.

820 A large number of interesting usage examples can be found, for example, in Kiselyov et al.  
821 [2010].

### 822 3.2.1 Data families

823 Data families allow choosing an algebraic type declaration depending on the type index. For  
824 example, for a more efficient implementation of data structures. This resembles template specialization  
825 in C++.

```
1 data family XList elem
2 data instance XList () = IntList Int
3 data instance XList Bool = BoolList ByteArray
```

826 The only way to work with a data family is to place the implementation in a type class so that  
827 it accompanies the value from the point when the type index is concrete.

```
1 class XListOp elem where
2   xelem :: elem -> XList elem -> Bool

3 instance XListOp () where
4   xelem () (IntList size) = size > 0

5 instance XListOp Bool where
6   xelem key (ByteArray bs) = ...

7 xelemAll :: XListOp elem => XList elem -> [elem] -> Bool
8 xelemAll xs = all (`xelem` xs)
```

---

<sup>50</sup>[https://ghc.gitlab.haskell.org/ghc/doc/users\\_guide/exts/type\\_families.html](https://ghc.gitlab.haskell.org/ghc/doc/users_guide/exts/type_families.html)

<sup>51</sup><https://serokell.io/blog/type-families-haskell>

828 Unfortunately, unlike template specialization, one cannot provide a default definition like `data instance XL`

### 829 3.2.2 Synonym families

830 Synonym families (type families) are effectively functions on types. They can be open, closed,  
831 and associated. Open families are open in the sense that instances can be written separately from  
832 the declaration, similar to type classes. Closed families, on the contrary, are fully described in one  
833 place, cases are ordered top to bottom, which somewhat relaxes restrictions on patterns.

```
1 type family Plus (n :: Nat) (m :: Nat) :: Nat where
2   Plus Zero m = m
3   Plus (Suc n) m = Suc (Plus n m)
```

834 To see what something at the type level evaluates to, one can use the following command in  
835 `ghci`:

```
1 ghci> :k! Plus (Suc Zero) (Suc (Suc Zero))
2 Plus (Suc Zero) (Suc (Suc Zero)) :: Nat
3 = Suc (Suc (Suc Zero))
```

836 Associated families work similarly, only they are declared within a certain type class, serving as  
837 a functional alternative to `FunctionalDependencies` Jones [2000] (which look more relational). In  
838 other words, they allow associating another type with the type for which an instance is written. For  
839 example, `collections` and their element type:

```
1 class Container c where
2   type Elem c
3   elements :: c -> [Elem c]
4 instance Container [a] where
5   type Elem [a] = a
6   elements = id
7 instance Container ByteString where
8   type Elem ByteString = Word8
9   elements = ByteString.unpack
```

840 In modern languages, associated families often appear as associated types<sup>52</sup>. Swift relies heavily  
841 on associated types, not supporting generics in protocols (interfaces) at all<sup>53</sup>. Meanwhile, Scala  
842 attempts to track, in the presence of existential types (interfaces), which specific value a given  
843 associated type came from, using path-dependent types Amin et al. [2014].

---

<sup>52</sup>Associated types are essentially existential types; they are connected to associated families by the logical process of Skolemization.

<sup>53</sup>(youtube) 2017 LLVM Developers' Meeting: "Implementing Swift Generics"

### 844 3.2.3 Injective families

845 Type families differ from type constructors roughly in the same way that functions differ  
846 from data constructors. Constructors are passive and do not reduce (`Maybe Int`), while functions  
847 compute to some result (e.g. `F Int ~ Bool`). In particular, like functions, families are not necessarily  
848 injective.

849 For type constructors, knowing that the types they construct are equivalent, one can infer that  
850 the type arguments are equivalent too. For example:

```
1 Maybe a ~ Maybe b ⇒ a ~ b
```

851 Obviously, for type families this property does not hold by default:

```
1 type family NonInjective a where
2   NonInjective Int = Double
3   NonInjective Char = Double
```

852 Haskell provides explicit syntax for declaring injective type families, resembling functional dependencies  
853 in type classes (TypeFamilyDependencies) Stolarek et al. [2015]. Of course, the compiler will verify  
854 that the implementation is injective. The syntax requires binding the result to a name via equality  
855 and specifying, analogous to FunctionalDependencies, that the result determines some of the type  
856 indices of the family:

```
1 type family InjectiveB a b = r | r -> b
2   ...
```

### 857 3.2.4 First-class families

858 Besides injectivity, type families also do not necessarily possess the property of **generativity**,  
859 which is critically important for type inference — the same result is not necessarily obtained from  
860 the same family:

```
1 f a ~ g a ⇒ f ~ g
```

861 Together, injectivity and generativity constitute **matchability**. When Haskell works with a type  
862 of arrow kind, it assumes that the type is matchable. Accordingly, families cannot be passed as  
863 parameters, and all their occurrences must be fully applied to all arguments (fully saturated). Or  
864 one needs to explicitly indicate that the family returns a constructor (in this specific case the  
865 semantics depend on moving arguments to the right of `::`):

```
1 type family ToCtor (s :: Symbol) :: Type -> Type where
2   ToCtor "maybe" = Maybe
3   ToCtor "identity" = Identity
```

866 One way to work around this restriction is defunctionalization of families Xia, Eisenberg and  
867 Stolarek [2014]. As we discussed earlier 3.1.7, instead of a first-class function, a symbol designating  
868 it is introduced, along with an interpretation function that can perform the action corresponding  
869 to that symbol. In this case, the interpretation function will be an open family `Apply` [Maguire, a,  
870 chapter 10].

871 There is ongoing work in Haskell<sup>54</sup> to remove the saturation restriction Kiss et al. [2019]. For  
872 this, one needs to distinguish matchable and unmatchable type functions. This is proposed to be  
873 done by additional indexing of arrow kinds:  $\rightarrow \equiv \rightarrow^M$  and  $\rightarrow \equiv \rightarrow^U$ . And, of course, these indices can  
874 be polymorphic.

```
1 data Matchability = Matchable | Unmatchable
2 hMap
3   :: forall (m :: Matchability) (c :: Type -> Constraint)
4   . forall (f :: Type ->m Type) (as :: [Type])
5   . All as c => (forall a. c a => a -> f a) -> HList as -> HList (Map f as)
```

### 875 3.3 The Constraint kind

876 There had long been proposals to add support for constraint synonyms, constraint families, etc.  
877 to GHC Orchard and Schrijvers [2010]. Eventually, a mechanism for unifying types and constraints  
878 was proposed<sup>55</sup> and implemented<sup>56</sup>. Thus, everything that worked for types now also works for  
879 constraints.

880 In GHC with ConstraintKinds, a special kind `Constraint` was added:

- 881 • A type class constructs a constraint: `Monad :: (Type -> Type) -> Constraint;`
- 882 • Equality is a constraint: `(a ~ b) :: Constraint;`
- 883 • An empty constraint tuple is a constraint: `() :: Constraint;`
- 884 • A constraint tuple is a constraint: `(Eq a, a ~ b) :: Constraint.`

885 Now, for example, we can reify a dictionary as a language-level object:<sup>57</sup>

```
1 data Dict (c :: Constraint) where
2   Dict :: c => Dict c
```

886 Recall the heterogeneous list considered earlier 2.1.5:

```
1 data HList (tys :: [Type]) where
2   HNil :: HList '[]
3   HCons :: ty -> HList tys -> HList (ty : tys)
```

---

<sup>54</sup>GHC proposal: Unsaturated Type Families.

<sup>55</sup><https://gitlab.haskell.org/ghc/ghc/-/wikis/kind-fact>

<sup>56</sup><http://blog.omega-prime.co.uk/2011/09/10/constraint-kinds-for-ghc/>

<sup>57</sup><https://hackage.haskell.org/package/constraints-0.14.2/docs/src/Data.Constraint.html#Dict>

887 This data structure is the first-class analogue of variadic generics in C++ or Swift<sup>58</sup> (the whole  
888 point of variadics is to avoid working with `HList`-like structures directly). For example, we can write  
889 `map` for such a structure if all types satisfy a certain constraint. For this, we first implement a family  
890 that generates a constraint tuple for each type in the list:

```
1 type family All (c :: k -> Constraint) (tys :: [k]) :: Constraint where
2   All c '[] = ()
3   All c (ty : tys) = (c ty, All c tys)

4 -- All Show [Int, Double] ~ (Show Int, (Show Double, ()))
```

891 Now we can implement `map`:

```
1 hmap :: forall c res tys . All c tys
2     => (forall ty . c ty => ty -> res) -> HList tys -> [res]
3 hmap f = \case
4   HNil -> []
5   HCons x xs -> f x : hmap @c f xs
```

```
6 ghci> hmap @Show show (HCons (1 :: Int) $ HCons 'a' HNil)
```

892 More exercises of this kind with heterogeneous constructions can be found in de Vries and Löh  
893 [2014].

894 Constraints can also be parametrically polymorphic<sup>59</sup> Bottu et al. [2017]:

```
1 data Rose f x = Rose x (f (Rose f x))

2 instance (Eq a, forall b. Eq b => Eq (f b)) => Eq (Rose f a) where
3   Rose x1 rs1 == Rose x2 rs2 = x1 == x2 && rs1 == rs2
```

## 895 3.4 Uses of ad-hoc polymorphism

896 Often languages that have something resembling type classes try to express as many other  
897 language features and useful techniques through them as possible. It turns out that this is a  
898 surprisingly powerful mechanism. In this section we will consider some selected examples.

### 899 3.4.1 Serialization

900 A classic example of using type classes is serialization. The problem is that deserialization is  
901 performed when the object does not yet exist (we do not consider the deficient Java approach of  
902 filling an object with default values followed by mutation). Therefore, it is not possible to write an  
903 OOP interface `Serializable`.

904 The standard serialization library in Kotlin<sup>60</sup> provides the entity `KSerializer`, which is an

---

<sup>58</sup><https://github.com/swiftlang/swift-evolution/blob/main/proposals/0398-variadic-types.md>

<sup>59</sup>[https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/quantified\\_constraints.html](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/quantified_constraints.html)

<sup>60</sup><https://github.com/Kotlin/kotlinx.serialization>

905 interface for a separate serializer object for our type (for efficiency, CPS in the form of encoder and  
906 decoder event streams is used here; we will return to this approach later 4.4.1):

```
1 interface KSerializer<T> {  
2     fun serialize(encoder: Encoder, value: T)  
3     fun deserialize(decoder: Decoder): T  
4 }
```

907 Obviously, to construct a serializer for a polymorphic type, serializers for the type parameters  
908 are needed. Which already clearly resembles type classes.

```
1 class PairSerializer(  
2     keySerializer: KSerializer<K>,  
3     valueSerializer: KSerializer<V>,  
4 ): KSerializer<Pair<K, V>> { ... }
```

909 However, Kotlin does not have type classes, so serializers must be created manually. A set of  
910 technologies helps with this: `inline fun, reified` generics, reflection, and a compiler plugin from  
911 the library, which however is not type-safe.

### 912 3.4.2 Existential types

913 The existential quantifier in types is an encapsulation mechanism [Pierce, 2002, chapter 24].  
914 We can replace occurrences of a subtype  $\sigma$  in a type  $\tau$  with a variable  $\alpha$  and obtain  $\exists\alpha. \tau'$ , where  
915  $[\alpha \mapsto \sigma] \tau' \equiv \tau$ . Thus, we have hidden part of the type from the user.

916 For example, one can use existential types to bring different types to a common form and put  
917 their values in a single list. In Haskell, an existential type is expressed through a `data` declaration  
918 in which the data constructor is polymorphic but the type is monomorphic<sup>61</sup>. Thus, the trivial  
919 existential type  $\exists\alpha. \alpha$  can be obtained in Haskell as the type `Any`, defined as follows:

```
1 data Any where  
2     Any :: forall a . a -> Any -- logically equivalent to (exists a . a) -> Any  
  
3 list :: [Any]  
4 list = [Any 42, Any "Hello", Any (Just Nothing)]
```

920 At the deconstruction site of `Any`, a value of some unknown type will be available. Obviously,  
921 nothing can be done with such a value. However, in addition to the value, one can place in the  
922 constructor a witness that this unknown type satisfies some type class<sup>62</sup>.

```
1 data Has (c :: Type -> Constraint) where  
2     Has :: c a => a -> Has c
```

<sup>61</sup>Similar to the type erasure idiom in C++.

<sup>62</sup>Similar to Rust fat pointers (a pair of a pointer to data and a pointer to a virtual function table) and OOP objects.

923 A value of type `Has` witnesses that there exists some inhabited type `a` that belongs to a certain  
924 type class. For example, consider `Show`:

```
1 showAll :: [Has Show] -> String
2 showAll = List.intercalate ", " . map \(Has x) -> show x
```

925 In general, to eliminate such a data type, higher-rank types are needed [Pierce, 2002, chapter  
926 24.3]:

```
1 foldHas :: Has c -> (forall a . c a => a -> b) -> b
2 foldHas (Has x) k = k x
```

927 More details can be found in [Maguire, a, chapter 7] and [Pierce, 2002, chapter 24].

### 928 3.4.3 Name resolution

929 The process of name resolution in programming languages determines which program entity a  
930 given use of a name is bound to. Name resolution considers imports, the hierarchy of namespaces  
931 and scopes, expression types. . . As a rule, this is a complex process, inseparable from type inference.

932 However, in GHC the name resolution stage is rather simple and runs before type inference. Its  
933 trivial nature is hinted at by its name — Renamer — it simply rewrites names in the program to  
934 fully-qualified names, relying on imports.

935 On one hand, simplicity is good. On the other, strict separation from type inference imposes  
936 an unpleasant restriction: types cannot participate in name resolution. This problem is most acute  
937 with record field labels. One has to name all fields in a module differently to avoid clashes.

938 To make name resolution depend on types, Haskell once again resorts to type classes. Namely,  
939 a type class `IsLabel` is defined that depends on a symbol and the expected type:

```
1 class IsLabel (s :: Symbol) a where
2   fromLabel :: a
```

940 There is syntactic sugar for calling `fromLabel` (`OverloadedLabels`):

```
1 #name ≡ fromLabel @"name"
```

941 Now resolution of the name `name` will take the type into account<sup>63</sup>:

```
1 data Pet = Pet { name :: String }
2 instance IsLabel "name" (Pet -> String) where
3   fromLabel Pet{ name } = name

4 data Person = Person { name :: String, pets :: [Pet] }
5 instance IsLabel "name" (Person -> String) where
6   fromLabel Person{ name } = name

7 ghci> #name pet
```

---

<sup>63</sup>To get rid of the redefinition error, you need to enable `NoFieldSelectors`.

### 942 3.4.4 Non-syntactic type equalities, System FC

943 Modern Haskell is a syntactically rich language that, despite its diversity of constructs, is  
944 translated into a small typed internal language. This is the language *System F<sub>C</sub>* Sulzmann et al.  
945 [2007b], which extends *System F* (2.1) with non-syntactic type equalities. It turns out that this  
946 is sufficient to support such Haskell features as generalized algebraic data types, associated type  
947 families, functional dependencies, and so on.

948 Specifically, a built-in constraint  $\sim$  is introduced, witnessing the equivalence of two types<sup>64</sup>. For  
949 example, the type of function `id` can be written in this unusual way:

```
1 f :: forall a b . a ~ b => a -> b
2 f = id
```

950 In fact, this is a function of four parameters: two type parameters, a coercion, and an argument. A  
951 coercion is a zero-sized value, automatically inferred by the compiler, that serves as a witness that  
952 the two corresponding types are equivalent.

953 For example, the GADT from 2.1.4 is desugared as follows:

```
1 data Expr ty where
2   Const :: Int -> Expr Int
3   IsZero :: Expr Int -> Expr Bool
4   If :: forall ty . Expr Bool -> Expr ty -> Expr ty -> Expr ty
5   -- translates to
6 data Expr ty where
7   Const :: forall ty . ty ~ Int => Expr ty
8   IsZero :: forall ty . ty ~ Bool => Expr Int -> Expr ty
9   If :: forall ty . Expr Bool -> Expr ty -> Expr ty -> Expr ty
```

954 And after pattern matching on the constructors, the equality constraint enters the branch and  
955 allows the type inference system to perform the necessary rewrites.

956 More on type inference in the presence of local assumptions can be found in the classical paper  
957 *OutsideIn(X)* Vytiniotis et al. [2011].

958 Obviously, Haskell can inhabit an equality constraint following reflexivity, symmetry, and transitivity.  
959 Also, the compiler can generate new axioms (the user cannot directly write their own axioms).  
960 For example, for a type family the compiler generates equality axioms between the applied family  
961 constructor and the resulting types:

```
1 type family Plus (n :: Nat) (m :: Nat) :: Nat where
2   Plus Zero m = m
3   Plus (Suc n) m = Suc (Plus n m)
4   -- expands to
5 axiom Plus Zero m ~ m
6 axiom Plus (Suc n) m ~ Suc (Plus n m)
```

---

<sup>64</sup>[https://ghc.gitlab.haskell.org/ghc/doc/users\\_guide/exts/equality\\_constraints.html](https://ghc.gitlab.haskell.org/ghc/doc/users_guide/exts/equality_constraints.html)

### 962 3.4.5 Coercions and roles

963 Haskell supports **generative type abstractions** in the form of `newtype` declarations. This  
964 technique allows defining domain-specific types that are indistinguishable from the wrapped types  
965 at runtime but allow distinguishing them at compile time. For instance, we can introduce wrappers  
966 for numbers that in the problem domain represent identifiers of different entities. Now the type  
967 system will prevent confusing them.

```
1 newtype ModuleId = ModuleId Int64
2 newtype CourseId = CourseId Int64
```

968 There is an extremely underappreciated programming practice<sup>65</sup> where our program has a clear  
969 boundary at which data from the external world is parsed. Beyond it, raw unstructured data is  
970 enriched with structure and takes on meaning within the problem domain. Or we reject this data  
971 as incorrect. In the remaining part of the program, we work with typed data whose properties have  
972 already been established and guaranteed. For example, we can be sure that the number `ModuleId`  
973 is strictly greater than zero.

974 However, if we have a collection of wrapped data and we want to work with it as a collection of  
975 raw data, we would have to transform the collection, even though this transformation does nothing.  
976 The Haskell optimizer cannot eliminate it because it works with a typed intermediate representation  
977 and cannot remove a conversion that changes the type.

```
1 newtype Csv = Csv { unCsv :: String }
2 concatC :: [Csv] -> Csv
3 concatC = Csv . concat . map unCsv
```

978 Therefore, Haskell has a mechanism of safe coercions between types with the same runtime  
979 representation. This is implemented using the magical type class `Coercible`. The compiler implements  
980 it automatically (see Fig. 12).

```
1 class Coercible from to where
2   coerce :: from -> to
```

981 Now we can eliminate the unnecessary list transformation:

```
1 concatC :: [Csv] -> Csv
2 concatC = coerce concat
```

982 The safety of coercions is ensured by the **role system**. Each type parameter has a special  
983 property — a role.

984 The phantom role is assigned to phantom type parameters. They can be freely coerced (there  
985 are no prerequisites for `Coercible` instances):

---

<sup>65</sup><https://lexi-lambda.github.io/blog/2019/11/05/parse-don-t-validate/>

The most important rules that GHC uses to solve `Coercible` constraints are as follows (the full rules are given in Figure 5):

(1) The *unwrapping rule*:

- ▶ For every **newtype** `NT = MkNT t`, we have `Coercible t NT` if and only if the constructor `MkNT` is in scope.

(2) The *lifting rule*:

- ▶ For every type constructor `TC r p n`, where
  - `r` stands for `TC`'s parameters at a representational role,
  - `p` for those at a phantom role and
  - `n` for those at a nominal role,if `Coercible r1 r2`, then `Coercible (TC r1 p1 n) (TC r2 p2 n)`.

(3) `Coercible` is an equivalence relation:

- ▶ The *reflexivity rule*: `Coercible a a`.
- ▶ The *symmetry rule*: If `Coercible a b` then `Coercible b a`.
- ▶ The *transitivity rule*: If `Coercible a b` and `Coercible b c` then `Coercible a c`.

Figure 12: Principles of constructing `Coercible` instances Breitner et al. [2014].

```
1 data Phantom h = Phantom
2 data NestedPhantom b = MkNP [Phantom b] | SomethingElse
3 instance Coercible (Phantom a) (Phantom b)
4 instance Coercible (NestedPhantom a) (NestedPhantom b)
```

986 A type parameter has the representational role if the type constructor can be coerced only when  
987 the arguments can be coerced:

```
1 data Maybe a = Nothing | Just a
2 instance Coercible a b => Coercible (Maybe a) (Maybe b)
```

988 A type parameter has the nominal role if the type constructor can be coerced only when the  
989 arguments are equal. This is required when the type argument indexes a family or a constraint.

```
1 type family F a
2 data Applied a = Applied (F a)
3 instance (a ~ b) => Coercible (Applied a) (Applied b)
4 data ShowDict a where
5   ShowDict :: Show a => a -> ShowDict a
6 instance (a ~ b) => Coercible (ShowDict a) (ShowDict b)
```

990 Sometimes the compiler infers the wrong role for a type parameter. For example, if the invariants  
991 of a structure depend on a specific implementation of some type class for the type argument, which  
992 is not visible in the type declaration itself. In such cases, roles can be specified explicitly:

```
1 type role Map nominal representational
2 data Map k v = ...
```

993 More details can be found in Breitner et al. [2014] and [Maguire, a, chapter 8].

### 994 3.4.6 Type reflection

995 Reflection is a language mechanism for obtaining information about types at runtime (at the  
996 term level). This sounds familiar, and indeed Haskell implements this mechanism through type  
997 classes Peyton Jones et al. [2016].

998 The library provides the magical type class `Typeable`, which is implemented by the compiler for  
999 each concrete type via `deriving`. To obtain information about a type, a `Typeable` instance for that  
1000 type must be in scope. The type structure is represented by the sum type `TypeRep`, which provides  
1001 the ability for additional type checking using generalized algebraic data types and type tags.

```
1 class Typeable a where
2   typeRep# :: TypeRep a
```

1002 For example, one can obtain the name of a type constructor as follows:

```
1 typeName :: forall a. Typeable a => String
2 typeName = tyConName $ typeRepTyCon $ typeRep $ Proxy @a
3 ghci> typeName @Int
```

1003 **Exercise 18** *Declare a type class that allows printing a list of types.*

1004 Using the type representation structure and existential types, one can emulate dynamic typing  
1005 in Haskell. Namely: any type can be converted to `Dynamic`, and then safely converted back.

```
1 data Dynamic where
2   Dynamic :: Typeable a => a -> Dynamic
3 fromDynamic :: Typeable a => Dynamic -> Maybe a
```

1006 This can be useful, for example, for defining a heterogeneous key-value store:

```
1 data Store = Map Key Dynamic
2 data Ref ty = Ref Key
3 get :: Typeable ty => Store -> Ref ty -> Maybe ty
```

### 1007 3.4.7 Data reflection

1008 As we discussed earlier, the coherence property guarantees that each type in Haskell corresponds  
1009 to exactly one instance of a given type class. And the only way to declare an instance in Haskell  
1010 is with a top-level declaration, meaning it cannot depend on any local data. However, Haskell has  
1011 the library `Data.Reflection`<sup>66</sup>, which allows creating local instances for fresh, magically generated<sup>67</sup>  
1012 types.

1013 It uses the idea of “lifting values into types” that we discussed earlier (see 3.1.4), but in a  
1014 somewhat more general form. Instead of defining type classes of the form `Known_`, a single type  
1015 class is introduced, indexed by the type of terms `terms` that are lowered from types:

```
1 class Reifies ty terms | ty -> terms where
2   reflect :: Proxy ty -> terms
```

1016 Also, using the following function, the library allows generating a fresh type and a `Reifies`  
1017 instance that, given this fresh type, returns the given value of type `a` (passed as the first argument).  
1018 Since it is passed into a higher-rank function, the fresh type cannot escape the scope 2.1.2:

```
1 reify :: a -> (forall fresh . Reifies fresh a => Proxy fresh -> res) -> res
```

1019 To use a non-standard instance of a type class for some type `a`, one needs to declare a new  
1020 type (for example, using `newtype`), containing the given one, and write the needed instance for it  
1021 (see, for example, `Down`). We do not want to declare a new declaration for each case, so we create  
1022 a wrapper similar to `Data.Tagged`, which allows adding a phantom type tag to the value type. By  
1023 varying the tag, we can obtain as many types wrapping the given one as we need.

```
1 newtype Wrapped tag a = Wrapped { unwrap :: a }
```

1024 We declare the wrapper type `Wrapped tag a` as a member of the needed type class. The  
1025 implementation code is obtained via `reflect` by the tag type in the form of an actual dictionary.

```
1 data ReifiedOrd a = ReifiedOrd { compare :: a -> a -> Ordering }
2 instance Reifies tag (ReifiedOrd a) => Ord (Wrapped tag a) where
3   compare = coerce $ compare $ reflect $ Proxy @tag
```

1026 Finally, we can call the `sort` function, locally replacing the ordering with the reverse one:

```
1 sort :: Ord a => [a] -> [a]
2 sortReverse :: forall a . Ord a => [a] -> [a]
3 sortReverse xs =
4   let dict = ReifiedOrd { compare = flip compare } in
5   reify dict \(Proxy :: Proxy fresh) ->
6     coerce $ sort @(Wrapped fresh a) $ coerce xs
```

<sup>66</sup><https://www.tweag.io/blog/2017-12-21-reflection-tutorial/>

<sup>67</sup><https://www.schoolofhaskell.com/user/thoughtpolice/using-reflection>

### 1027 3.4.8 Open structures

1028 In dynamic languages, one can create objects on the fly, sequentially adding content to them  
1029 without introducing a declaration beforehand. In Haskell, this is also possible using pairs for products  
1030 and `Either` for sums. For example, one can add a new field by creating a new pair: `(oldObj, newField)`.

1031 However, this implementation is suboptimal both in terms of efficiency (more efficient implementations  
1032 can be found in [Maguire, a, chapter 11]) and in terms of usability. Namely, the order of fields matters  
1033 and there is no subtyping relation on types in Haskell (for example, one cannot pass a value with  
1034 fewer fields or variants). But one can notice that constraints lack these drawbacks. Therefore, the  
1035 type of a data structure can be organized, for example, as follows:

```
1 (Int, Double) replaced by (Member Int d, Member Double d) => Prod d
```

### 1036 3.4.9 Exceptions and the open hierarchy

1037 An important aspect of error handling is that many errors are handled uniformly. Thus, errors  
1038 should form a hierarchy similar to the one obtained in OOP languages through inheritance, so that  
1039 one can react to a group of errors with a single handler. This gives rise to the task of modeling  
1040 such a hierarchy in Haskell.

1041 Moreover, statically typed errors are an active area of research; we will discuss this in the context  
1042 of effect systems (see 8). Classical exceptions, however, are dynamically typed. This variant is  
1043 especially well-suited for programmer errors, which ideally should not be handled in the program  
1044 except for resource cleanup.

1045 Exception support is present in the Haskell runtime system as a simple and familiar way to handle  
1046 exceptional situations: programmer errors, incomplete pattern match execution, asynchronous system  
1047 signals Marlow et al. [2001]. . . Exceptions are dynamically typed and form a hierarchy Marlow [2006].  
1048 While pure code can throw an exception, since  $\perp$  by Haskell's semantics inhabits every type, catching  
1049 an exception is only possible<sup>68</sup> in `IO` Jones [2001], using special language primitives.

1050 To make a type an exception, one needs to declare an `Exception` instance for it:

```
1 class (Typeable a, Show a) => Exception a where  
2   toException :: a -> SomeException  
3   toException = SomeException  
  
4   fromException :: SomeException -> Maybe a  
5   fromException (SomeException e) = cast e
```

1051 Where `SomeException` is an existential wrapper similar to `Dynamic` (see 3.4.6), into which the  
1052 concrete exception type is wrapped. By catching `SomeException`, one can catch any exception  
1053 (`cast` will always succeed).

```
1 data SomeException where
```

---

<sup>68</sup>(stackoverflow) Why can Haskell exceptions only be caught inside the IO monad?

```

2   SomeException :: Exception a => a -> SomeException

3   instance Exception SomeException

```

1054 The Haskell runtime system provides intrinsics for throwing and catching exceptions; we wrap  
1055 them to support any `Exception` type:

```

1   throw :: Exception e => e -> a
2   throw = primThrow (toException e)

3   catch :: Exception e => IO a -> (e -> IO a) -> IO a
4   catch io handler = io `primCatch` \e -> case fromException e of
5     Nothing -> throw e
6     Just e'  -> handler e'

7   ghci> throw "error" `catch` \(e :: String) -> putStrLn e

```

1056 In the simplest case, a custom exception type can be implemented in two lines. To catch it, one  
1057 must either catch this type itself or `SomeException`, because for both of them `fromException` on  
1058 an object of the form `SomeException MyError` will return `Just`<sup>69</sup>.

```

1   data MyError = MyError deriving (Show, Typeable)
2   instance Exception MyError

```

1059 Let us add the exception `ArithException` and one more, more general exception type between it  
1060 and `SomeException` — `SomeArithException` (thus, there will be three ways to catch `ArithException`).  
1061 For this, we make `SomeArithException` an existential wrapper, and automatically wrap each  
1062 exception of type `ArithException` in it. In `fromException`, at each level of nesting, we try to  
1063 obtain the wrapping constructor by a recursive call.

```

1   data SomeArithException where
2     SomeArithException :: Exception a => a -> SomeArithException

3   -- SomeException - base (default implementation)
4   instance Exception ArithException

5   data DivisionByZero = DivisionByZero deriving Show
6   instance Exception DivisionByZero where
7     toException = toException . SomeArithException
8     fromException e = do
9       SomeArithException e' <- fromException e
10      cast e'

```

---

<sup>69</sup>`instance Exception SomeException where fromException = Just`

1064 Thus, when throwing `DivisionByZero`, an object of the following form will be constructed:

```
1 SomeException (SomeArithException DivisionByZero)
```

1065 The implementation of `fromException` for a concrete type can verify the presence of the corresponding  
1066 constructor in the resulting exception object.

### 1067 3.4.10 Lightweight partial stack traces

1068 An amusing exploitation of type classes in Haskell is lightweight partial stack traces<sup>70</sup>. In general,  
1069 collecting traces requires runtime support, which in Haskell's case is further complicated by the fact  
1070 that the computation model, graph reduction, does not contain real traces and they have to be  
1071 emulated. Here we obtain traces without runtime support.

1072 The standard library defines the constraint `GHC.Stack.HasCallStack`, allowing one to obtain  
1073 information about the call site of a function. This information is effectively placed by the compiler  
1074 during instance inference. If information from the level above is available at the call site, the compiler  
1075 propagates it further. Thus, information is available only to a certain depth of the call stack.

```
1 myHead :: HasCallStack => [a] -> a
2 myHead []      = error "empty"
3 myHead (x:xs) = x
```

```
4 bad :: Int
5 bad = myHead []
```

```
6 ghci> bad
7 *** Exception: empty
8 CallStack (from HasCallStack):
9   error, called at Bad.hs:8:15 in main:Bad
10  myHead, called at Bad.hs:12:7 in main:Bad
11  -- no information about bad call site here
```

1076 `HasCallStack` is simply an implicit (see 3.1.2) that the compiler knows about:

```
1 type HasCallStack = (?callStack :: CallStack)
```

### 1077 3.4.11 Customizable type errors

1078 When developing libraries that are complex from the type perspective, it is desirable to provide  
1079 users with more informative type errors than the defaults. For this, GHC has a mechanism in  
1080 `GHC.TypeLits` that allows constructing a special type whose information will appear in the error  
1081 message. For example, this type can be returned from a synonym family with an incorrect set of  
1082 arguments. Or one can use the constraint trick (see 3.1.3) and place such a type as a premise in an

---

<sup>70</sup>[https://downloads.haskell.org/ghc/latest/docs/users\\_guide/exts/callstack.html](https://downloads.haskell.org/ghc/latest/docs/users_guide/exts/callstack.html)

1083 instance. If the instance matched and the compiler started processing the constraints on the left,  
1084 something went wrong [Maguire, a, chapter 12].

```
1 instance (TypeError
2   ( Text "Attempting to show a function of type "
3     :<>: Text "" :<>: ShowType (a -> b) :<>: Text ""
4     :$$: Text "Did you forget to apply an argument?"
5   )) => Show (a -> b) where
6   show = undefined -- the implementation does not matter, execution will not reach here
```

## 1085 4 Data Types

1086 This chapter collects some general knowledge about types. We will also obtain various equivalent  
1087 representations of recursive data types (in other words, collections). Many concepts are special cases  
1088 of this variety.

1089 Sections 4.1, 4.2 mostly follow [Maguire, a, chapter 1].

### 1090 4.1 Variance

1091 In this section we will consider the topic from the programming perspective [Maguire, a, chapter  
1092 3], without doing justice to category theory. The gap can be filled with the help of a wonderful  
1093 article written in the style of a play Hinze et al. [2012].


1094 A **covariant functor** is a pair of some type constructor  $F$  and an operation on functions  $\text{fmap}$   
1095  $:: (a \rightarrow b) \rightarrow (F\ a \rightarrow F\ b)$ . Plus laws stating that  $\text{fmap}$  respects  $\text{id}$  and composition.

```
1 class Functor f where
2   fmap :: (a -> b) -> (f a -> f b)
```



1096 A **contravariant functor** is a pair of a type constructor and an operation on functions that  
1097 reverses the arrow. Plus corresponding laws.

```
1 class Contravariant f where
2   contramap :: (a -> b) -> (f b -> f a)
```



1098 A type constructor can be declared a covariant or contravariant functor (or neither) with respect  
1099 to a given type parameter depending on the form of the declarations of the corresponding data

1100 constructors. Namely, on the sign of the positions in which that type parameter appears in the  
1101 type.

1102 Let us develop an intuitive understanding of position signs. Type **A** appears in a positive position  
1103 in **B** if its value can be extracted from **B**. Conversely, type **A** appears in a negative position if its value  
1104 must, on the contrary, be provided. Let us consider position signs of types in basic type constructors:

Type	sign of position <b>A</b>	sign of position <b>B</b>
<b>Either A B</b>	+	+
<b>(A, B)</b>	+	+
<b>A -&gt; B</b>	-	+

1106 Indeed, from a sum and a product one can extract components via pattern matching, and from  
1107 an arrow one can obtain the right type by applying it to an argument. At the same time, a value of  
1108 the type on the left of the arrow must be provided.

1109 Plus and minus obey intuitive algebraic rules when considering more complex types. Let us  
1110 consider the example `f :: ((A, B) -> C) -> (D, E)`.

- 1111 • Plus times plus gives plus. Indeed, one simply needs to apply two eliminations instead of one  
1112 to obtain the desired type. In our example, to obtain **D**, one first applies the function, then  
1113 destructs the pair.
- 1114 • Plus times minus (and vice versa) gives minus. Indeed, **C** must be provided by us: `f (\ab -> provideC)`.
- 1115 • Minus times minus gives plus. The pair **(A, B)** is provided to us: `f (\ab -> ...)`.

1116 **Exercise 19** *Verify that plus times minus gives minus.*

1117 Returning to functors, if a type parameter appears in the declaration only in positive positions,  
1118 the type constructor can be declared a covariant functor with respect to that parameter. If only  
1119 in negative positions — a contravariant functor. If in both, then it cannot be declared any functor.  
1120 Accordingly, we call type parameters covariant, contravariant, and invariant.

1121 **Exercise 20** *Declare `instance Contravariant F` for `data F a = L (a -> ()) | R Int`.*

1122 Thus, a covariant functor can be understood as a computation whose result can be post-  
1123 processed, and a contravariant functor as a computation whose arguments can be pre-processed.

1124 A type with two positive parameters can be declared a **bifunctor**:

```
1 class Bifunctor f where
2   bimap :: (a -> c) -> (b -> d) -> f a b -> f c d
```

1125 A type with two parameters, one positive and one negative, is a **profunctor**:

```
1 class Profunctor p where
2   dimap :: (c -> a) -> (b -> d) -> p a b -> p c d
```

1126 Profunctors are certain generalizations of the function arrow. For example, if we have an SQL  
1127 query that takes data and returns a result, it can be declared a profunctor with the semantics of  
1128 adding pre-processing of input data and post-processing of output data:

```
1  dimap serialize deserialize (query :: Sql Text Text) :: Sql Age [User]
```

1129 The concept of variance is also frequently encountered in object-oriented languages (and in  
1130 subtyping theory in general) to denote the possibility of extending the subtyping relation to polymorphic  
1131 types.

1132 Indeed, the **subtyping relation**  $B <: A$  says that a value of type  $B$  can safely be used in a  
1133 position where a value of type  $A$  is expected. In other words, there exists a function `upcast :: B`  
1134 `-> A`. If the type constructor  $F$   $a$  is covariant with respect to parameter  $a$ , then from `upcast` one  
1135 can obtain `upcast' :: F B -> F A`. That is, the subtyping relation also automatically includes  $F$   
1136  $B <: F A$ . The contravariant case is analogous.

1137 **Exercise 21** *Verify in your favorite language with subtyping and variance support that minus times*  
1138 *minus gives plus.*

## 1139 4.2 Isomorphism

1140 Suppose we need to design a function or a data model. We start with a type declaration —  
1141 how should we choose it and from which alternatives? First, let us understand when two types are  
1142 interchangeable; for this, we consider the notion of isomorphism.

1143 Two types  $A$  and  $B$  are called **isomorphic** (denoted  $A \cong B$ ) if and only if there exists a pair of  
1144 functions `to :: A -> B` and `from :: B -> A` such that<sup>71</sup>

```
1  to . from = id  
2  from . to = id
```

1145 In other words, one can establish a bijection between the inhabitants of such types. It is easy to  
1146 see that from a semantic point of view it does not matter which of the isomorphic types is used —  
1147 they can be substituted for each other by adding calls to the conversion functions. Two such types  
1148 contain the same “amount of information.” For example, the types `Bool` and `Maybe ()` are in this  
1149 sense completely interchangeable. Let us demonstrate this by exhibiting a pair of mutually inverse  
1150 functions<sup>72</sup>:

```
1  to :: Bool -> Maybe ()  
2  to b = if b then Just () else Nothing  
  
3  from :: Maybe () -> Bool  
4  from m = case m of Nothing -> False; Just () -> True
```

---

<sup>71</sup>Equality of terms can mean different things, for example  $\alpha\beta\gamma$ -equivalence. We will use **extensional equality** for functions — two functions are equal when their results are equal on all inputs. <https://ncatlab.org/nlab/show/function+extensionality>

<sup>72</sup>One must not forget to show that the functions are mutually inverse, but this is done trivially by enumeration of inputs (possibly with induction) and reduction.

1151 Despite semantic interchangeability, to encode information about whether the user passed a  
 1152 particular flag to the program, we would most likely use the type `Bool` due to non-functional  
 1153 considerations of code readability. Analogously, one can consider efficiency considerations.

1154 A categorical perspective on the matter can be found in Hinze and James [2010]. We, however,  
 1155 adhere to the set-theoretic interpretation of types.

## 1156 4.2.1 Cardinality: sums, products, exponentials

1157 Types can be thought of as syntax for writing sets, and the terms inhabiting them as syntactic  
 1158 representations of elements of those sets. Thus the term `(True, False)` is a representation of an  
 1159 element of the set of pairs, written in type syntax as `(Bool, Bool)` (instead of the mathematical  
 1160  $\mathbb{B} \times \mathbb{B}$ ). Or the term `\x -> x + 1` is a representation of the function adding one from the set of  
 1161 functions over integers, written as `Integer -> Integer` (instead of the mathematical  $\mathbb{Z} \rightarrow \mathbb{Z}$ ).

1162 Note that two types are isomorphic if the corresponding sets have the same number of elements.  
 1163 Moreover, there are  $n!$  such isomorphisms in the case of finite sets. Let us learn to determine the  
 1164 number of such elements. Using  $|\cdot|$  we will denote the **cardinality** of a type — the number of  
 1165 elements in the corresponding set.

Type and its declaration	cardinality
<code>data Void</code>	0
<code>data Unit = Unit</code> <sup>73</sup>	1
<code>data Bool = False   True</code>	2

1167 The idea of algebraic data types is that complex types can be built from simple ones using the  
 1168 operation  $+$  (“or”) and the operation  $\times$  (“and”)<sup>74</sup>:

Type	cardinality
<code>data Either a b = Left a   Right b</code>	$ a  +  b $
<code>data Pair a b = Pair a b</code>	$ a  \times  b $

1170 Let us count the number of inhabitants of various types (you can verify the conclusion by  
 1171 enumerating all terms manually):

- 1172 •  $|Either\ Unit\ (Either\ Bool\ Bool)| = |Unit| + (|Bool| + |Bool|) = 5$ .
- 1173 •  $|Pair\ (Either\ Bool\ Unit)\ (Pair\ Unit\ Void)| = 0$  — the type `Void` is uninhabited, as  
 1174 is a tuple containing it.
- 1175 • If `data Example = FirstAlternative Bool | AnotherOne Unit Bool Bool`, then  
 1176  $|Example| = |Bool| + |Unit| * |Bool| * |Bool| = 2 + 1 * 2 * 2 = 6$ .

1177 The function arrow is called the exponential type. Indeed, combinatorially the number of inhabitants  
 1178 of `A -> B` is computed as

$$|A \rightarrow B| = |B|^{|A|}$$

1179 So how should one design types? There are several considerations:

<sup>73</sup>`Unit` is written in Haskell using special syntax `()`, meaning something like an empty tuple.

<sup>74</sup><https://stanford-cs242.github.io/f18/lectures/02-2-algebraic-data-types.html>

- 1180 • The type should have no fewer elements than the domain, so that all necessary objects are  
1181 representable.
- 1182 • The type should have as few elements as possible that are not in the domain, so that the  
1183 error space is minimal.
- 1184 • Then among isomorphic types, the optimal one is chosen based on non-functional requirements.

1185 Per the second rule, before manipulating a domain object one should first convert the data  
1186 describing it into the most structured representation possible — the one that provides the strongest  
1187 guarantees<sup>75</sup>.

## 1188 4.2.2 Algebraic representation of a type

1189 As we saw above, to show the existence of an isomorphism between two types one can either  
1190 exhibit a pair of mutually inverse functions, or show that the cardinalities of the two types coincide.  
1191 In this section we will learn to associate an algebraic expression with a type, reflecting its structure  
1192 and cardinality. Thus, we will be able to obtain equivalent expressions via syntactic transformations  
1193 of formulas, from which we will recover types that are guaranteed to be isomorphic to the given  
1194 one<sup>76</sup>.

1195 The algebraic representation is based on computing cardinalities of types. In effect, we forget  
1196 information inessential for isomorphism about the names of data constructors and type constructors,  
1197 i.e. we transition to structural typing.

Type	algebraic formula
<code>data Void</code>	0
<code>data Unit = Unit</code>	1
<code>data Bool = False   True</code>	1 + 1 (denote as 2)
<code>data Maybe a = Nothing   Just a</code>	1 + a
<code>data Either a b = Left a   Right b</code>	a + b
<code>data Pair a b = Pair a b</code>	a × b
<code>a -&gt; b</code>	b <sup>a</sup>

1199 **Exercise 22** Write the following type in algebraic form:

```
1 data T a b = Undefined | Defined a (a -> b)
```

1200 As the equivalence relation, we use isomorphism of the corresponding types. Under this interpretation,  
1201 the classical properties of algebraic operations are preserved (Fig. 13). Indeed, for example:

```
1 -- (cb)a ≅ ca×b
2 to :: (a -> b -> c) -> (a, b) -> c
3 to = uncurry
4 from :: ((a, b) -> c) -> a -> b -> c
5 from = curry
```

$$\begin{array}{c}
\text{sums} \\
X + 0 \cong X \quad X + Y \cong Y + X \\
(X + Y) + Z \cong X + (Y + Z) \\
\text{products} \\
X \times 1 \cong X \quad 0 \times X \cong 0 \quad X \times Y \cong Y \times X \\
(X \times Y) \times Z \cong X \times (Y \times Z) \\
\text{product over sum} \\
(X + Y) \times Z \cong (X \times Z) + (Y \times Z) \\
\text{exponentials} \\
1^X \cong 1 \quad X^0 \cong 1 \quad X^1 \cong X \\
Z^{X+Y} \cong Z^X \times Z^Y \quad (Z^Y)^X \cong Z^{X \times Y} \\
\text{exponential over product} \\
(Y \times Z)^X \cong Y^X \times Z^X
\end{array}$$

Figure 13: Laws of school algebra for nostalgia's sake Hinze and James [2010].

1202 **Exercise 23** Show that  $(a + b) + c \cong a + (b + c)$ .

1203 **Exercise 24** Show that  $c^{a+b} \cong c^a \times c^b$ .

1204 An interesting observation is that functions can be used as data structures, in accordance with  
1205 the isomorphism  $c^{a+b} \cong c^a \times c^b$ . Indeed, in that case the function argument serves as an index (its  
1206 cardinality must coincide with the size of the collection).

```

1  -- a × a ≅ a2
2  get :: (a, a) -> (Bool -> a)
3  get (x, y) idx = if idx then x else y
4  tabulate :: (Bool -> a) -> (a, a)
5  tabulate f = (f True, f False)

```

1207 The **canonical representation of a type** is a sum of products of types:

$$\sum_i \prod_j t_{ij}$$

1208 The canonical representation is a kind of normal form in which algebraic types can be written (any  
1209 algebraic type can be reduced to it using the rules). It is easy to recognize in it the form of **data**  
1210 declarations in Haskell.

1211 Remarkably, even the derivative of an algebraic type has a perfectly understandable semantics.  
1212 It is the context of a zipper, a data structure similar to an iterator that allows navigating through  
1213 a data structure of arbitrary shape Huet [1997], McBride [2001], Abbott et al. [2003].

<sup>75</sup><https://lexi-lambda.github.io/blog/2019/11/05/parse-don-t-validate/>

<sup>76</sup><https://codewords.recurse.com/issues/three/algebra-and-calculus-of-algebraic-data-types>

## 1214 4.3 Recursive types

1215 Recursive types, similarly to recursive terms, can include themselves in their definitions. In other  
1216 words, a recursive type is isomorphic to some other type in which it itself appears as a subtype.

```
1 fac n = if n <= 1 then 1 else n * fac (n - 1)
2 data Nat = Zero | Suc Nat
```

1217 We will mostly look at recursive types from a practical perspective. However, their formal  
1218 type-theoretic description, category-theoretic and set-theoretic interpretations are of independent  
1219 interest [Pierce, 2002, part 4].

### 1220 4.3.1 Just a list

1221 Consider the classical functional list. A list is either a collection of zero elements, or one, or  
1222 two... Algebraically this can be written as follows:

$$L = 1 + a + a^2 + a^3 + \dots$$

1223 In effect we obtained a type with an infinite expression. Let us work with it as a formal series.  
1224 Factor out  $a$ :

$$L = 1 + a \times (1 + a + a^2 + \dots)$$

1225 Notice that the expression in parentheses represents a list, yielding the following recursive equation<sup>77</sup>:

$$L = 1 + a \times L$$

1226 It is easy to see that this is in fact the familiar list definition from Haskell:

```
1 data List a = Nil | Cons a (List a)
```

1227 Let us obtain a finite non-recursive representation of the type  $L$ . In effect, we need to obtain a  
1228 type isomorphic to a type that includes the original:

$$L \cong 1 + a \times L$$

1229 Let us extend the type language with abstraction (polymorphism) and solve the resulting recursive  
1230 equation in the style of  $\lambda$ -calculus, using a fixpoint combinator:

$$L = \text{FIX } \lambda r. 1 + a \times r$$

1231 Let us encode this in Haskell. As the recursion combinator we take `data FixList a`:

---

<sup>77</sup>Alternatively, one can obtain the same result by noticing that we are dealing with a Taylor series <https://codewords.recurse.com/issues/three/algebra-and-calculus-of-algebraic-data-types>.

```

1 type Shape a r = Either () (a, r) -- λar.1 + a × r
2 data FixList a = In (Shape a (FixList a))
3 -- FixList a ≅ Shape a (FixList a) ≅ Shape a (Shape a (FixList a)) ≅ ...

4 example :: FixList Int -- [1, 2, 3]
5 example = In (Right (1, In (Right (2, In (Right (3, In (Left ())))))))

```


1232 Thus, we have separated the list definition into two parts: one is responsible for the shape of  
 1233 the type, and the other for recursion<sup>78</sup>. The shape can be reused in the definition of a fold:

```

1 foldr :: (Either () (a, r) -> r) -> FixList a -> r
2 foldr phi (In shape) = case shape of
3   Left () -> phi (Left ())
4   Right (x, xs) -> phi (Right (x, foldr phi xs))
5 -- compare with the classical definition
6 foldr :: r -> (a -> r -> r) -> [a] -> r
7 foldr ini f list = case list of
8   [] -> ini
9   x:xs -> f x (foldr ini f xs)

```

1234 Indeed, a fold in the general sense allows replacing each data constructor in the tree with some  
 1235 function. The result is a computation that has access to the entire contents of the data structure  
 1236 and returns an aggregated result:



1237

1238 That is, the number and types of arguments to the fold depend on the “shape” of the type being  
 1239 folded. In the case of a list, a nullary function is needed in place of Nil and a binary one in place  
 1240 of Cons. We represent these two functions as one using the shape type. Compare:

```

1 foldr          0 (      \e r -> e + r)
2 foldr (\case Left () -> 0; Right (e, r) -> e + r)

```

### 1241 4.3.2 Fixpoint of a functor

1242 Let us abstract FixList over the shape type:

<sup>78</sup>The technique of abstracting over the “recursive self-reference” is called **open recursion**. This gives the user more control. For example, at the term level, the recursion combinator can count the number of recursive calls and abort the computation when some limit is exceeded.

```

1 newtype Fix :: (Type -> Type) -> Type
2 newtype Fix f = In { out :: f (Fix f) }

3 data ListF a r = Nil | Cons a r
4 type List a = Fix (ListF a)

```

1243 **Exercise 25** *What will the types of `In` and `out` be?*

1244 One can show that  $[a] \cong \text{List } a$ :

```

1 to :: [a] -> List a
2 to = \case
3   [] -> In Nil
4   x:xs -> In $ Cons x (to xs)

5 from :: List a -> [a]
6 from (In shape) = case shape of
7   Nil -> []
8   Cons x xs -> x : from xs

```

1245 The shape type can be made a functor in its last parameter. This will allow us to subsequently  
1246 replace occurrences of subtrees with something useful.

```

1 instance Functor (ListF a) where
2   fmap :: (rec -> other) -> ListF a rec -> ListF a other
3   fmap f = \case
4     Nil -> Nil
5     Cons x xs -> Cons x (f xs)

```

1247 **Exercise 26** *Express the following tree as a fixpoint of a functor. Declare a functor instance for*  
1248 *the shape type.*

```

1 data Tree a = Leaf a | Node a (Tree a) (Tree a)

```

### 1249 4.3.3 Recursion schemes

1250 Just as structured imperative programming, compared to undisciplined use of `goto`, helps in  
1251 reasoning about programs, so **recursion schemes** allow one to algebraically describe properties  
1252 of recursive functions, as opposed to “unstructured” recursion<sup>79</sup> Meijer et al. [1991], Meijer and  
1253 Hutton [1995].

1254 The overarching idea was to formulate required properties and compute the needed programs  
1255 similarly to how mathematicians find solutions to differential equations. However, this idea did not  
1256 gain the necessary development and adoption (though mastering the algebraic approach is generally

<sup>79</sup><https://reasonablypolymorphic.com/blog/recursion-schemes/index.html>




Figure 14: A catamorphism reuses occurrences of recursive subtrees to store the fold results of the corresponding subtrees.

1257 useful Maguire [b]). Nevertheless, this knowledge, on the one hand, gives a deeper understanding  
 1258 of recursion, and on the other hand, will serve us when examining and solving the main problem of  
 1259 this course — the expression problem.

1260 The universal fold is called a **catamorphism**. A catamorphism first recursively folds the subtrees,  
 1261 reaching them via `fmap`, and leaves fold results of type `a` in place of the former occurrences of  
 1262 subtrees. The result is a value of type `f a`, where `f` is some shape functor. Then a function of type  
 1263 `f a -> a` is applied, which determines how to fold one layer of the recursive structure when the  
 1264 subtrees have already been folded.

```
1 cata :: Functor f => (f a -> a) -> Fix f -> a
2 cata phi = phi . fmap (cata phi) . out
```

1265 In other words, `cata` replaces all `In` constructors with the function `phi`, which contains information  
 1266 about how to react to different constructor variants in the tree:

```
1 example          = In (Cons 1 (In (Cons 2 (In (Cons 3 (In Nil))))))
2 cata phi example =β phi (Cons 1 (phi (Cons 2 (phi (Cons 3 (phi Nil)))))
```

1267 For example, summing a list looks as follows (Fig. 14):

```
1 sum :: List Int -> Int
2 sum = cata \case
3   Nil -> 0
4   Cons x result -> x + result
```

1268 The function `f a -> a` is called an **f-algebra**. Indeed, if we take the signature of an algebra as  
 1269 the functor `f` and the carrier as `a`, then the `f`-algebra defines some interpretation of the signature:

```
1 data MonoidSig carrier = Mempty | Mappend carrier carrier
2 interpretSig :: MonoidSig Int -> Int
3 interpretSig = \case Mempty -> 0; Mappend l r -> l + r
```

1270 **Exercise 27** Which algebra can be used to copy a data structure?

1271 **Exercise 28** Which algebra can be used to print a list as a string?

1272 In contrast to the universal fold, one can construct an **anamorphism** — a universal unfold  
1273 (analogous to `unfold` for a list). Here the f-coalgebra (the arrow in the opposite direction) shows  
1274 how to obtain one layer of a data structure from some seed value, where in place of recursive  
1275 references there will be seeds from which subtrees will subsequently grow. The anamorphism first  
1276 unfolds one layer, and then recursively unfolds all the subtrees:

```
1 ana :: Functor f => (s -> f s) -> s -> Fix f
2 ana psi = In . fmap (ana psi) . psi
3 -- compare with the classical definition of list unfold
4 unfoldr :: (s -> Maybe (a, s)) -> s -> [a]
```

1277 **Exercise 29** Implement an anamorphism that builds a list from 0 to a given *n*.

1278 One also introduces the **hylomorphism**, which allows describing an arbitrary recursive computation.  
1279 A hylomorphism is defined as a composition of an anamorphism and a catamorphism. First, the  
1280 anamorphism builds an explicit tree representing the call tree of some recursive procedure, then the  
1281 catamorphism folds it into the result.

```
1 hylo :: Functor f => (a -> f a) -> (f b -> b) -> a -> b
2 hylo psi phi = cata phi . ana psi
```

1282 For example, computing factorial can be implemented as follows:

```
1 fac n = hylo
2   (\n -> if n > 0 then Cons n (n - 1) else Nil)
3   (\case Nil -> 1; Cons n acc -> n * acc)
```

1283 One can introduce many more different recursion schemes<sup>80</sup> and describe their properties.  
1284 However, we will stop here for now.

1285 An interesting fact: a catamorphism represents the set-theoretic principle of induction, and an  
1286 anamorphism represents coinduction, if one takes the function as the set embedding [Pierce, 2002,  
1287 chapter 21].

## 1288 4.4 Everything through folds

1289 It turns out that using a catamorphism one can obtain an isomorphism between data structures  
1290 and their folds:  $\text{Fix } f \cong \text{forall } a . (f a \rightarrow a) \rightarrow a$ .

```
1 to :: Functor f => Fix f -> (forall a . (f a -> a) -> a)
2 to = flip cata

3 from :: (forall a . (f a -> a) -> a) -> Fix f
4 from g = g In
```

---

<sup>80</sup>[https://wiki.haskell.org/Zygomorphic\\_prepromorphisms](https://wiki.haskell.org/Zygomorphic_prepromorphisms)

1291 For example, the following two lists are equivalent (all `In` constructors are replaced with the  
 1292 given algebra):

```

1 data ListF elem rec = Nil | Cons elem rec

2 xs1 :: Fix (ListF Int)
3 xs1 = In (Cons 1 (In (Cons 2 (In (Cons 3 (In Nil))))))

4 xs2 :: (ListF Int a -> a) -> a
5 xs2 = \alg -> alg (Cons 1 (alg (Cons 2 (alg (Cons 3 (alg Nil))))))

6 ghci> xs2 @Int \case Nil -> 0; Cons x acc -> x + acc
7 6
  
```

1293 Now let us get rid of the shape functor. It is a non-recursive type that can be presented in  
 1294 canonical form (see 4.2.2):

$$f\ a \cong \sum_i \prod_j (t_{ij}\ a)$$

1295 Then the algebra can be written as follows:

$$f\ a \rightarrow a \cong a^{\sum_i \prod_j (t_{ij}\ a)} \cong \prod_i a^{\prod_j (t_{ij}\ a)}$$

1296 Products remain, which can be eliminated using currying, yielding **Church encoded** data structures.  
 1297 For example, for a list we have:

$$\begin{aligned}
 (\text{ListF elem } a \rightarrow a) \rightarrow a &\cong a^{a^{1 + \text{elem} \times a}} \cong a^a \times a^{\text{elem} \times a} \cong (a^a)^{\text{elem}} \\
 &\cong a \rightarrow (\text{elem} \rightarrow a \rightarrow a) \rightarrow a
 \end{aligned}$$

1298 We have obtained nothing other than a Church list<sup>8182</sup>. A data structure without a single  
 1299 constructor!<sup>83</sup> Let us rewrite the list we already know once more:

```

1 xs3 :: a -> (Int -> a -> a) -> a
2 xs3 = \ini f -> f 1 (f 2 (f 3 ini))
  
```

1300 **Exercise 30** Which familiar standard list function on a `data` list returns a Church list?

1301 Let us try to intuitively understand what all this means. Notice that a Church list accepts  
 1302 functions corresponding to pattern-matching branches or arguments of the folding function. Instead

<sup>81</sup>[https://en.wikipedia.org/wiki/Church\\_encoding](https://en.wikipedia.org/wiki/Church_encoding)

<sup>82</sup><https://okmij.org/ftp/tagless-final/course/Boehm-Berarducci.html>

<sup>83</sup>In fact this is not at all surprising, if one recalls that first-class functions are represented as closures that contain data. We have obtained the same singly linked list, but using closures.

1303 of constructors, the corresponding functions are called directly. That is, instead of creating a  
1304 data structure and delivering it to the site of deconstruction (pattern matching), we *deliver the*  
1305 *deconstruction site to the construction site*, and it turns out that there is nothing to construct  
1306 after all.

1307 We have already worked with the reverse process, defunctionalization, when first-class functions  
1308 were turned into constructor calls, and their bodies into pattern-matching branches 3.1.7. One  
1309 can notice that here we are dealing with the reverse process, **refunctionalization**, when instead of  
1310 constructor calls the corresponding interpreting code is called directly Danvy and Millikin [2009].

1311 From a technical perspective, we have understood that conditional branching (if, pattern matching,  
1312 switch constructs) and virtual calls (closure calls, method calls on interfaces) are interchangeable.

1313 It is also worth noting that the isomorphism  $\text{Fix } f \simeq \text{forall } d . (f \ d \ \rightarrow \ d) \ \rightarrow \ d$  is a  
1314 generalization of the isomorphism  $a \sim \text{forall } r . (a \ \rightarrow \ r) \ \rightarrow \ r$  (with  $f = \text{Const } a$ ), which  
1315 is the foundation of the CPS transformation that we consider later in the chapter 6.

#### 1316 4.4.1 Deforestation & list fusion

1317 In functional programming we build new functions by composing existing, simpler functions. This  
1318 approach allows reusing implemented functionality, reducing code complexity and the probability  
1319 of errors. However, it can lead to unnecessary overhead from allocation and deconstruction of  
1320 intermediate data structures.

1321 For example, compare the following two implementations. The first is preferable from a code  
1322 quality standpoint, but it creates an intermediate list during execution:

```
1  all p xs = and (map p xs)
2  -- or the fused version
3  all p [] = True
4  all p (x:xs) = p x && all p xs
```

1323 Obviously, this is the job of an optimizing compiler to turn good code with abstractions into  
1324 fast code. The optimization that eliminates (intermediate) data structures (trees) from programs  
1325 is called **deforestation**. As a result, two functions are said to “fuse.”

1326 The term and the first deforestation algorithm were proposed by Philip Wadler Wadler [1988];  
1327 it is based on several simple rewriting rules aimed at obtaining the situation `case K args of ...`,  
1328 and aggressive inlining (see examples of operation in Fig. 15). However, this algorithm can lead  
1329 to exponential code blowup and may not terminate in the presence of recursive calls<sup>84</sup>. For the  
1330 algorithm to terminate, programs must be written in a certain strict form called treeless.

1331 Let us focus on lists and consider more practical solutions<sup>85</sup>. The first obvious solution would be  
1332 to add a special rewrite rule for each pair of list functions (see 3.1.6), deforesting using algebraic  
1333 properties of list transformations. However, there would be exponentially many such rules.

---

<sup>84</sup>Deforestation is a special case of supercompilation Романенко, which in turn is a generalization of a large number of compiler optimizations.

<sup>85</sup><https://markkarpov.com/tutorial/ghc-optimization-and-fusion.html>

*append (append xs ys) zs*

transforms to

*h<sub>0</sub> xs ys zs*

where

*h<sub>0</sub> xs ys zs = case xs of*  
                  *Nil              : h<sub>1</sub> ys zs*  
                  *Cons x xs      : Cons x (h<sub>0</sub> xs ys zs)*  
*h<sub>1</sub> ys zs      = case ys of*  
                  *Nil              : zs*  
                  *Cons y ys      : Cons y (h<sub>1</sub> ys zs)*

---

*flip (flip zt)*

transforms to

*h<sub>0</sub> zt*

where

*h<sub>0</sub> zt = case zt of*  
          *Leaf z              : Leaf z*  
          *Branch xt yt      : Branch (h<sub>0</sub> xt) (h<sub>0</sub> yt)*

Figure 15: Examples of the deforestation algorithm in action Wadler [1988].

```

1  std::variant<Msg1, Msg2>           1  template<class Impl>
2      deserialize(bytes bs) {       2  auto deserialize(bytes bs) {
3      if (...) {                    3      if (...) {
4          return std::variant{Msg1(...)}; 4          return Impl::processMsg1(...);
5      else {                          5      else {
6          return std::variant{Msg2(...)}; 6          return Impl::processMsg2(...);
7      }                                7      }
8  }                                    8  }

```

Figure 16: Manual deforestation in C++.

```

1  {-# RULES
2  "map/map" forall f g xs. map f (map g xs) = map (f . g) xs
3  #-}

```

1334 The modern deforestation technique in Haskell — fold/build list fusion Gill et al. [1993] —  
1335 instead of using many algebraic rules, defines a universal way of constructing and deconstructing a  
1336 list. We deconstruct a list using `foldr`. We construct using the function `build`:

```

1  build :: (forall b . (a -> b -> b) -> b -> b) -> [a]
2  build g = g (:) []

```

1337 For example, the list `[1, 2, 3]` and the function `map` are now written as follows<sup>86</sup>:

```

1  list123 :: [Int]
2  list123 = build \s z -> s 1 (s 2 (s 3 z))

3  map :: (a -> b) -> [a] -> [b]
4  map f xs = build \s z -> foldr (\x acc -> s (f x) acc) z xs

```

1338 One can notice that the input to `build` is a Church list, hence the unsurprising law:

```

1  foldr f ini (build g) ≡ g f ini

```

1339 Thus, we eliminated list construction by replacing constructor calls with calls to folding functions.  
1340 Deforestation can also be performed manually (see 16).

#### 1341 4.4.2 Visitor pattern

1342 Consider some tree and its fold:

```

1  data Tree a = Leaf | Node a [Tree a]
2  foldTree :: Tree a -> r -> (a -> [r] -> r) -> r

```

<sup>86</sup>In the Haskell standard library, list functions are written normally, but alongside them are RULES rules (see 3.1.6) that substitute their implementation with fold/build versions.

1343 Let us rewrite and rename:

```
1 data Visitor a r = Visitor { onLeaf :: r, onNode :: a -> [r] -> r }
2 visitTree :: Tree a -> Visitor a r -> r
```

1344 To make this look even more like OOP style, the tree itself should be given by a fold (as an  
1345 interface with a visit function), and different nodes by concrete implementations (descendant  
1346 objects):

```
1 data Tree a = Tree { visit :: forall r . Visitor a r -> r }

2 leaf :: Tree a
3 leaf = Tree { visit = \Visitor{onLeaf} -> onLeaf }

4 node :: a -> [Tree a] -> Tree a
5 node x ts = Tree { visit = \v@Visitor{onNode} -> onNode x (map (`visit` v) ts) }
```

1347 For clarity, let us show this code in Kotlin as well:

```
1 interface Visitor<a, r> {
2     fun onLeaf(): r
3     fun onNode(x: a, subtrees: List<r>): r
4 }

5 interface Tree<a> {
6     fun <r> visit(visitor: Visitor<a, r>): r
7 }
8 class Leaf : Tree<Nothing> {
9     override fun <r> visit(visitor: Visitor<Nothing, r>): r = visitor.onLeaf()
10 }
11 class Node<a>(val value: a, val subtrees: List<Tree<a>>) : Tree<a> {
12     override fun <r> visit(visitor: Visitor<a, r>): r =
13         visitor.onNode(value, subtrees.map { it -> it.visit(visitor) })
14 }
```

## 1348 4.5 Everything through unfolds

1349 Recall that there exists a universal unfold — the anamorphism, which from a generating procedure  
1350 produces an entire data structure (see 4.3.3).

```
1 ana :: Functor f => forall s . (s -> f s) -> s -> Fix f
2 ana psi = In . fmap (ana psi) . psi
```

1351 Let us bring the function `ana` to a type of the form `A -> B`, so that it is then easier to study the  
1352 arrow `B -> A`. To do this, we uncurry and move the quantifier to the left of the arrow (it changes  
1353 to the opposite):

```
1 ana :: Functor f => (exists s . (s, s -> f s)) -> Fix f
```

1354 Let us encode the existential quantifier using a new type (see 3.4.2) and rewrite the anamorphism  
1355 to work with it:

```
1 data Box f where
2   --      exists s. (s,      s -> f s)
3   Box :: forall s . s -> (s -> f s) -> Box f

4 ana' :: Functor f => Box f -> Fix f
5 ana' (Box currSeed psi) =
6   In $ (\nextSeed -> ana' (Box nextSeed psi)) <$> psi currSeed
```

1356 Now let us build an isomorphism between data structures and their trivial unfolds, which each time  
1357 return the next layer of the given structure,  $\text{Fix } f \simeq \text{Box } f$ :

```
1 to :: Fix f -> Box f
2 to x = Box x out

3 from :: Functor f => Box f -> Fix f
4 from = ana'
```

1358 Thus, we have obtained evidence that any recursive data structure can be, at least trivially,  
1359 represented as  $\text{Box } f$ . Such a representation is sometimes called **co-Church encoding** Gibbons  
1360 [2008].

1361 For example, an infinite lazy list of natural numbers can be defined as follows. Note that here  
1362 we *do not rely on Haskell's laziness*, which means we can use this technique in eager languages as  
1363 well.

```
1 nats :: Box (ListF Int)
2 nats = Box 0 \curr -> Cons curr (curr + 1)
```

1364 **Exercise 31** *Implement a lazy function*

```
1365 take :: Int -> Box (ListF Int) -> Box (ListF Int).
```

1366 In other words, infinite data structures can be represented as procedures that lazily generate  
1367 one layer of a recursive structure at a time on demand.

### 1368 4.5.1 Abstract data types

1369 Consider the case when the shape functor is a product:

$$s \rightarrow f \ s \cong \left( \prod_i (t_i \ s) \right)^s \cong \prod_i (s \rightarrow t_i \ s)$$

1370 That is, the coalgebra is equivalent to a tuple of functions. Thus, `Box` is nothing other than an  
 1371 *abstract data type (ADT)*: it contains hidden state of unknown nature and a set of operations for  
 1372 working with it Gibbons [2008]. In fact, we have already encountered such a construction earlier,  
 1373 when we discussed existential types 3.4.2 (there we used type class instances as tuples of functions).  
 1374 OOP objects can also be viewed through coalgebras. More about the relationship between  
 1375 objects and abstract data types can be found in Cook [2009].

## 1376 4.5.2 Stream fusion

1377 Earlier we considered the `foldr/build` list fusion optimization that eliminates intermediate lists  
 1378 (see 4.4.1). However, this technique does not work for many popular functions, such as `zip` (one  
 1379 needs to coroutine between two algebras) or `take` (one needs to abort the fold at a certain point).

1380 Later, a technique using the co-structure of a list — an `unfold` or `Stream` — was proposed<sup>87</sup> Coutts  
 1381 et al. [2007]:

```

1 data ListF a r = Nil | Cons a r
2 type MyStream a = Box (ListF a) -- ∃s . (s -> ListF a s, s)
3 -- ≅
4 data Step a s = Done | Yield a s
5 data Stream a where
6   Stream :: forall s . (s -> Step a s) -> s -> Stream a

7 stream :: [a] -> Stream a
8 stream xs = Stream (\case [] -> Done; x:xs -> Yield x xs) xs

9 unstream :: Stream a -> [a]
10 unstream (Stream next s) = case next s of
11   Done -> []
12   Yield a s' -> a : unstream (Stream next s')
```

1382 The idea is that now *stream processing functions are non-recursive* and easily undergo basic  
 1383 compiler deforesting transformations in the style of Wadler [1988]. Indeed, instead of a recursive  
 1384 call we remember the state for the next iteration.

```

1 mapS :: (a -> b) -> Stream a -> Stream b
2 mapS f (Stream next s) = Stream next' s
3   where
4     next' s = case next s of
5       Done -> Done
6       Yield x s' -> Yield (f x) s'
```

1385 The overall pipeline looks as follows:

<sup>87</sup><https://markkarpov.com/tutorial/ghc-optimization-and-fusion.html>

- 1386 1. We write all list functions through streams:
- ```
1387 1 map :: (a -> b) -> [a] -> [b]
1388 2 map f = unstream . mapS f . stream
```
- 1387 2. Using RULES (see 3.1.6) we define a rewrite rule for stream/unstream, dual to foldr/build:
- ```
1388 stream . unstream = id.
```
- 1389 3. Standard compiler optimizations take care of the rest.<sup>88</sup>

1390 To support a non-recursive filter function, a separate kind of step is usually added:

```
1391 1 data Step a s = Done | Skip s | Yield a s

1392 2 filterS :: (a -> Bool) -> Stream a -> Stream a
1393 3 filterS p (Stream next s) = Stream next' s
1394 4   where
1395 5     next' s = case next s of
1396 6       Done -> Done
1397 7       Skip s' -> Skip s'
1398 8       Yield a s' -> if p a then Yield a s' else Skip s'
```

1391 The Haskell standard library still uses foldr/build fusion, but there exist many industrial stream  
1392 libraries [Bragilevsky, chapter 14].

1393 More about streams can be found at Oleg Kiselyov's page<sup>89</sup>.

## 1394 4.6 Ubiquitous duality

1395 We have considered three forms of recursive data:

- 1396 • **Fix**  $f$  — ordinary recursive data types, finite in eager languages;
- 1397 •  $\forall a . (f\ a \rightarrow a) \rightarrow a$  — Church encoding, finite structures defined by folds;
- 1398 •  $\exists s . (s \rightarrow f\ s, s)$  — co-Church encoding, potentially infinite data structures, data abstractions  
1399 (working through an interface over a hidden representation).

1400 In this section we consider some practical examples in which the solution space is defined by  
1401 dual representations of collections.

### 1402 4.6.1 Push vs pull streaming

1403 Let us consider various representations of lists or, as they are sometimes called, event streams.  
1404 A more detailed comparison can be found here [Kiselyov et al., 2017, section 3].

1405 Church encoding for lists gives a **push stream**, or **internal iteration**, or **Observer pattern**<sup>90</sup>  
1406 (in a sense a special case of the Visitor pattern 4.4.2), where we register an event handler and the

<sup>88</sup>Full deforestation can be guaranteed using staging (see later 9) Kiselyov et al. [2017].

<sup>89</sup><https://okmij.org/ftp/Streams.html>

<sup>90</sup><https://learn.microsoft.com/en-us/dotnet/api/system.iobserver-1?view=net-9.0>

1407 collection calls (pushes) methods on it. Examples of push streams include the Java Stream API<sup>91</sup>,  
1408 Clojure Transducers<sup>92</sup>.

```
1  events.observe(object : Observer<Event> {  
2      override fun onComplete() { .. }  
3      override fun onNext(elem: Event) { .. }  
4  })
```

1409 co-Church encoding for lists gives a **pull stream**, or **external iteration**, or **Iterator pattern**.  
1410 To get the next element, the consuming code must explicitly request (pull) it. This representation  
1411 is often much more convenient from a usage perspective (see 4.5.2)<sup>93</sup>, but is more complex on  
1412 the implementation side, because one needs to identify and maintain the existential state in the  
1413 coalgebra. Generators are designed to automate the identification of this state, since one can notice  
1414 that it is the continuation of the computation (see later 6.4.3).

1415 Working with streams is often called **reactive programming**<sup>94</sup>.

## 1416 4.6.2 Data vs codata

1417 Recall the shape functor for lists:

```
1  data ListF a r = Nil | Cons a r
```

1418 Using the `data` keyword we define a new algebraic data type through the ways to construct  
1419 it. Then, using pattern matching we can deconstruct values of an algebraic data type and use the  
1420 information stored in it.

```
1  data List a where  
2      List :: ListF a (List a) -> List a  
  
3  fold :: (ListF a r -> r) -> List a -> r  
4  fold alg xs = case xs of  
5      List Nil -> alg Nil  
6      List (Cons y ys) -> alg (Cons y (fold alg ys))
```

1421 However, one can also define a data type through the ways to deconstruct it using the `codata`  
1422 keyword (in an imaginary language). Thus, from a co-list one can request the next layer of the  
1423 structure. Such a data type can be constructed using co-pattern matching, describing how the  
1424 deconstructing function should behave on a constructing term (`unfold coalg s`)<sup>95</sup>.

---

<sup>91</sup><https://docs.oracle.com/javase/8/docs/api/java/util/stream/Stream.html>

<sup>92</sup><https://clojure.org/reference/transducers>

<sup>93</sup><https://github.com/pulldown-cmark/pulldown-cmark/?tab=readme-ov-file#why-a-pull-parser>

<sup>94</sup>(youtube) React 2014 : Erik Meijer - What does it mean to be Reactive?

<sup>95</sup>(youtube) CS410 2017 Lecture 15 Coinduction and Coalgebras

```

1  codata CoList a where
2    force :: CoList a -> ListF a (CoList a)

3  unfold :: (s -> ListF a s) -> s -> CoList a
4  force (unfold coalg s) = case coalg s of
5    Nil -> Nil
6    Cons x s' -> Cons x (unfold coalg s')

```

1425 Essentially, `codata` represents a dictionary of functions that capture some state necessary for  
1426 generating subsequent layers of the structure. Thus, `codata` can be compared to OOP interfaces,  
1427 and co-pattern matching to an anonymous interface implementation. The two preceding examples  
1428 can be written, for example, in Kotlin as follows (the list shape is encoded as `Pair<a, r>?`, verify  
1429 for yourself that this type is isomorphic to `ListF`).

```

1  class List<a>(val layer: Pair<a, List<a>>?)

2  fun <a, r> fold(alg: (Pair<a, r>?) -> r, xs: List<a>): r =
3    if (xs.layer == null) alg(null)
4    else alg(Pair(xs.layer.first, fold(alg, xs.layer.second)))

5  interface CoList<a> { fun force(): Pair<a, CoList<a>>? }

6  fun <a, s> unfold(coalg: (s) -> Pair<a, s>?, ini: s): CoList<a> =
7    object : CoList<a> {
8      override fun force(): Pair<a, CoList<a>>? {
9        val layer = coalg(ini)
10       return if (layer == null) null else
11         Pair(layer.first, unfold(coalg, layer.second))
12     }
13 }

```

1430 The functional style of programming primarily operates on directly observable data `data`, while  
1431 OOP hides data behind interfaces `codata`. However, both constructs have their clear areas of  
1432 application and should not be neglected regardless of preferred style Downen et al. [2019]<sup>96</sup>.

1433 In languages with totality checking, inductive and coinductive definitions are distinguished. For  
1434 them, pattern matching and co-pattern matching, the induction principle (catamorphism) and the  
1435 coinduction principle (anamorphism) are used. Inductive definitions specify finite data structures;  
1436 totality checking verifies that recursive calls are made on structurally smaller subterms. Coinductive  
1437 definitions specify potentially infinite data structures; for them, productivity is checked — the  
1438 generating function always produces a new constructor after a recursive call<sup>97</sup>.

<sup>96</sup><https://reasonablypolymorphic.com/blog/review-codata/>

<sup>97</sup><https://rocq-prover.org/doc/V8.18.0/refman/language/core/coinductive.html>

## 1439 4.7 Appendix: category of algebras

1440 This is an optional section that is not necessary for understanding the rest of the course.  
1441 However, it is useful for understanding frequently used terminology.

1442 A **category** is a collection of objects and a collection of arrows. For each object  $X$  there exists  
1443 an identity arrow, and for each pair of arrows there exists a way to obtain their composition:  
1444  $f : Y \rightarrow Z, g : X \rightarrow Y \Rightarrow f \circ g : X \rightarrow Z$ .

1445 One defines a category corresponding to Haskell — *Hask*. In reality this is a poor category from  
1446 the theoretical standpoint, but it will suffice for our informal reasoning<sup>98</sup>. Objects in *Hask* are types  
1447 of the Haskell language, and morphisms are terms defining functions between the corresponding  
1448 types. The identity morphism is `id`, and composition is defined as `f . g = \x -> f (g x)`.

1449 **Exercise 32** *How does one represent constants in such a category?*

1450 A **functor** is a mapping between categories that maps objects of one category to objects of  
1451 another, and arrows of one to arrows of the other. In Haskell, type constructors define the mapping  
1452 between objects, and `fmap` defines the mapping between arrows. A functor must preserve the identity  
morphism and composition.

$$\begin{array}{ccc} A & \xrightarrow{f} & B \\ F A & \xrightarrow{\text{fmap } f} & F B \end{array}$$

1453 An **algebra** in a category  $C$  is a pair of an object of the category  $X \in \text{Obj}(C)$  and a morphism  
1454  $\phi : F X \rightarrow X$ , where  $F$  is a functor. The morphism  $F X \rightarrow X$  itself is called an **f-algebra**. Algebras  
1455 in the categorical sense can describe algebras. Thus, as the object  $X$  we take the carrier of the  
1456 algebra. As the functor  $F$  — the signature of the algebra in the form of a shape type. Then the  
1457 morphism is the interpretation of the signature.  
1458

1459 An **algebra morphism** is a morphism between carriers  $h : X \rightarrow Y$  such that the following  
diagram commutes. A **diagram commutes** if all possible paths along the arrows in it are equal.

$$\begin{array}{ccc} F X & \xrightarrow{\text{fmap } h} & F Y \\ \text{phi} \downarrow & & \downarrow \text{psi} \\ X & \xrightarrow{h} & Y \end{array}$$

<sup>98</sup><https://math.andrej.com/2016/08/06/hask-is-not-a-category/>

1460

1461 In an algebra morphism one can recognize familiar features of homomorphisms, that is, operations  
1462 between carriers that “respect” the operations of the algebraic theory’s signature.

1463 Algebras over a category  $C$  form a **category of algebras**, in which the objects are algebras and  
1464 the morphisms are algebra morphisms.

1465 An **initial object** of a category is an object from which there exists a unique arrow to every  
1466 other object. A **terminal (final) object** of a category is an object into which there exists a unique  
1467 arrow from every other object.

1468 Initial and terminal objects of a category do not necessarily exist in a single instance. But all  
1469 initial objects are isomorphic to each other, as are all terminal objects.


1470 **Exercise 33** Give the initial and terminal objects of the category *Hask*.

1471 A recursive type is a type, which means it corresponds to an object in the category *Hask*.  $X$  is  
1472 a recursive type with shape  $F$  if the following isomorphism holds:

$$X \simeq F X$$

1473 One can notice that the isomorphism witness from right to left resembles an  $f$ -algebra, and  
1474 from left to right an  $f$ -coalgebra (everything is the same, but with all arrows reversed). And indeed,  
1475 the appropriate object  $X$  must be either the initial object of the category of algebras, or the  
1476 terminal object of the category of coalgebras (with the corresponding morphisms). The first variant  
1477 corresponds to finite data structures, and the second to potentially infinite ones.

1478 The initial object of the category of algebras over *Hask* for the functor  $f$  is the following  
1479 algebra:  $(\text{Fix } f, \text{In})$ <sup>99</sup><sup>100</sup>. Indeed, for every type  $a$  and for every  $f$ -algebra  $\text{phi}$  we can construct  
a morphism  $\text{cata } \text{phi} :: \text{Fix } f \rightarrow a$  such that the following diagram commutes:




1480

1481 Let us show that  $(\text{Fix } f, \text{out})$  is the terminal object of the category of coalgebras (thanks  
1482 to Haskell’s laziness). Analogously, for every object  $a$  and  $f$ -coalgebra  $\text{psi}$  there exists a morphism  
1483  $\text{ana } \text{psi} :: a \rightarrow \text{Fix } f$ :

<sup>99</sup><https://bartoszmilewski.com/2017/02/28/f-algebras/>

<sup>100</sup><https://ncatlab.org/nlab/show/catamorphism>



## 1484 5 Interpreters

1485 The art of programming largely consists of the ability to manage complexity, and within this  
1486 course we will consider building languages and interpreters as the primary tool for this.

### 1487 5.1 Interpreters as the foundation

1488 We begin with an overview of the role of interpreters in programming.

#### 1489 5.1.1 The tower of interpreters

1490 The most basic interpreter is the processor, physically embodied in hardware. It receives a  
1491 program in some language, such as x86, reads instructions, and turns them into actions on memory.  
1492 However, it is difficult for a human to program in this language; a new language is needed that  
1493 encapsulates part of the complexity and hides unnecessary details.

1494 To obtain a new language, we build a software interpreter. A **software interpreter**  $U_M^N$  is a  
1495 program in language  $M$ <sup>101</sup> that takes as input a program in language  $N$  and its input — data from  
1496  $D$ , and returns the result of executing that program on that data:

$$U_M^N : N \times D \rightarrow D$$

1497 We will call the implementation language of the interpreter  $M$  the **meta-language**, and  $N$   
1498 the **object language**. One can intuitively think of an interpreter as follows: it is an explanation,  
1499 understandable to the meta-language, of what the constructs of the object language mean. In other  
1500 words, which meta-language instructions need to be executed to obtain the desired semantics of  
1501 the object language instructions.

1502 For example, given a program  $p_N$  and its input data  $d_{in}$ , the result of executing this program  
1503  $d_{out}$  can be obtained as

$$d_{out} = U_M^N \left( \underbrace{\langle p_N, d_{in} \rangle}_{\in N \times D} \right)$$

1504 But an interpreter is also a program. How do we run it? Take our basic interpreter  $U^{x86}$ ; it has  
1505 no implementation language since it is implemented in hardware, not software. Take an assembly  
1506 language interpreter implemented in x86 codes,  $U_{x86}^{Asm}$ , an assembly program  $p_{Asm}$ , and its input  $d_{in}$ .  
1507 Recall that a program is also data, just in a special format. Then we obtain the result of applying  
1508  $p_{Asm}$  to the data as follows:

$$d_{out} = U^{x86} \left( \left\langle \underbrace{U_{x86}^{Asm}}_{\in x86}, \underbrace{\langle \underbrace{p_{Asm}}_{\in Asm}, \underbrace{d_{in}}_{\in D} \rangle}_{\in D} \right\rangle \right)$$

---

<sup>101</sup>By a language we mean the set of programs in that language, in other words, the set of trees of a certain form.

1509 But assembly language is also not very pleasant for programming. However, in it one can already  
1510 write an interpreter for a more complex language. And so on. We obtain a *tower of interpreters*,  
1511 at the top of which is the language in which we want to solve our actual problem:

$$d_{out} = U^{x86} (\langle U_{x86}^{Asm}, \langle U_{Asm}^C, \langle U_C^{Has}, \langle p_{Has}, d_{in} \rangle \rangle \rangle \rangle)$$

1512 In practice, a language is often defined via translation (compilation) into another. However,  
1513 we will focus on embedded languages going forward, and the notion of an interpreter will be more  
1514 useful to us. Moreover, there exist universal theoretical methods for deriving a compiler from an  
1515 interpreter<sup>102</sup> (see 9), so for us they are similar.

### 1516 5.1.2 Interpreters everywhere

1517 Very well, we have arrived at the language of our heart (Haskell); why do we continue talking  
1518 about interpreters? Because for solving concrete business problems, general-purpose languages are  
1519 still too ceremonious — the programmer has to think about a large number of things irrelevant to  
1520 their domain and the problem at hand. Complexity is the programmer's main enemy, because the  
1521 human brain's resources are no match for the complexity of the reality our programs must describe.  
1522 Thus, in practice one constantly has to describe new languages most suited for solving specific  
1523 applied problems. And new languages are defined using interpreters.

1524 What does a classical recursive interpreter look like? It receives a program in the form of some  
1525 tree and recursively traverses it, computing the results of subtrees. When it visits a tree node, it  
1526 determines its type and understands what actions need to be performed. That is, the node type  
1527 dispatches, navigates, the interpreter's execution to the appropriate code. Thus, an interpreter for  
1528 a simple expression language looks as follows:

```
1  eval :: Expr -> Int
2  eval prog = case prog of
3    Const x -> x
4    Plus l r -> eval l + eval r
```

1529 One can see that this resembles, for example, a command-line utility — we parse arguments,  
1530 determine what needs to be done and how, and do it. Surprisingly, the Unix philosophy, in particular,  
1531 consists of building small languages (utilities with a text API) that solve one problem well Bentley  
1532 [1986]. It also resembles handling a request by a web server — determining the route the request  
1533 arrived at and performing the corresponding action. That is, we write interpreters in real life more  
1534 often than we think. We simply do not see that what we write is actually an interpreter for some  
1535 language. In general, folding a data structure can already be considered as interpretation Gibbons  
1536 and Wu [2014].

1537 Moreover, as we will see later 5.4, writing any function is already defining a new language.  
1538 There was a language in which one could not add a user to an application. We wrote a function  
1539 `registerUser` — a new command appeared in the language — add a user. Below we will formally

---

<sup>102</sup>[https://en.wikipedia.org/wiki/Partial\\_evaluation](https://en.wikipedia.org/wiki/Partial_evaluation)

1540 show that this approach is equivalent to adding a new node to the language's syntax tree (5.4.1).  
1541 Using functions is an example of language embedding, when instead of creating a new separate  
1542 language, we implement it as a library for an already existing language Gibbons [2013].

1543 As we will see throughout the course, many problems can be viewed as inventing a language and  
1544 implementing an interpreter. This means that if we learn to write interpreters well, we will learn to  
1545 do a lot of things at once! And our main efforts will be directed at studying the tools for building  
1546 interpreters of embedded languages.

1547 During the exposition we frequently use the Hutton's Razor technique, which involves considering  
1548 a ridiculously small language for studying complex concepts. It is argued that for studying most  
1549 questions one can construct such a language that makes everything important maximally clear.

### 1550 5.1.3 Interpreters and programming language semantics

1551 Programming language semantics<sup>103</sup> is a discipline studying the properties of languages and  
1552 the meaning of programs, its properties, and methods of description. An excellent introduction is  
1553 the Software Foundations book series<sup>104</sup> Pierce et al. [2010]. There are many different styles of  
1554 describing program semantics; the most important for us will be denotational semantics.

1555 **Denotational semantics**<sup>105106107</sup> describes the meaning of programs by associating them with  
1556 objects of some set, the **semantic domain**. In other words, the denotational semantics of a language  
1557  $L$  is a total function from a program in that language to an element of the domain  $D$ :

$$\llbracket \bullet \rrbracket : L \mapsto D$$

1558 The domain is chosen based on the language and the information one wants to extract from  
1559 programs. For example, to find out the size of a program (here, expressions with addition), one can  
1560 take natural numbers as the domain:

$$\begin{aligned} \llbracket n \rrbracket &= 1 \\ \llbracket l + r \rrbracket &= 1 + \max(\llbracket l \rrbracket, \llbracket r \rrbracket) \end{aligned}$$

1561 If we are interested in the final result, we can compute it:

$$\begin{aligned} \llbracket n \rrbracket &= n \\ \llbracket l + r \rrbracket &= \llbracket l \rrbracket + \llbracket r \rrbracket \end{aligned}$$

1562 If the program has input, the domain will be a function  $\mathbb{N} \rightarrow \mathbb{N}$ :

$$\begin{aligned} \llbracket n \rrbracket(m) &= n \\ \llbracket l + r \rrbracket(m) &= \llbracket l \rrbracket(m) + \llbracket r \rrbracket(m) \\ \llbracket input \rrbracket(m) &= m \end{aligned}$$

---

<sup>103</sup>[https://en.wikipedia.org/wiki/Semantics\\_\(computer\\_science\)](https://en.wikipedia.org/wiki/Semantics_(computer_science))

<sup>104</sup><https://coq.vercel.app/ext/sf/>

<sup>105</sup>[https://en.wikipedia.org/wiki/Denotational\\_semantics](https://en.wikipedia.org/wiki/Denotational_semantics)

<sup>106</sup>[https://en.wikibooks.org/wiki/Haskell/Denotational\\_semantics](https://en.wikibooks.org/wiki/Haskell/Denotational_semantics)

<sup>107</sup>(youtube) The Lost Art of Denotational Semantics — Eric Meyer.

1563 Thus, *a program is merely a syntactic notation for some element of the semantic domain*. There  
1564 are many possible domains; they can even be games<sup>108</sup>.

1565 **Exercise 34** *Into which domain is it reasonable to interpret programs in a language with mutable*  
1566 *integer variables? And in a nondeterministic language?*

1567 It is easy to see that the denotational semantics of a language is simply an interpreter, only  
1568 written in the language of mathematics. Such interpreters are placed at the base of the tower of  
1569 interpreters when the goal is to study properties of languages and programs rather than to execute  
1570 them.

1571 Also, we can implement an interpreter in some real language, and it too will define the semantics  
1572 of the object language. However, the formality of such a definition will depend on the formality of  
1573 the meta-language's semantics description. Such interpreters are called **definitional**; they define  
1574 the semantics of a language, typically sacrificing efficiency for clarity. The relationships between the  
1575 object language and the meta-language are studied in the classical papers Reynolds [1972, 1998]<sup>109</sup>.

1576 We will use definitional interpreters to define the semantics of new languages, using Haskell as  
1577 the meta-language. And as domains we will take Haskell types. And we will interpret programs not  
1578 into the set of functions between natural numbers, but, say, into the type `Nat -> Nat` in Haskell<sup>110</sup>.  
1579 Thus, the denotational semantics of the language of sums with input will be written as follows:

```
1 eval :: Prog -> (Nat -> Nat)
2 eval = \case
3   Val n -> \_ -> n
4   Plus l r -> \m -> eval l m + eval r m
5   Input -> \m -> m
```

1580 A semantics is called **compositional** if the meaning of constructs depends only on the meaning  
1581 of subconstructs. In other words, if the denotational semantics is a fold over the program (Fig. 17)  
1582 and can be written in terms of a catamorphism 4.3.3.

1583 Another popular style of describing semantics is **operational semantics**, which represents the  
1584 meaning of a program as a sequence of computation steps. This can be either sequential rewriting  
1585 of the expression itself, or rewriting the state of some abstract machine. Operational semantics,  
1586 in turn, is defined as an unfold (or anamorphism 4.3.3) of the sequence of computation steps  
1587 from the program. Here one can clearly see a certain duality between denotational and operational  
1588 semantics Hutton [1998].

#### 1589 5.1.4 Embedded domain-specific languages (eDSL)

1590 By **domain-specific languages (DSL)**<sup>111</sup> one often means specialized languages for specific  
1591 domains, such as database queries or document formatting. As a rule, such languages are not  
1592 Turing-complete.

---

<sup>108</sup>[https://en.wikipedia.org/wiki/Game\\_semantics](https://en.wikipedia.org/wiki/Game_semantics)

<sup>109</sup>The notion of continuation, actively used by the author, will be considered later in this course (section 6).

<sup>110</sup><https://okmij.org/ftp/Denotational.html>

<sup>111</sup>[https://en.wikipedia.org/wiki/Domain-specific\\_language](https://en.wikipedia.org/wiki/Domain-specific_language)




Figure 17: Denotational semantics defines the meaning of syntactic constructs through operations on the domain Hutton [1998].

1593 In this course, however, we will consider any domain-specific specialization of a general-purpose  
 1594 language to be a domain-specific language<sup>112</sup>. This follows from the consideration that code should  
 1595 read like well-written prose with appropriate word usage, providing the reader with only the necessary  
 1596 amount of detail, hiding the inessential behind conventions and terminology.

1597 **Standalone domain-specific languages** are languages that have their own concrete syntax as  
 1598 well as programming tools (IDE, runtime... ). Examples: SQL, AWK, Antlr...<sup>113</sup>

1599 **Embedded domain-specific languages (eDSL)** are languages that leverage the infrastructure  
 1600 of other languages. They are usually implemented as libraries for programs in an already existing  
 1601 general-purpose language. They do not have a fully independent concrete syntax. Examples: ORM,  
 1602 string processing functions, parser combinator libraries... .

1603 **Deep eDSL** — terms in such a language build an abstract syntax tree for further interpretation:

```
1 f :: Int -> Int
2 f = eval $ Const 1 `Plus` Input
```

1604 Interpreters that take trees as input and interpret them into a semantic domain are called  
 1605 **initial**<sup>114</sup>; we will also call them “classical.” We will start with them. Classical interpreters are  
 1606 useful, for example, for implementing the “last language” — the program’s interface to the outside  
 1607 world, and as a foundation for our further constructions.

1608 However, one can notice that the intermediate tree that is produced typically does not interest  
 1609 us. We only care about obtaining the domain element, which we already know how to use directly.

1610 **Shallow eDSL** skip the tree construction stage and directly build the value in the semantic domain.

1611 We will call such interpreters **final**.

```
1 cnst :: Int -> (Int -> Int)
2 cnst x _ = x
```

<sup>112</sup>[https://en.wikipedia.org/wiki/Language-oriented\\_programming](https://en.wikipedia.org/wiki/Language-oriented_programming).

<sup>113</sup>There is an excellent practical book Nystrom that covers building interpreters and simple virtual machines. At the same time it covers the creation of a full-fledged language in all its aspects, from parsing to memory management.

<sup>114</sup>The syntax type defined using `data` is the initial object of the category of interpretations.

```

3 input :: Int -> Int
4 input env = env

5 plus :: (Int -> Int) -> (Int -> Int) -> (Int -> Int)
6 plus l r env = l env + r env

7 f :: Int -> Int
8 f = cnst 1 `plus` input

```

1612 Interpreters are often called **observers**, which analyze terms and give them some meaning Gibbons  
1613 [2013]. One can notice that for deep eDSL one can write arbitrarily many different observers.  
1614 However, in the case of shallow embedding, observers are always id. We will discuss possible  
1615 solutions to this problem later in section 5.4.1.

1616 A discussion of terminology and comparison of approaches to DSL construction can be found  
1617 in Gibbons [2013]. A brief description of terms is in the Language Engineering course notes Hutchinson.

1618 Let us introduce one more important concept. A **meta-circular interpreter**<sup>115</sup> is an interpreter  
1619 that defines constructs of the object language through constructs of the meta-language Reynolds  
1620 [1972]. For example:

```

1 interpret term = case term of
2   App f t -> (interpret f) (interpret t)
3   If c t e -> if interpret c then interpret t else interpret e
4   ...

```

1621 The properties of the meta-language in this case largely determine the properties of the object  
1622 language Reynolds [1972, 1998]. In this course we will strive to reuse the capabilities of the meta-  
1623 language as much as possible.

1624 **Exercise 35** *Guess which properties the object language from the example above inherits.*

### 1625 5.1.5 Example: the Accelerate library

1626 An interesting example of an embedded language that lies somewhere between deep and shallow  
1627 is the Accelerate library<sup>116</sup> [Marlow, 2011, chapter 6]. It allows describing computations in Haskell  
1628 that will be executed on a GPU<sup>117</sup>.

1629 To execute something on a GPU, one needs to generate and compile Cuda code. Thus, the  
1630 Accelerate language must be a deep embedding in order to have the computation tree for translation  
1631 into Cuda in the most efficient manner.

<sup>115</sup>[https://en.wikipedia.org/wiki/Meta-circular\\_evaluator](https://en.wikipedia.org/wiki/Meta-circular_evaluator)

<sup>116</sup><https://hackage.haskell.org/package/accelerate>

<sup>117</sup>Another approach: Java code reflection, to extract a model of the code at runtime. However, such an approach does not provide static guarantees to the programmer and requires deep integration into the meta-language.

1632 At the same time, describing numerical computations as a tree is extremely inconvenient. It  
1633 would be nice to have familiar operators and higher-order functions for working with arrays on  
1634 the GPU. Therefore, Accelerate actually provides a shallow interface for building trees. Thus, for  
1635 expression trees, implementations of numeric type classes are defined, e.g. `Num`, where operations  
1636 simply extend the tree:

```
1 example :: Acc (Vector Int) -> Acc (Vector Int) -> Acc (Vector Int)
2 example xs ys = A.zipWith (+) xs ys
```

## 1637 5.2 Value types

1638 Let us consider the possible ways to implement languages whose values can have different types.

### 1639 5.2.1 Untyped tagless interpreters

1640 First, let us consider a trivial untyped language. By untyped we mean the absence of type  
1641 checking both before and during program execution. We define the abstract syntax of this language  
1642 as follows:

```
1 data Expr = Const Int | IsZero Expr | If Expr Expr Expr
```

1643 Values arising during the execution of programs in this language will be represented by values  
1644 of types `Bool` and `Int` of the Haskell language. Accordingly, the semantic domain of a program in  
1645 this language is either `Bool` or `Int`, depending on the program itself.

```
1 evalUnsafe :: Expr -> forall res . res
2 evalUnsafe = \case
3   Const val -> unsafeCoerce val
4   IsZero cond -> unsafeCoerce $ evalUnsafe @Int cond == 0
5   If c t e -> if evalUnsafe c then evalUnsafe t else evalUnsafe e
```

1646 Here `unsafeCoerce` is used to circumvent Haskell's static type system and simply execute  
1647 programs in our untyped language. We are entitled to do this because `Int` and `Bool` in Haskell have  
1648 the same size. Writing an incorrect program in this language or choosing the wrong interpretation  
1649 domain leads to a crash.

### 1650 5.2.2 Typed tagged interpreters

1651 To achieve some safety during execution, we will attach tags to values that will be available at  
1652 runtime. We introduce the following algebraic type:

```
1 data RtValue = RtBool Bool | RtInt Int
```

1653 Now the semantic domain will be the type `RtValue`, and the interpreter can check types at  
1654 runtime:

```

1  evalRt :: Expr -> RtValue
2  evalRt (IsZero expr) = case evalRt expr of
3    RtBool value -> error "Type error"
4    RtInt value -> RtBool (value == 0)
5  -- ...

```

1655 The safety situation has definitely improved, but type checking at runtime is already too late:  
 1656 it requires additional performance overhead and makes testing more expensive.

1657 This approach is often called **dynamic typing**, when we attribute values with some type  
 1658 information for use at runtime (see also 3.4.6).

### 1659 5.2.3 Typed tagless interpreters

1660 Let us describe the type system of our small language.

$$\frac{n : \text{Int}}{\text{Const } n : \text{int}} \text{Const} \quad \frac{n : \text{int}}{\text{IsZero } n : \text{bool}} \text{IsZero} \quad \frac{c : \text{bool} \quad t : \tau \quad e : \tau}{\text{If } c \text{ t e} : \tau} \text{If}$$

1661 We can reuse Haskell types as type tags:

$$\text{int} \rightsquigarrow \text{Int}, \text{bool} \rightsquigarrow \text{Bool}$$

1662 Notice that with the help of a generalized algebraic data type `Expr ty` (2.1.4), we have precisely  
 1663 encoded these typing rules. In other words, we have obtained a statically typed programming  
 1664 language by reusing Haskell's type system.

```

1  data Expr ty where
2    Const :: Int -> Expr Int
3    IsZero :: Expr Int -> Expr Bool
4    If :: forall ty . Expr Bool -> Expr ty -> Expr ty -> Expr ty

5  eval :: Expr ty -> ty
6  eval = \case
7    Const n -> n
8    IsZero e -> eval e == 0
9    If c t e -> if eval c then eval t else eval e

```

1665 Thanks to static typing, we can dispense with tagging values during interpretation without losing  
 1666 safety.

## 1667 5.3 Bindings and first-class functions

1668 In this section we consider techniques and concepts related to implementing first-class functions  
 1669 and bindings in general. Since this functionality is typically already implemented in the meta-  
 1670 language, we will strive to reuse it as much as possible.

1671 `let`-bindings can be represented through first-class functions as follows:

$$\mathbf{let } x := N \mathbf{ in } M \equiv (\lambda x. M) N$$

1672 Recall that first-class functions can be eliminated using defunctionalization, considered earlier 3.1.7.

### 1673 5.3.1 Name semantics

1674 There are several ways to assign semantics to identifiers.

1675 **Dynamic scoping** — the values of free variables of a function depend on the scope at the call  
1676 site. That is, name resolution happens at the moment the variable is accessed. For example, the  
1677 following code prints 42:

```
1  val x = 4
2  val f = () => x + 1
3  val x = 41
4  println(f())
```

1678 This approach is simpler to implement and was used in early versions of Lisps, for example.  
1679 However, in this case functions are not a reliable abstraction barrier; ideally all free variables should  
1680 be part of the signature (we will return to this in the chapter on effect systems 8).

1681 **Lexical/static scoping** — variables are bound to values at the point of function declaration; at  
1682 the point of invocation the result depends only on the parameters (modulo mutable state<sup>118</sup>). The  
1683 word “lexical” is often used in languages when we can determine something from the source code  
1684 without running the program. Thus, the code from the example above would print 5.

1685 In the rest of this section we will discuss various ways of implementing first-class functions with  
1686 static variable scoping, of which there are, in fact, a great many<sup>119</sup>.

### 1687 5.3.2 Substitutions

1688 As one can notice, in classical lambda calculus, substitutions from beta-reduction (recalled  
1689 in section 1.1) provide static scoping. Indeed, the argument is immediately substituted into all  
1690 occurrences of the variable, so it does not remain free but simply disappears.

$$(\lambda x. (\lambda x. \lambda y. x + y) 4) 41 \rightsquigarrow (\lambda x. (\lambda y. 4 + y)) 41$$

1691 This approach is not the most efficient, because each application requires rewriting the entire  
1692 function body (!). At the same time, it is fairly simple to implement for certain representations of  
1693 lambda terms. Let us consider an example of such a representation — **locally nameless** Charguéraud  
1694 [2012].

---

<sup>118</sup>For example, in Kotlin, mutable variables can be captured in lambdas. Changes outside are observable inside the lambda, and vice versa. This can sometimes be very convenient, but often leads to very non-obvious behavior.

<sup>119</sup><https://jesper.cx/posts/1001-syntax-representations.html>

```

1 data Term var
2   = Var var
3   | App (Term var) (Term var)
4   | Lam (Term (Maybe var))

```

1695 In this representation one can choose any type for naming free variables:

```

1 example :: Term String
2 example = Var "x" `App` Var "y" -- x y

```

1696 Adding each bound variable adds a new inhabitant `Nothing` to the variable type for referring to the  
1697 nearest bound variable:

```

1 -- λx. x y
2 example1 = Lam $ Var Nothing `App` Var (Just "y")
3 -- λx y. x y z
4 example2 = Lam $ Lam $ Var (Just Nothing) `App` Var Nothing `App` Var (Just (Just "z"))

```

1698 Monadic bind is the implementation of substitution for such terms:

```

1 instance Monad Term where
2   (>>=) :: Term var -> (var -> Term var') -> Term var'
3   Var var >>= subst = subst var
4   App l r >>= subst = App (l >>= subst) (r >>= subst)
5   Lam t >>= subst = Lam $ t >>= \case
6     Nothing -> Var Nothing
7     Just var -> Just <$> subst var

```

1699 **Exercise 36** *Think about why `fmap Just` is needed in the last line.*

1700 Accordingly, a call-by-name interpreter for this lambda calculus will look as follows:<sup>120</sup>

```

1 eval :: Term var -> Term var
2 eval = \case
3   Var var -> Var var
4   App f arg -> case eval f of
5     Lam body -> eval $ body >>= maybe arg Var
6     t -> App t (eval arg)
7   Lam t -> Lam (eval t)

```

---

<sup>120</sup>Since we are considering classical  $\lambda$ -calculus, the resulting value is also a term, but in normal form.

### 1701 5.3.3 Environment

1702 One can perform substitution of variable values lazily, propagating an environment that maps free  
1703 variables to terms. Things have not yet improved much overall, but we have obtained compositional  
1704 semantics from non-compositional semantics by making context dependencies explicit (more on this  
1705 later 5.4.3).

```
1 data Term1 = Var1 String | App1 Term1 Term1 | Lam1 String Term1
2 type Env = Map String Term1

3 eval1 :: Term1 -> Env -> Term1
4 eval1 term env = case term of
5   Var1 name -> Map.findWithDefault (Var1 name) name env
6   App1 f arg -> case eval1 f env of
7     Lam1 name body -> eval1 body (Map.insert name arg env)
8     t -> App1 t (eval1 arg env)
9   Lam1 name body ->
10     let env' = Map.delete name env in
11     Lam1 name (eval1 body env')
```

1706 **Exercise 37** *Explain why the environment is modified on line 10?*

1707 If the `Lam1` branch does not recursively traverse the subterm and substitute variable values, the  
1708 information about free variable values in it will be lost and we will get dynamic scoping instead  
1709 of static scoping. To restore static scoping, the `Lam1` branch of the interpreter must construct a  
1710 closure that includes the current environment (see 5.3.4).

### 1711 5.3.4 Closures

1712 To avoid eagerly performing substitution into function bodies while preserving static scoping, we  
1713 add one more construct, a **closure**<sup>121122</sup> [Nystrom, chapter 11]. It will store the context in which  
1714 the corresponding term should be executed.

```
1 data Term1 = Var1 String | App1 Term1 Term1 | Lam1 String Term1
2             | Closure Env String Term1 -- only for computations
3 type Env = Map String Term1

4 eval1 :: Term1 -> Env -> Term1
5 eval1 term env = case term of
6   Var1 name -> Map.findWithDefault (Var1 name) name env
7   App1 f arg -> case eval1 f env of
```

---

<sup>121</sup>[https://en.wikipedia.org/wiki/Closure\\_\(computer\\_programming\)](https://en.wikipedia.org/wiki/Closure_(computer_programming))

<sup>122</sup>The term closure was proposed by Peter Landin, along with a bunch of other things.

```

8     Closure env' name body ->
9         let arg' = eval1 arg env in
10        eval1 body (Map.insert name arg' env')
11    t -> App1 t (eval1 arg env)
12    Lam1 name body -> Closure env name body

```

1715 Closures are typically used in industrial languages as the runtime representation of higher-order  
1716 functions. During compilation, **closure conversion** is performed first — higher-order functions are  
1717 represented as a pair of an environment and a pointer to a function that takes the environment as  
1718 an extra argument. Now that the function contains no free variables, **lambda lifting**<sup>123</sup> is done —  
1719 it is lifted to the top level. Detailed examples can be found in the Harvard slides Chong.

### 1720 5.3.5 Typed context

1721 Let us consider the encoding described, for example, in Kiselyov [2012a].

1722 First, let us learn to use Haskell's type system to check the validity of environment lookups. We  
1723 represent the environment as a list of types encoded using nested pairs:

```

1 (4, (4.0, "hello")) :: (Int, (Double, String))

```

1724 Environment lookups will be encoded as numbers in unary notation. The type of the number (a  
1725 typed reference into the context) specifies the set of environments *env* from which a value of type  
1726 *ty* can be extracted at the given position.

```

1 data Ref env ty where
2   Here :: Ref (ty, env) ty
3   There :: Ref env ty -> Ref (ty', env) ty

```

1727 For example, the type of the number 1 asserts that it can be used to extract a value of type *ty*  
1728 from a context where a value of the corresponding type is at position one (zero-indexed):

```

1 There Here :: Ref (ty', (ty, env)) ty

```

1729 Now we can encode a typed safe lookup into the context:

```

1 envLookup :: env -> Ref env ty -> ty
2 envLookup env ref = case (ref, env) of
3   (Here, (x, _)) -> x
4   (There ref', (_, env')) -> envLookup env' ref'

```

1730 **Exercise 38** *Can the pair be deconstructed directly on line 2? Explain.*

<sup>123</sup>[https://en.wikipedia.org/wiki/Lambda\\_lifting](https://en.wikipedia.org/wiki/Lambda_lifting)

### 1731 5.3.6 Meta-circular interpretation

1732 It would be extremely undesirable to implement bindings and first-class functions for an eDSL  
1733 by ourselves. Let us build a meta-circular interpreter (see 5.1.4) that will reuse the meta-language's  
1734 first-class functions to implement them in the object language.

1735 Terms will now be annotated not only with resulting types but also with the types of environments  
1736 needed for interpretation, considered earlier 5.3.5. An abstracted term has access to a larger  
1737 environment.

```
1 data Term2 env ty where
2   Var2 :: Ref env ty -> Term2 env ty
3   App2 :: Term2 env (arg -> res) -> Term2 env arg -> Term2 env res
4   Lam2 :: Term2 (arg, env) res -> Term2 env (arg -> res)
```

1738 Now we can interpret abstraction as a Haskell function and application as application:

```
1 eval2 :: Term2 env ty -> env -> ty
2 eval2 term env = case term of
3   Var2 ref -> env `envLookup` ref
4   App2 f arg -> (eval2 f env) (eval2 arg env)
5   Lam2 t -> \arg -> eval2 t (arg, env)
```

1739 **Exercise 39** *How did it happen that an extra argument needs to be accepted in the last line?*

1740 **Exercise 40** *Is this a call-by-value interpreter or call-by-name? What does this depend on?*

1741 **Exercise 41** *Think about which solution should be more performant, this one or the previous?*

1742 Note that now the functions of the object language at runtime are simply functions of the meta-  
1743 language. This means that in programs in the object language we *can fully reuse the meta-language!*  
1744 Let us add a construct that allows storing an arbitrary meta-language value in the tree:

```
1 data Term2 env ty where
2   Val2 :: ty -> Term2 env ty
3   -- ...
4 eval2 :: Term2 env ty -> env -> ty
5 eval2 term env = case term of
6   Val2 x -> x
7   -- ...
8 example :: Term2 env (Int -> Int)
9 example = Lam (Val2 (+) `App2` Val2 1 `App2` Var2 Here)
```

### 1745 5.3.7 Higher-order abstract syntax

1746 Another thing we are still doing manually is defining binders (and in unary notation, no less).  
1747 We want to reuse them from the meta-language. For this, we will store meta-language functions  
1748 directly in the syntax tree — use **higher-order abstract syntax (HOAS)**<sup>124</sup><sup>125</sup> Pfenning and Elliott  
1749 [1988]:

```
1 data Term3 ty where
2   Val3 :: ty -> Term3 ty
3   Plus :: Term3 Int -> Term3 Int -> Term3 Int
4   App3 :: Term3 (arg -> res) -> Term3 arg -> Term3 res
5   Lam3 :: (Term3 arg -> Term3 res) -> Term3 (arg -> res)

6 example3 :: Term3 Int
7 example3 = (Lam3 \x -> x `Plus` Val3 41) `App3` Val3 1
```

1750 The interpretation is very simple and completely meta-circular:

```
1 eval3 :: Term3 ty -> ty
2 eval3 term = case term of
3   Val3 x -> x
4   Plus l r -> eval3 l + eval3 r
5   App3 f arg -> (eval3 f) (eval3 arg)
6   Lam3 f -> \arg -> eval3 (f (Val3 arg))
```

1751 **Exercise 42** *Could Lam3 have been declared as follows?*

```
1 Lam3 :: (arg -> Term3 res) -> Term3 (arg -> res)
```

### 1752 5.3.8 Serialization of functions

1753 In this section we discussed possible implementations of first-class functions, i.e. functions that  
1754 can be used as flexibly as data. A natural question arises: can we serialize a first-class function and  
1755 send it for execution on another machine?

1756 A function consists of code and captured free variables in the case of static scoping. Accordingly,  
1757 if the code is in a serializable form (e.g. position-independent bytecode), it can in principle be sent  
1758 over the network and executed on another virtual machine instance. This is what Erlang does,  
1759 for example. However, this approach is inefficient, since the bytecode needs to be interpreted or  
1760 pre-compiled. Thus, Erlang sacrifices execution speed for horizontal scalability.

1761 If we guarantee that the same code runs on different cluster nodes, as is usually the case in  
1762 practice, a more efficient implementation can be achieved. For example, using defunctionalization (see 3.1.7),  
1763 we can serialize only the algebraic type objects encoding functions. Since the same code runs on the

---

<sup>124</sup>[https://en.wikipedia.org/wiki/Higher-order\\_abstract\\_syntax](https://en.wikipedia.org/wiki/Higher-order_abstract_syntax)

<sup>125</sup>What is higher-order in higher-order abstract syntax?

1764 other cluster node, we can deserialize the object there and execute it using `apply`. However, this  
1765 approach does not support modularity well (it is hard to split one algebraic type into many), and  
1766 `apply` performs decoding before executing the code each time (which, admittedly, can be neglected  
1767 given network overhead).

1768 The approach implemented in Haskell<sup>126127</sup> allows assigning each function without free variables  
1769 a statically known address, the same for all application instances. Then one can construct a  
1770 serializable closure through a sequence of partial applications:

```
1 data Closure a where
2   StaticPtr :: StaticPtr b -> Closure b
3   Encoded :: ByteString -> Closure ByteString
4   Ap :: Closure (b -> c) -> Closure b -> Closure c

5 main = send "some-node" $
6   closure (static factorial) `closureAp` closurePure 10
```

1771 More details can be found in the foundational paper on Cloud Haskell Epstein et al. [2011].  
1772 From a practical perspective — in the book [Marlow, 2011, chapter 16].

1773 **Exercise 43** *Is it necessary to explicitly add the free variable (\*) (the multiplication operator) to*  
1774 *the closure in the factorial implementation?*

## 1775 5.4 Tagless final interpreters

1776 As we established earlier (5.1.2), programming consists of writing interpreters on top of each  
1777 other. Interpreters define the semantics of new languages (5.1.3). In the classical form, a language  
1778 is defined as a set of trees, and an interpreter maps trees to meta-language objects. If the language  
1779 is embedded, this approach is called deep embedding (see 5.1.4), and the corresponding interpreter  
1780 is initial.

$$[[\bullet]] : L \rightarrow D$$

1781 One can notice that ultimately we only use the domain element to which the interpreter maps  
1782 the program. The program itself is merely a convenient syntactic notation for a domain element and  
1783 is an intermediate step, not an end in itself. At the same time, the domain for embedded languages  
1784 defined by interpreters consists of meta-language objects. Can we bypass the stage of interpreting  
1785 our own syntax and directly build the domain object in the meta-language syntax? Yes, such an  
1786 embedding is called shallow embedding (5.1.4); this chapter is about it.

---

<sup>126</sup><https://blog.ocharles.org.uk/blog/guest-posts/2014-12-23-static-pointers.html>

<sup>127</sup><https://hackage.haskell.org/package/distributed-closure>

### 1787 5.4.1 Different interpretations for shallow embedding

1788 As we learned earlier (4.4), any data structure can be represented as a fold. And, moreover, in  
1789 the end one can do without a single data constructor (as in Church lists, for example): the algebra  
1790 is represented by a set of functions, each responsible for folding a particular constructor.

```
1 Fix f  $\cong$  forall a . (f a -> a) -> a
2 List e  $\cong$  forall a . (e -> a -> a) -> a -> a
```

1791 Thus, instead of constructing a language tree and then interpreting (folding) it, we can directly  
1792 construct a term of type  $\forall a . (f a \rightarrow a) \rightarrow a$ . By providing it with a domain type  $a$  and an  
1793 algebra  $f a \rightarrow a$  (or as a bundle of functions), we immediately obtain an element of the desired  
1794 domain. To specify a different interpretation, one needs to pass a different domain type and algebra:

```
1 example :: (Int -> a) -> (a -> a -> a) -> a
2 example cnst plus = cnst 1 `plus` cnst 41

3 ghci> example show (\l r -> l ++ " + " ++ r)
4 "1 + 41"
```

1795 If we fix the interpretation, the argument functions can be implemented statically and simply  
1796 referenced in the term. Thus, declaring functions can be thought of as extending some embedded  
1797 domain-specific language. General discussions about shallow embeddings, folds, and libraries can be  
found in Gibbons [2013], Gibbons and Wu [2014]. Compare:

Deep	Shallow
The language syntax is defined by the set of valid tree nodes	A function declaration defines a new tree node: a call to that function
The interpreter, upon encountering each node, executes the corresponding meta-language code (a pattern-match branch) after computing subtrees	The interpreter, upon encountering a call, executes the function body code after computing arguments

1798

### 1799 5.4.2 Going all the way

1800 Let us return to the uncurried version of folds: this is now a type accepting a tuple of functions.  
1801 And a tuple of functions can be replaced by a type class. Then the class declaration will define the  
1802 syntax of the embedded language, and instances for domains — the implementation. This approach  
1803 is called **tagless final encoding**<sup>128</sup> and is effectively Church encoding of data with type classes and  
1804 a set of tricks Carrette et al. [2007], Kiselyov [2012a].

1805 Let us again consider the language with addition:

<sup>128</sup><https://okmij.org/ftp/tagless-final/>

```

1 data Expr = Const Int | Plus Expr Expr
2 eval :: Expr -> Int
3 eval = \case Const x -> x; Plus l r -> eval l + eval r

```

1806 Or via catamorphism:

```

1 data ExprF rec = Const Int | Plus rec rec
2 eval :: Fix ExprF -> Int
3 eval = cata \case Const x -> x; Plus l r -> l + r

```

1807 The corresponding tagless final encoding will look as follows:

```

1 class Expr domain where
2   cnst :: Int -> domain
3   plus :: domain -> domain -> domain
4 instance Expr Int where
5   cnst x = x
6   plus l r = l + r

```

1808 Now we can construct an ordinary Haskell term, specify the domain, and the type class machinery  
1809 will supply the appropriate algebra automatically:

```

1 example :: forall domain . Expr domain => domain
2 example = cnst 1 `plus` cnst 41
3 ghci> example :: Int
4 42

```

1810 To add another interpretation, we implement an instance for a different domain:

```

1 instance Expr String where
2   cnst x = show x
3   plus l r = l <> " + " <> r
4 ghci> example :: String
5 "1 + 41"

```

### 1811 5.4.3 Recovering compositionality of semantics

1812 A **semantics** is called **compositional** if the semantics of constructs depends only on the  
1813 semantics of subconstructs (see 5.1.3). In other words, it can be defined by a catamorphism  
1814 and, accordingly, by a type class instance in tagless final. That is, to be able to build a tagless

1815 final implementation for any semantics, one needs to be able to universally turn non-compositional  
1816 semantics into compositional ones.

1817 Various code transformations are typically not compositional. As an example, consider pushing  
1818 through unary negations:

```
1 data Expr1 = Lit Int | Add Expr1 Expr1 | Neg Expr1

2 transform1 :: Expr1 -> Expr1
3 transform1 = \case
4   Lit x -> Lit x
5   Add l r -> Add (transform1 l) (transform1 r)
6   Neg (Lit x) -> Lit (-x) -- problem
7   Neg (Neg e) -> transform1 e -- problem
8   Neg (Add l r) -> Add (transform1 (Neg l)) (transform1 (Neg r)) -- problem
```

1819 To recover compositional semantics, one needs to make context dependencies explicit using an  
1820 arrow domain Kiselyov [2012a]:

```
1 data Ctx = CtxPos | CtxNeg
2 flipCtx = \case CtxPos -> CtxNeg; CtxNeg -> CtxPos

3 transform1' :: Expr1 -> (Ctx -> Expr1)
4 transform1' expr = case expr of
5   Lit x -> \case CtxNeg -> Lit (-x); CtxPos -> Lit x
6   Neg e -> \ctx -> transform1' e (flipCtx ctx)
7   Add l r -> \ctx -> Add (transform1' l ctx) (transform1' r ctx)
```

1821 From here one can obtain the tagless final version:

```
1 class Expr2 d where
2   lit :: Int -> d
3   add :: d -> d -> d
4   neg :: d -> d

5 instance Expr2 d => Expr2 (Ctx -> d) where
6   lit x = \case CtxNeg -> lit (-x); CtxPos -> lit x
7   neg e = \ctx -> neg e (flipCtx ctx)
8   add l r ctx = add (l ctx) (r ctx)
```

#### 1822 5.4.4 Typed tagless final interpreter

1823 Let us consider our example of tagless initial encoding 5.2.3:

```
1 data Expr ty where
2   Const :: Int -> Expr Int
```

```

3   IsZero :: Expr Int -> Expr Bool
4   If :: forall ty . Expr Bool -> Expr ty -> Expr ty -> Expr ty

5   eval :: Expr ty -> ty
6   eval = \case
7     Const x   -> x
8     IsZero t  -> eval t == 0
9     If c t e  -> if eval c then eval t else eval e

```

1824 To obtain a final encoding, we parameterize the domain by the resulting type of the expression:

```

1   class Expr (domain :: Type -> Type) where
2     cnst :: Int -> domain Int
3     isZero :: domain Int -> domain Bool
4     if' :: forall ty . domain Bool -> domain ty -> domain ty -> domain ty

5   instance Expr Identity where
6     cnst x = Identity x
7     isZero (Identity x) = Identity (x == 0)
8     if' (Identity c) t e = if c then t else e

```

1825 **Exercise 44** Which domain would be suitable for printing expressions?

### 1826 5.4.5 Meeting old friends: `Applicative`, `Monad`

1827 Consider the following language in initial encoding with higher-order abstract syntax (see 5.3.6,  
1828 5.3.7). On the right we rewrite it in tagless final, choosing appropriate names for pieces of syntax.

<pre> 1   data Expr s ty where 2     Val   :: ty -&gt; Expr s ty 3     App   :: Expr s (arg -&gt; res) 4           -&gt; Expr s arg 5           -&gt; Expr s res 6 7 8     LetIn :: Expr s ty 9           -&gt; (ty -&gt; Expr s ty') 10          -&gt; Expr s ty' 11 12 13    Get   :: Expr s s 14    Put   :: s -&gt; Expr s () </pre>	<pre> 1   class Applicative domain where 2     pure  :: ty -&gt; domain ty 3     (&lt;*&gt;) :: domain (arg -&gt; res) 4           -&gt; domain arg 5           -&gt; domain res 6 7   class Monad domain where 8     (&gt;=&gt;) :: domain ty 9           -&gt; (ty -&gt; domain ty') 10          -&gt; domain ty' 11 12   class MonadState s domain where 13     get  :: domain s 14     put  :: s -&gt; domain () </pre>
--	---

1830 We implement a function that modifies a value in our new language:

```
1  --
2  modify :: (s -> s) -> Expr s s
1831 3  modify f =
4    Get `LetIn` \x ->
5    Put (f x) `LetIn` \() ->
6    Val x

1  --
2  modify :: (s -> s) -> State s s
3  modify f =
4    get >>= \x ->
5    put (f x) >>= \() ->
6    pure x
```

1832 And the painfully familiar interpretation:

```
1  eval :: Expr s ty -> s -> (s, ty)
2  eval = \case
3
4
5  Val x -> \s -> (s, x)
6  App fs xs -> \s1 ->
7    let (s2, f) = eval fs s1 in
8    let (s3, x) = eval xs s2 in
1833 9    (s3, f x)
10
11
12  LetIn comp k -> \s ->
13    let (s', x) = eval comp s in
14    eval (k x) s'
15
16
17  Get -> \s -> (s, s)
18  Put s -> \_ -> (s, ())

1  newtype State s a = State
2    { runState :: s -> (s, a) }
3
4  instance Applicative (State s) where
5    pure x = State \s -> (s, x)
6    fs <*> xs = State \s1 ->
7      let (s2, f) = runState fs s1 in
8      let (s3, x) = runState xs s2 in
9      (s3, f x)
10
11  instance Monad (State s) where
12    comp >>= k = State \s ->
13      let (s', x) = runState comp s in
14      runState (k x) s'
15
16  instance MonadState s (State s) where
17    get = State \s -> (s, s)
18    put s' = State \s -> (s', ())
```

1834 Thus, *applicative functors are a meta-circular language with application*, and *monadic bind is*  
1835 *essentially let-in in higher-order syntax*.

1836 If we recall that  $\text{let } x = M \text{ in } N \equiv (\lambda x. N) M \equiv [x \mapsto M] N$ , then we understand that monadic  
1837 bind is the substitution  $a \rightarrow m b$  of the computation  $m b$  in place of results of type  $a$  in the  
1838 computation  $m a$ . Or, if we take an expression tree as  $m$ , then bind is the substitution of subtrees  
1839 in place of variables:

```
1  data Expr var = Var var | Empty | Append (Expr var) (Expr var)
2  instance Monad Expr where
3    (>>=) :: Expr var -> (var -> Expr var') -> Expr var'
4    Var name >>= subst = subst name
5    Empty >>= _ = Empty
```

```
6 Append l r >>= subst = Append (l >>= subst) (r >>= subst)
```

1840 The idea to use the category-theoretic notion of a monad<sup>129</sup> in  $\lambda$ -calculus was originally proposed  
1841 in Moggi [1988] for more conveniently writing denotational semantics<sup>130</sup>. This turned out to be so  
1842 convenient at the time that monads began to be used ubiquitously in functional programming to  
1843 extend simple functional languages with various powerful capabilities without changing the languages  
1844 themselves Wadler [1990, 1992]. Moggi’s principle was formulated:

1845 “Computations of type  $\alpha$  correspond to values of type  $f \alpha$ .”

1846 The use of applicative functors was proposed significantly later, to avoid naming intermediate  
1847 computation steps when this is unnecessary McBride and Paterson [2008]. Thus, applicatives  
1848 provide an embedded language of expressions, while monads provide a language of statements.

## 1849 5.5 Expression problem

1850 The **expression problem**<sup>131</sup> is a criterion for the expressiveness of a programming language,  
1851 formulated by Wadler in 1998<sup>132</sup>. The question is: how easy is it to extend the syntax of an  
1852 embedded language and add new interpretations? In other words, how easy is it to add new kinds  
1853 of data and processing methods.

1854 By “easy” we mean locality: does one need to modify various pieces of code for this? For example,  
1855 if the language syntax is defined by an ordinary algebraic data type, then adding a new interpretation  
1856 is “easy” — just add a new recursive function, but adding a new syntactic construct is “hard” — all  
1857 interpreters must be changed:

```
1 data Expr
2   = Const Int
3   | Plus Expr Expr -- adding
4
4 eval :: Expr -> Int
5 eval = \case Const x -> x; Plus l r -> eval l + eval r
6
6 show :: Expr -> String
7 show = \case Const x -> show x; Plus l r -> show l ++ " + " ++ show r
```

1858 If the language is defined, for example, using inheritance, then conversely, extending the syntax  
1859 is easy — add a new class, but adding an interpretation is hard — add a method implementation in  
1860 each class:

---

<sup>129</sup>[https://ncatlab.org/nlab/show/monad+\(in+computer+science\)](https://ncatlab.org/nlab/show/monad+(in+computer+science))

<sup>130</sup>The notion of effect, often accompanying monads, will be considered later 7.

<sup>131</sup>[https://en.wikipedia.org/wiki/Expression\\_problem](https://en.wikipedia.org/wiki/Expression_problem)

<sup>132</sup><https://homepages.inf.ed.ac.uk/wadler/papers/expression/expression.txt>

```

1 interface Lang {
2     fun eval(): Int
3     fun show(): String
4 }
5 class Const(val x: Int) : Lang {
6     override fun eval() = x
7     override fun show() = x.toString()
8 }
9 class Plus(val l: Lang, val r: Lang) : Lang {
10    override fun eval() = l.eval() + r.eval()
11    override fun show() = "$l + $r"
12 }

```

1861 It turns out there exist approaches that achieve “ease” along both dimensions. We will devote  
1862 much attention to them in this course.

1863 Indeed, as we discussed earlier 5.1.1, programs are a series of interpreters. For the same fight  
1864 against complexity, it is important to be able to describe these interpreters modularly — defining  
1865 parts of languages independently of each other and assembling the needed languages on the spot  
1866 from ready-made blocks. This helps compose programs from simple reusable components, each  
1867 with a clear area of responsibility.

1868 The expression problem has arisen and been solved many times: expression problem, stable  
1869 denotations, extensible (modular) interpreters. It took a long time before the understanding emerged  
1870 that all of these are about the same thing<sup>133</sup>.

### 1871 5.5.1 Coproduct of functors

1872 Let us use the representation of data as a fixed point of a functor (see 4.3.2). As a model  
1873 language we take the expression language with addition:

```

1 data Basic rec = Const Int | Plus rec rec
2 algBasic :: Basic Int -> Int
3 algBasic = \case Const x -> x; Plus l r -> l + r
4 evalBasic :: Fix Basic -> Int
5 evalBasic = cata algBasic

```

1874 Note that the sum (coproduct) of shape functors gives a shape functor whose algebra is obtained  
1875 from the algebras of the components<sup>134</sup>:

```

1 data (l :+: r) rec = L (l rec) | R (r rec)

```

<sup>133</sup><https://okmij.org/ftp/Computation/having-effect.html>

<sup>134</sup>We use the TypeOperators extension.

```

2  (\/) :: (l a -> a) -> (r a -> a) -> ((l :+: r) a -> a)
3  phi \ / psi = \case L l -> phi l; R r -> psi r

```

1876 Let us extend our language with reading a number from the environment. Following the denotational  
1877 semantics considered earlier 5.1.3, we choose the function `Int -> Int` as the domain:

```

1  data Basic rec = Const Int | Plus rec rec
2  data Input rec = Input

3  algBasic' :: Basic (Int -> Int) -> Int -> Int
4  algBasic' = \case Const x -> \_env -> x; Plus l r -> \env -> l env + r env

5  algInput :: Input (Int -> Int) -> Int -> Int
6  algInput = \case Input -> \env -> env

```

1878 Thus, we have achieved the ability to separately define parts of the syntax and semantics of a  
1879 language and assemble the needed language on the spot. Let us write a program in our language  
1880 with implicit immutable state:

```

1  f :: Int -> Int
2  f = cata (algBasic' \ / algInput) $
3    In (L (Plus (In (L (Const 1))) (In (R Input)))) -- 1 + input

```

1881 However, note that we have not yet fully solved the problem, since interpreting the new construct  
1882 `Input` required a more complex domain and we had to rewrite the interpretation of the old one:  
1883 from `algBasic` to `algBasic'`. Now we understand why stable denotations is another name for the  
1884 expression problem 5.5. We will complete this solution later 7.

## 1885 5.5.2 Product of algebras

1886 Let us build a tagless final representation for the interpreter from the previous section. Now the  
1887 syntax is defined by type classes:

```

1  class Basic domain where
2    cnst :: Int -> domain
3    plus :: domain -> domain -> domain

4  class Input domain where
5    input :: domain

```

1888 As before, we take functions `Int -> Int` as the domain and implement the interpretation:

```

1  instance Basic (Int -> Int) where
2    cnst value = \_ -> value

```

```
3   plus l r = \env -> l env + r env
4   instance Input (Int -> Int) where
5     input = \env -> env
```

1889 Now, to assemble the language on the spot, one can simply take advantage of the composability  
1890 of type class contexts; all the rest of the work:

```
1   example :: forall domain . (Basic domain, Input domain) => domain
2   example = cst 1 `plus` input
3   ghci> (example :: Int -> Int) 41
4   42
```

## 1891 6 Continuations

1892 Since the early 1960s, continuations have appeared in the literature multiple times in various  
1893 forms and diverse applications Reynolds [1993], Landin [1997], until in the 1970s Wadsworth coined  
1894 a general term and a unified concept — **continuation**<sup>135</sup> — “the meaning of the rest of the program”.

1895 The initial impetus for these reflections was the Algol 60 language, which had a non-trivial  
1896 mechanism of labels and jumps. The problem was both the implementation of the semantics and  
1897 its denotational description together with translation into lambda calculus. Indeed, how does one  
1898 mathematically describe `goto`? In which domain should one seek the semantics of such programs?  
1899 How does one write a definitional interpreter mapping a program into this domain? The solution  
1900 turned out to be the ability to refer to the semantics of the remainder of the program, the  
1901 continuation, at a given point (for example, at a label).

1902 The language designer must always think about continuations, since the evaluator at each  
1903 point of a program needs to know what to execute next. And as we have already understood,  
1904 every programmer is the author of many languages that bring them closer to solving their problem.  
1905 Moreover, many everyday language constructs explicitly or implicitly directly manipulate continuations.

### 1906 6.1 The concept of continuations

1907 Let us consider the semantics of an expression from the operational point of view, as a sequence  
1908 of rewriting steps (see 5.1.3). If we carefully examine each step, we discover that it consists of two  
1909 stages: finding a subexpression (redex) in which an elementary computation step can be performed,  
1910 executing that step, and so on<sup>136</sup>. During the search for a redex, the expression is split into two  
1911 parts (see 18):

- 1912 • The focus — the subexpression in which we search for a redex;
- 1913 • The continuation — the remainder of the expression with a “hole,” indicating the place where  
1914 the result of the computation step should be plugged in.

1915 As a rule, continuations exist outside user code as the state of the interpreter, which at each  
1916 successive moment in time needs to remember what code to execute next and in what state.  
1917 However, languages provide users with many constructs that allow controlling continuations (see  
1918 19):

- 1919 • The function `exit` discards the entire program continuation;
- 1920 • The construct `try-catch` allows discarding part of the continuation up to the exception  
1921 handler and executing the remainder;
- 1922 • The construct `return` allows resuming execution at the point where the function was called;
- 1923 • The constructs `break` and `continue` restore the continuation after and before the loop,  
1924 respectively. . .

---

<sup>135</sup><https://en.wikipedia.org/wiki/Continuation>

<sup>136</sup>In fact, computing with continuations is taught in elementary school, when students learn to evaluate expressions “step by step.”




Figure 18: The expression is split into a focus (red) and a continuation (blue); when a step is performed in the focus, they are combined into a new expression, and the process repeats.

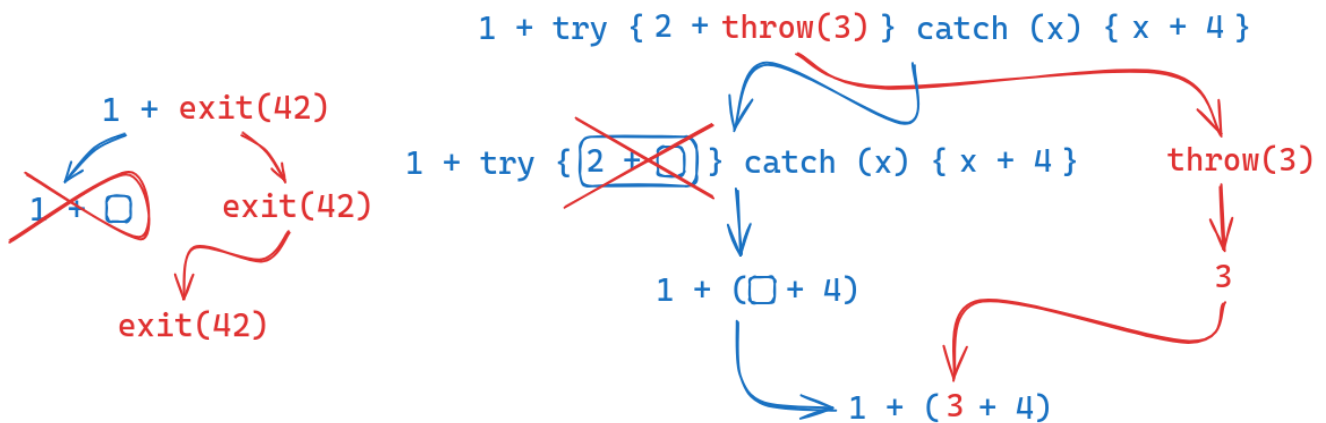


Figure 19: The construct `exit` discards the continuation and the program halts, while `throw` discards only part of the continuation up to the nearest `try-catch`.

## 1925 6.1.1 Reduction semantics

1926 The style of specifying operational semantics through the description of searching for primitive  
1927 (“**head**”) redexes and their reductions is called **reduction semantics (reduction context semantics)**<sup>137</sup>.  
1928 As usual, following Hutton’s Razor, let us consider the interpretation of a small simple language  
1929 with subtraction (an asymmetric operation for testing the implementation):

```
1 data Expr = Const Int | Diff Expr Expr
```

1930 First, we need to define the syntax of continuations, “expressions with a hole.” That is, from a  
1931 technical standpoint, a continuation is a data structure containing all the information necessary to  
1932 continue execution (here — only the remainder of the expression):

```
1 data K  
2   = Hole          -- hole □  
3   | LDiff K Expr  -- focus went into the left subexpression, remembered the right  
4   | RDiff Int K   -- computed the left, focus went right
```

1933 Let us define the operation of splitting an expression into a focus and a continuation:

```
1 split :: Expr -> (Expr, K)  
2 split e = case e of  
3   Const _ -> (e, Hole)  
4   Diff (Const _) (Const _) -> (e, Hole) -- primitive redex  
5   Diff (Const l) r ->                    -- left subexpression already computed  
6     let (focus, k) = split r in          -- search for redex in the right  
7     (focus, RDiff l k)  
8   Diff l r ->                             -- nothing computed yet  
9     let (focus, k) = split l in         -- start by searching for redex on the left  
10    (focus, LDiff k r)
```

```
11 ghci> split (Diff (Diff (Const 1) (Const 2)) (Const 3)) -- (1 - 2) - 3  
12 (Diff (Const 1) (Const 2), LDiff Hole (Const 3)) -- (1 - 2, □ - 3)
```

1934 The primitive reduction step only knows how to subtract numbers:

```
1 headReduction :: Expr -> Expr  
2 headReduction = \case  
3   Diff (Const l) (Const r) -> Const (l - r)  
4   e -> e
```

1935 After a primitive reduction step, we need to plug the result back into the continuation, so that  
1936 from an “expression with a hole” we obtain a full expression that can be further reduced:

---

<sup>137</sup>[https://en.wikipedia.org/wiki/Operational\\_semantics](https://en.wikipedia.org/wiki/Operational_semantics)

```

1 plugIn :: Expr -> K -> Expr
2 plugIn e k = case k of
3   Hole -> e
4   LDiff k' r -> Diff (e `plugIn` k') r
5   RDiff l k' -> Diff (Const l) (e `plugIn` k')

```

1937 As a rule, plugging a term  $t$  into a continuation  $E$  is denoted as  $E[t]$ .

1938 Now we can define a full step:

```

1 transition :: Expr -> Expr
2 transition e =
3   let (focus, k) = split e in    -- split into focus and context
4   headReduction focus `plugIn` k -- perform computation in the focus and plug in

```

1939 Then the operational semantics is the unfolding of the list of intermediate expressions:

```

1 eval :: Expr -> [Expr]
2 eval = List.unfoldr \prev ->
3   let next = transition prev in
4   if prev == next then Nothing else Just (next, next)

```

1940 One can notice that `split` and `plugIn` are uniquely determined by the syntax of continuations  
1941 and the form of head reductions, while `transition` and `eval` are the same for any language.  
1942 Therefore, reduction semantics is usually specified as the syntax of continuations, a list of head  
1943 reductions, and a single inference rule — a step in context:<sup>138</sup>

Values	$v ::= \mathbb{Z}$
Terms	$t ::= v \mid t \ominus t$
Evaluation context	$E ::= \square \mid E \ominus t \mid \mathbb{Z} \ominus E$
(diff)	$v_1 \ominus v_2 \longrightarrow v_1 - v_2$
	$\frac{t \longrightarrow t'}{E[t] \longrightarrow E[t']}$
step	

1944 There exists a standard tool PLT Redex<sup>139</sup> for describing and testing semantics in the reduction  
1945 context style.

## 1946 6.1.2 Continuation semantics

1947 Let us write the denotational semantics of our language:

```

1 evalDirect :: Expr -> Int
2 evalDirect = \case
3   Const n -> n
4   Diff l r -> evalDirect l - evalDirect r

```

<sup>138</sup> $\ominus$  here denotes syntactic minus (a node in the tree).

<sup>139</sup><https://redex.racket-lang.org/>

1948 Let us rewrite the denotational semantics in a style with explicit continuations. But first, let us  
1949 work with the type `K` and simplify handling it:

```
1 data K = Hole | LDiff K Expr | RDiff Int K
2 -- K = 1 + (K × Expr) + (Int × K) = 1 + (Expr + Int) × K
3 data Frame = LDiff Expr | RDiff Int
4 type K = [Frame]
```

1950 We have obtained a representation of the continuation as a stack of frames.

1951 We will write the denotational semantics for the domain `K -> Int`. Such a variant of denotational  
1952 semantics with an explicit representation of the meaning of the rest of the program (in the form  
1953 of `K`) is sometimes called **continuation semantics**. And the corresponding programming style with  
1954 passing continuations is called **continuation passing style (CPS)**.

1955 In our case, the continuation semantics will look like a pair of mutually recursive functions:

```
1 evalK :: Expr -> K -> Int
2 evalK e k = case e of
3   Const n -> k `appK` n           -- execute the rest of the program
4   Diff l r -> evalK l (LDiff r : k) -- remember to compute the right next

5 appK :: K -> Int -> Int
6 appK k result = case k of
7   [] -> result                    -- nothing more to do
8   LDiff r : k' -> evalK r (RDiff result : k') -- go compute the right
9   RDiff l' : k' -> k' `appK` (l' - result)  -- continue on the result
```

1956 Here we again descend through the expression searching for a primitive redex (here — a constant),  
1957 while remembering what needs to be done afterwards.

1958 The first recursive implementation `evalDirect` did not worry about continuations. However,  
1959 continuations are an integral part of the computation process: the evaluator needs to know what  
1960 to do next at every moment. In fact, `evalDirect` merely delegates the management of the object  
1961 language's continuations to the meta-language, but how? Notice that the resulting implementation  
1962 `evalK` is tail-recursive, meaning it can be compiled into a loop that does not consume the call  
1963 stack, while `K` is the stack. Thus, in the first case we made recursive calls and the continuation  
1964 was allocated on the (hardware) stack of the meta-language, while in the second case we manually  
1965 allocate the stack on the heap.

1966 In the case where the meta-language's stack is implemented on top of the hardware stack, there  
1967 is a risk of a stack overflow error. To avoid this, the **trampolining** technique is used, which consists  
1968 precisely in manually allocating the continuation on the heap Ganz et al. [1999], Bjarnarson [2012].

1969 Just as the type of a zipper's context can be computed as the derivative of the algebraic  
1970 representation of the corresponding type Huet [1997], McBride [2001], Abbott et al. [2003], one  
1971 can compute the type of a fold's continuation McBride [2008].

### 1972 6.1.3 First-class continuations

1973 In the examples above, language constructs manage continuations implicitly (see Fig. 19).  
1974 However, sometimes operators are introduced that allow explicitly manipulating continuations. With  
1975 their help one can implement both control flow manipulation capabilities such as generators and  
1976 coroutines<sup>140141</sup>, as well as all other effects such as state (see 7).

1977 **First-class continuations** are continuations that can be represented as values in the program.  
1978 Given that a continuation has a vacant slot for a not-yet-computed subexpression, first-class  
1979 continuations are represented as first-class functions.

1980 To obtain a first-class continuation in code, one needs to either write the code in CPS, or use a  
1981 built-in language operator, of which a great many have been invented [Hillerström, 2022, appendix  
1982 A]. For example, *J*, *escape* Reynolds [1972], *call/cc*...

1983 As an example, let us implement in our language an operation *Cont* that allows capturing the  
1984 current continuation. It will take a user function and pass to it the current continuation from itself  
1985 to the end of the program:

$$E[\text{cont } f] \longrightarrow f (\lambda x. E[x])$$

1986 First, let us extend the language with lambda calculus:

```
1 data Expr
2   = Const Int | Diff Expr Expr
3   | Var String | Lam String Expr | App Expr Expr
4
4 data Frame
5   = LDiff Expr | RDiff Value
6   | LApp Expr | RApp Value
7   | SetEnv Env -- after executing a closure, restore the current environment
```

1987 The interpreter implementation is fairly straightforward. The only thing is that the remainder  
1988 of the program is now represented as a linear list that is effectively traversed sequentially in a loop.  
1989 We must not forget to restore the original environment after executing a closure body for the  
1990 continuation to proceed:<sup>142</sup>

```
1 evalK :: (?env :: Env) => Expr -> K -> Value
2 evalK e k = case e of
3   Const n -> k `appK` Number n
4   Diff l r -> evalK l (LDiff r : k)
5   Var name -> k `appK` (?env ! name)
6   Lam name body -> k `appK` Closure name ?env body
7   App f arg -> evalK f (LApp arg : k)
```

<sup>140</sup>[https://en.wikibooks.org/wiki/Haskell/Continuation\\_passing\\_style#Example:\\_coroutines](https://en.wikibooks.org/wiki/Haskell/Continuation_passing_style#Example:_coroutines)

<sup>141</sup><https://kotlinlang.org/api/core/kotlin-stdlib/kotlin.coroutines/suspend-coroutine.html>

<sup>142</sup>For convenience, we will pass the environment as an implicit parameter.

```

8 appK :: (?env :: Env) => K -> Value -> Value
9 appK k result = case k of
10   [] -> result
11   LDiff r : k' -> evalK r (RDiff result : k')
12   RDiff l' : k' -> k' `appK` Number (unwrapNumber l' - unwrapNumber result)
13   LApp arg : k' -> evalK arg (RApp result : k')
14   RApp f : k' -> case f of
15     Closure name env body ->
16       let currEnv = ?env in
17       let ?env = Map.insert name result env in
18       evalK body (SetEnv currEnv : k')
19     K env k'' ->
20       let currEnv = ?env in
21       let ?env = env in
22       (k'' ++ SetEnv currEnv : k') `appK` result
23   other -> error $ "Expected callable, got " <> show other
24   SetEnv env : k' -> let ?env = env in k' `appK` result
25   where
26     unwrapNumber = \case
27       Number n -> n
28       other -> error $ "Expected number, got " <> show other

```

1991 Finally, let us implement the cont operator:

```

1 data Expr = ... | Cont Expr
2 data Frame = ... | ContFrame

3 evalK e k = case e of
4   ...
5   Cont f -> evalK f (ContFrame : k)

6 appK k result = case k of
7   ...
8   ContFrame : k' -> [RApp result] `appK` K ?env k' -- reuse the RApp branch

```

1992 Let us consider several examples (denoting `contLam name body = Cont (Lam name body)`,  
1993 `(-.) = Diff`, `c = Const`, `v = Var` and `(@) = App`):

- 1994 ● `c 10 -.` `contLam "k" (c 1)` — discard the continuation and return 1 ( $k = 10 - \square$ );
- 1995 ● `c 10 -.` `contLam "k" (v "k" @ c 1)` —  $\rightsquigarrow k 1 \rightsquigarrow 10 - 1 \rightsquigarrow 9$ ;
- 1996 ● `c 10 -.` `contLam "k" (v "k" @ c 1 - . c 2)` —  $\rightsquigarrow k 1 - 2 \rightsquigarrow 9 - 2 \rightsquigarrow 7$ ;
- 1997 ● `c 10 -.` `contLam "k" (v "k" @ c 1 - . v "k" @ c 2)` —  $\rightsquigarrow k 1 - k 2 \rightsquigarrow 9 - 8 \rightsquigarrow 1$ .

1998 Thus, we have obtained a language that allows the user to obtain the current continuation as  
1999 a function at any point in the program.

2000 The ability to obtain first-class continuations is almost never provided by industrial languages,  
 2001 as it is a rather dangerous tool. Indeed, if the user does not invoke the continuation, resources may  
 2002 not be released. If they invoke it multiple times, there may again be incorrect resource handling or  
 2003 corruption of mutable memory (continuations that can be safely invoked multiple times are called  
 2004 **multi-shot**). Such behaviors can be excluded by special handling of these situations Muhcu et al.  
 2005 [2025], compile-time checks (e.g., using linear types) or runtime checks. Currently, with active  
 2006 research on integrating effect handlers (see 7), first-class continuations may get a new chance.

## 2007 6.2 Continuations by hand

2008 Let us consider what CPS looks like in ordinary code, rather than in the context of deep  
 2009 embedding. For this, we will first switch to shallow embedding and directly construct elements of  
 2010 the target domain. Second, we will refunctionalize the continuations **K** and generalize: instead of  
 2011 **K -> Int** we will use `forall r . (a -> r) -> r`, where `a -> r` is the functional representation  
 2012 of the continuation.

2013 Such CPS exploits the following isomorphism:

```

1  to :: a -> (forall r . (a -> r) -> r)
2  to x k = k x

3  from :: (forall r . (a -> r) -> r) -> a
4  from comp = comp id
  
```

2014 In other words, instead of providing a value of type `a`, one can ask the caller how they intend to  
 2015 use this value `a -> r`, do it oneself, and return `r` to the caller.

2016 Practitioners know this isomorphism through the technique of using callbacks. Theoreticians  
 2017 know that it is a special case of the Yoneda lemma Hinze and James [2010]. Compare CPS with  
 2018 the instance of this lemma in preorder theory (above we took `a = Unit`; `b = a`):

$$a \rightarrow b \cong \forall r. (b \rightarrow r) \rightarrow (a \rightarrow r)$$

$$a \leq b \iff \forall r. (b \leq r) \Rightarrow (a \leq r)$$

2019 For example, we can rewrite factorial in CPS. Note that the code has access to a first-class  
 2020 continuation (however, it does not yet use it in any non-trivial way), and is again tail-recursive.

```

1  facCps :: Int -> (forall r . (Int -> r) -> r)
2  facCps n k
3    | n <= 1 = k 1
4    | otherwise = facCps (n - 1) \res -> k (n * res)

5  facCps 3 id ~> facCps 2 \res -> id (3 * res)
6  ~> facCps 1 \res -> id (3 * (2 * res)) ~> id (3 * (2 * 1)) ~> 6
  
```

2021 **Exercise 45** How much stack memory does the function `facCps` consume?

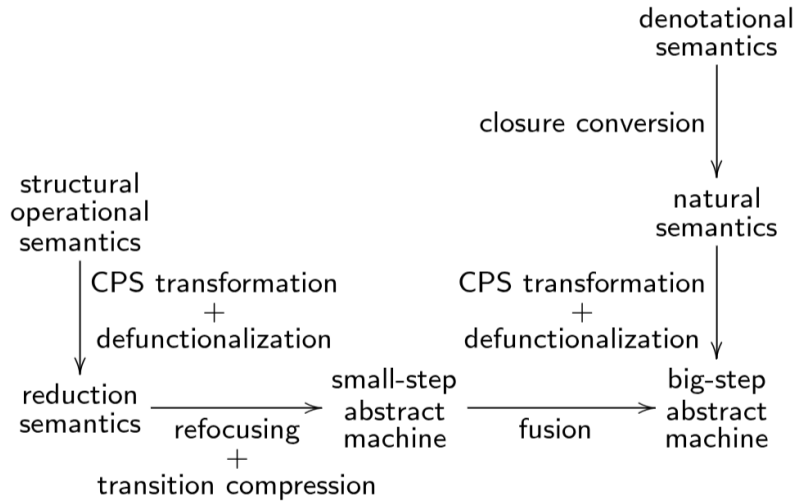


Figure 20: Relationship between semantics in various styles Danvy [2008].

## 2022 6.2.1 Defunctionalization and accumulators

2023 Let us defunctionalize the continuations in `facCps`. We use two higher-order functions, `id` and  
 2024 `\res -> k (n * res)`, which give us two constructors:

```

1  data K = Id | Times Int K
2  -- ≅
3  type K = [Int]

4  runK :: K -> Int
5  runK = product

6  facCps :: Int -> K -> Int
7  facCps n k
8    | n <= 1 = runK k
9    | otherwise = facCps (n - 1) (snoc k n)

10 facCps 3 [] ~> facCps 2 [3] ~> facCps 1 [3, 2] ~> runK [3, 2] ~> 3 * (2 * 1)
  
```

2025 We have again obtained the representation of the continuation as a stack of frames. This means  
 2026 that the interpreter `evalK` we obtained earlier (6.1.2) is a CPS version of a regular interpreter, but  
 2027 with defunctionalized continuations (which are convenient for debugging — they can be printed,  
 2028 unlike Haskell functions). If we extend our earlier language and write factorial in it, then perform  
 2029 fusion or refunctionalization (4.4.1) of the program tree and the continuation tree, we obtain  
 2030 the current factorial implementation. And in general, correspondences can be established between  
 2031 various styles of semantics implementations (see Fig. 20).

2032 This technique can be used in more complex cases as well, to easily obtain tail-recursive

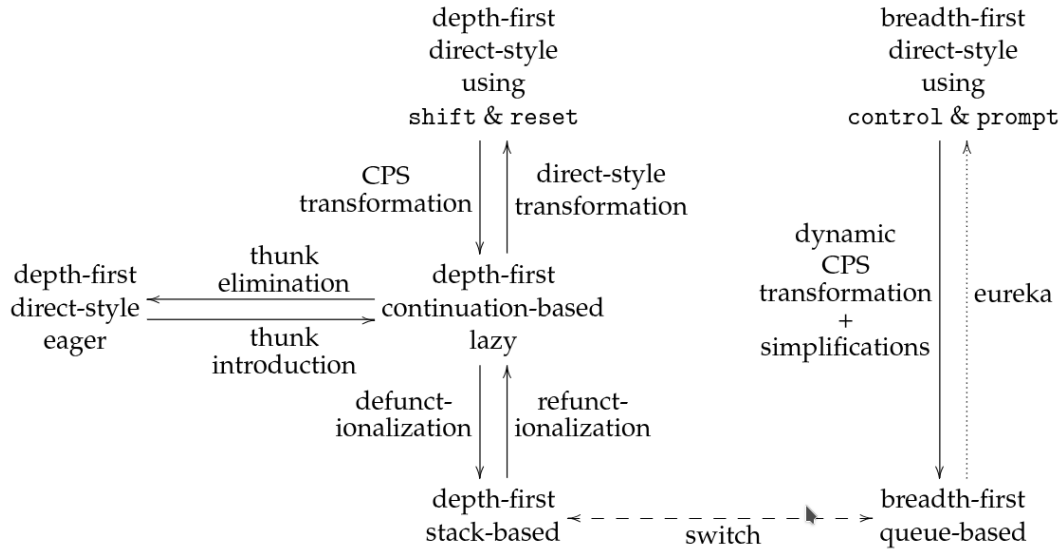


Figure 21: Code transformations based on CPS and de(re)functionalization Danvy [2006].

2033 (iterative) implementations<sup>143</sup> Gibbons [2021]<sup>144</sup>. In general, CPS and defunctionalization are a  
 2034 rich source of various refactorings (see Fig. 21).

2035 Now let us notice that the multiplication operation is associative, so we can replace `snoc` with  
 2036 multiplication and represent the continuation as a single number. We obtain the familiar factorial  
 2037 implementation with an accumulator:

```

1 facAcc :: Int -> Int -> Int
2 facAcc n acc
3   | n <= 1 = acc
4   | otherwise = facAcc (n - 1) (acc * n)

5 facAcc 3 1 ~> facAcc 2 (1 * 3) ~> facAcc 1 ((1 * 3) * 2) ~> (1 * 3) * 2 ~> 6

```

2038 Also, one can represent frames as endomorphisms and use function composition as the associative  
 2039 operation Ploeg and Kiselyov [2014]<sup>145</sup>.

### 2040 6.2.2 Monad Cont

2041 Due to CPS, the code has lost its familiar structure where functions directly return their results  
 2042 (i.e. **direct style**). With a large number of calls to transformed functions, the code becomes hard  
 2043 to read (a problem known as **callback hell**):

```

1 fibCps :: Int -> (forall r . (Int -> r) -> r)

```

<sup>143</sup><https://www.pathsensitive.com/2019/07/the-best-refactoring-youve-never-heard.html>

<sup>144</sup>(youtube) Jeremy Gibbons - Continuation-passing style, defunctionalization, and associativity.

<sup>145</sup>[https://wiki.haskell.org/Difference\\_list](https://wiki.haskell.org/Difference_list)

```

2 fibCps n k = if n <= 2 then k 1 else
3   fibCps (n - 1) \res1 ->
4   fibCps (n - 2) \res2 ->
5   k (res1 + res2)

```

2044 The domain `(a -> r) -> r` can be made into a monad, restoring direct-style code inside `do-`  
2045 notation. Let us introduce a `newtype` wrapper for declaring instances:

```

1 newtype Cont r a = Cont { runCont :: (a -> r) -> r }

```

2046 The functor adds post-processing to the result before passing it to the continuation:

```

1 instance Functor (Cont r) where
2   -- fmap :: (a -> b) -> ((a -> r) -> r) -> ((b -> r) -> r)
3   fmap f (Cont comp) = Cont \k -> comp (k . f)

```

2047 The applicative simply passes the value to the continuation:

```

1 instance Applicative (Cont r) where
2   pure x = Cont \k -> k x
3   (<*>) = ap

```

2048 One can notice that monadic bind also takes a continuation as its second argument, but a  
2049 “small” one, up to the end of the `do`-block. Thus, the meaning of the monadic bind implementation  
2050 for `Cont` is composing the “small” continuation with the “large” continuation passed from outside:

```

1 instance Monad (Cont r) where
2   (>>=) :: Cont r a -> (a -> Cont r b) -> Cont r b
3   Cont comp >>= k = Cont \k' -> comp \x -> runCont (k x) k'

```

2051 Now we can write linear code, and the monadic machinery automatically constructs continuations  
2052 and feeds them into previous computations:

```

1 fibCont :: Int -> Cont r Int
2 fibCont n = if n <= 2 then pure 1 else do
3   res1 <- fibCont (n - 1)
4   res2 <- fibCont (n - 2)
5   pure (res1 + res2)

```

2053 **Exercise 46** *Abort the computation if res1 is greater than 50.*

2054 **Exercise 47** *Abort the computation as soon as the total result exceeds 50.*

2055 The monad `Cont` provides an implementation of an embedded language in which one can obtain  
2056 the continuation of a computation:

```

1 Cont :: ((a -> r) -> r) -> Cont r a

```

2057 Practitioners know this technique of writing CPS code via monads through the concepts of  
2058 `Future/Promise`<sup>146147</sup>.

<sup>146</sup>[https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global\\_Objects/Promise](https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/Promise)

<sup>147</sup><https://github.com/promises-aplus/promises-spec/issues/94>

### 2059 6.2.3 call/cc

2060 The most well-known classical operator used in Scheme for obtaining first-class continuations  
2061 is `call/cc` (**call with current continuation**)<sup>148149</sup>. The provided continuation is **undelimited**  
2062 (**abortive**), since it contains “the end of the program” — no code will execute after its invocation.  
2063 Undelimited continuations are, de facto, not quite functions, since they do not return a result (it is  
2064 already “beyond the grave”); consequently, they also do not compose (just as it would be strange  
2065 to compose `abort` with `exit`)<sup>150151</sup>. They are more like co-values: while a part of the program is  
2066 executing, they await its result Curien and Herbelin [2000].

2067 Let us emulate `call/cc` in the `Cont` monad. We cannot capture the entire program continuation,  
2068 since it is only assembled within `Cont`, but we can ignore the continuation of the call to the captured  
2069 continuation:

```
1 callCC :: ((a -> Cont r b) -> Cont r a) -> Cont r a
2 callCC f = Cont \k -> runCont (f \x -> Cont \_ -> k x) k

3 foo :: Int -> Cont r String
4 foo x = callCC \k -> do
5   let y = x ^ 2 + 3
6       when (y > 20) $ k "over twenty" -- throws next line away
7   pure (show $ y - 4)
```

### 2070 6.3 Delimited continuations

2071 In modern practice, continuations are typically used not up to the end of the program, but only up  
2072 to a certain point. They are called **delimited** or **composable continuations**, **subcontinuations**<sup>152153</sup>.

2073 Constructs for working with such continuations come in pairs: an operator is introduced that  
2074 delimits the current continuation (possibly with a label); as well as an operator for capturing a  
2075 fragment of the current continuation (possibly up to a specific delimiter with a given label). Many  
2076 such operators have been invented [Hillerström, 2022, appendix A], but they are more or less all  
2077 reducible to each other.

2078 For example, working with exceptions involves two constructs: the continuation is delimited by  
2079 `try-catch`, and `throw` discards the corresponding partial continuation without capturing it (see  
2080 Fig. 19):

$$E_1[\text{try}\{E_2[\text{throw}(v)]\}\text{catch}(x)\{t\}] \rightarrow E_1[[x \mapsto v] t]$$

2081 We will consider the universal operators from Dyvbig et al. [2007]. The paper introduces the

<sup>148</sup><https://en.wikipedia.org/wiki/Call-with-current-continuation>

<sup>149</sup>[https://en.wikibooks.org/wiki/Haskell/Continuation\\_passing\\_style#callCC](https://en.wikibooks.org/wiki/Haskell/Continuation_passing_style#callCC)

<sup>150</sup><https://okmij.org/ftp/continuations/undelimited.html>

<sup>151</sup><https://okmij.org/ftp/continuations/against-callcc.html>

<sup>152</sup><https://www.cl.cam.ac.uk/teaching/2324/R277/handout-delimited-continuations.pdf>

<sup>153</sup>(youtube) Keynote: Delimited Continuations, Demystified by Alexis King | Lambda Days 2023.

(Variables)  $x, y, \dots$   
 (Expressions)  $e ::= x \mid \lambda x. e \mid e e$   
 $\quad \quad \quad \mid \text{newPrompt} \mid \text{pushPrompt } e e$   
 $\quad \quad \quad \mid \text{withSubCont } e e \mid \text{pushSubCont } e e$

Figure 22: Syntax of  $\lambda$ -calculus with primitives for working with continuations.

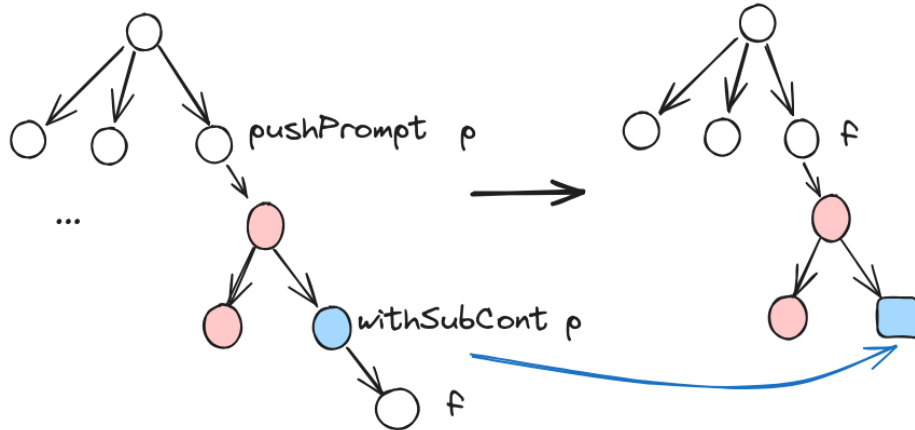


Figure 23: Example of `withSubCont` in action.

2082 following set of syntactic constructs (Fig. 22) for working with delimited continuations, in addition  
 2083 to pure call-by-value lambda calculus:

- 2084 • `newPrompt`<sup>154</sup> — creates a fresh identifier (label) for a delimiter;
- 2085 • `pushPrompt p e` — installs a delimiter with label  $p$  and executes expression  $e$ ;
- 2086 • `withSubCont p f` — captures the partial continuation up to the delimiter with label  $p$  and  
 2087 passes it to function  $f$ , returns the result of  $f$  (Fig. 23);
- 2088 • `pushSubCont k v` — executes the composition of the current continuation and  $k$  on value  
 2089  $v$ .

2090 Delimited continuation operators can be easily understood as **resumable exceptions**, i.e. exceptions  
 2091 that can be caught and the program resumed from the point where the exception was thrown (with  
 2092 some value). What is unusual is that the classical delimited continuation operators accept code  
 2093 that works with the partial continuation at the “throw site,” rather than at the “catch site.” That  
 2094 is, the handling block is written not with `catch`, but with `throw`:

$$E_1[\text{pushPrompt } p \{E_2[\text{withSubCont } p f]\}] \rightarrow E_1[f E_2]$$

$$E_1[\text{try}\{E_2[\text{throw } v]\}\text{catch}(x, k)\{t\}] \rightarrow E_1[[x \mapsto v] [k \mapsto \lambda y. E_2[y]] t]$$

<sup>154</sup>Historically in Lisps, undelimited continuations were delimited only by the REPL, hence the name for delimiters — “prompt.”

```

    (λp. 2 + pushPrompt p
      if (withSubCont p
        (λk. (pushSubCont k False) + (pushSubCont k True)))
      then 3
      else 4)
    newPrompt

```

Figure 24: Example expression (result is 9).

2095 **Exercise 48** Reduce the example in Fig. 24.

### 2096 6.3.1 Implementing the operators

2097 Let us extend the language from 6.1.3 with the operators for working with delimited continuations  
 2098 discussed above. For simplicity, instead of fresh prompt labels we will use names, and instead of  
 2099 pushSubCont ordinary application:

```

1 data Expr = ... | PushPrompt String Expr | WithSubCont String Expr
2 data Frame = ... | PushPromptFrame String | WithSubContFrame String

```

2100 The redex search is extended in the obvious way:

```

1 evalK :: (?env :: Env) => Expr -> K -> Value
2 evalK e k = case e of
3   -- ...
4   PushPrompt promptName body -> evalK body (PushPromptFrame promptName : k)
5   WithSubCont promptName f -> evalK f (WithSubContFrame promptName : k)

```

2101 The entire job of PushPrompt is to get stuck in the continuation as PushPromptFrame, but  
 2102 when execution reaches it, it is simply ignored:

```

1 appK :: (?env :: Env) => K -> Value -> Value
2 appK k result = case k of
3   -- ...
4   PushPromptFrame _ : k' -> k' `appK` result

```

2103 To capture a fragment of the continuation, we need an operation that finds the first prompt  
 2104 with the matching name in the continuation and returns the **subcontinuation** — up to that prompt,  
 2105 and the **metacontinuation** — after it:

```

1 splitByPrompt :: K -> String -> (K, K)
2 splitByPrompt k targetPromptName = go [] k
3   where
4     go _ [] = error $ "Prompt " ++ targetPromptName ++ " not found"

```

```

5     go subcont (PushPromptFrame promptName : metacont)
6     | promptName == targetPromptName = (subcont, metacont)
7     go subcont (frame : metacont) = go (subcont ++ [frame]) metacont

```

2106 Now WithSubCont must simply apply the user function to the subcontinuation, while not  
2107 forgetting to execute the metacontinuation afterwards:

```

1  appK :: (?env :: Env) => K -> Value -> Value
2  appK k result = case k of
3      -- ...
4      WithSubContFrame promptName : k' ->
5          let (subcont, metacont) = k' `splitByPrompt` promptName in
6              (RApp result : metacont) `appK` K ?env subcont

```

2108 **Exercise 49** Reduce the following example:

```

1  exampleDelimited = let ?env = Map.empty in flip evalK [] $
2      c 10 -. PushPrompt "p"
3          (c 5 -. WithSubCont "p"
4              (lam "k" $ v "k" @ c 1 -. v "k" @ c 3))

```

2109 Many classical operators, such as shift/reset, prompt/control, etc., can be obtained by  
2110 leaving the prompt frame in the subcontinuation or in the metacontinuation in the implementation  
2111 of splitByPrompt Dyvbig et al. [2007].

### 2112 6.3.2 In Monad Cont

2113 As an example, let us implement two classical operators for working with delimited continuations  
2114 — **shift-reset**. We will use the continuations assembled by the Cont monad. shift simply captures  
2115 the current continuation and passes it to the user computation similarly to cont:

```

1  shift :: ((a -> r) -> Cont r r) -> Cont r a
2  shift f = Cont \k -> runCont (f k) id

```

2116 reset passes id as the continuation to the argument computation, so that computation does not  
2117 have access to the continuation after reset (while the Haskell continuation remembers that after  
2118 comp it needs to execute k):

```

1  reset :: Cont a a -> Cont r a
2  reset comp = Cont \k -> k (runCont comp id)

```

2119 **Exercise 50** What will be the result of executing the following function:

```

1  exampleShiftReset = flip runCont id $
2      (1 +) <$> reset ((2 +) <$> shift \k -> pure (k 3 + k 5))

```

## 2120 6.4 Applications of continuations

2121 Continuations are useful for understanding the meaning of programs we write every day (see  
2122 107). Furthermore, as we will see in this chapter, first-class continuations can be used as a means  
2123 of building powerful embedded languages.

### 2124 6.4.1 Everything through continuations

2125 Let us implement a number of useful embedded languages using first-class continuations. To  
2126 gain access to continuations, we will use the shallow embedded language `Cont`<sup>155</sup>. Thus, we will  
2127 work with the tower of languages Haskell, `Cont`, `X`, where `X` is the embedded language under  
2128 consideration.

2129 Let us start with a trivial embedded imperative language:

```
1 runIdentityC :: (forall r . Cont r a) -> a
2 runIdentityC comp = runCont comp id

3 exampleIdentity :: Int
4 exampleIdentity = runIdentityC do
5   x <- pure 4
6   y <- pure 5
7   pure (x + y)
```

2130 Let us implement a language with exceptions. For this we modify the first type parameter of  
2131 `Cont`, the **answer (response) type**<sup>156</sup>. The function `abort` ignores the program continuation up  
2132 to `runExn`, which concludes the continuation by wrapping the successful result in `Just`:

```
1 abort :: Cont (Maybe r) a
2 abort = Cont $ const Nothing

3 runExn :: (forall r . Cont (Maybe r) a) -> Maybe a
4 runExn comp = runCont comp Just

5 exampleExn :: Int -> Maybe Int
6 exampleExn n = runExn do
7   when (n < 0) abort
8   pure (n + 1)
```

2133 Notice that the semantics of the `abort` operation is entirely concentrated in the function `abort`,  
2134 whereas in `Monad Maybe` the corresponding function would only be responsible for constructing a  
2135 domain element, while monadic `bind` would handle the continuations. Having access to first-class

---

<sup>155</sup><https://blog.poisson.chat/posts/2019-10-26-reasonable-continuations.html>

<sup>156</sup>[https://wiki.haskell.org/Cont\\_computations\\_as\\_question-answering\\_boxes](https://wiki.haskell.org/Cont_computations_as_question-answering_boxes)

2136 continuations, we no longer need to implement monadic bind manually. It is said that for arbitrary  
2137 `m`, `Cont (m r) a` is a “free” implementation of a monad<sup>157</sup>.

2138 Error handling through returning a special result (`Either` in Haskell, `Result` in Rust, `nil`  
2139 in Go...) has monadic semantics (conditional branches at each point discard only part of the  
2140 continuation), while a full-fledged exception mechanism has access to the entire delimited continuation  
2141 up to the corresponding `try-catch` block, which resembles our current implementation. From the  
2142 perspective of syntactic noise at the source code level, the exception variant is clearly preferable, and  
2143 we will consider the question of typing exceptions later (see 8). From a performance perspective,  
2144 things are less obvious<sup>158159</sup>, though stack unwinding by the runtime seemingly has more room for  
2145 optimization.

2146 Now let us try invoking the continuation multiple times. We obtain a language with nondeterminism  
2147 (or backtracking). Here, choice continues the rest of the program twice and accumulates all  
2148 obtained results in a list:

```
1 choice :: Cont [r] Bool
2 choice = Cont \k -> k True ++ k False

3 runNondet :: (forall r . Cont [r] a) -> [a]
4 runNondet comp = runCont comp (:[ ])

5 exampleNondet :: Int -> [Int]
6 exampleNondet n = runNondet do
7   b <- choice
8   if b then pure n else pure (n - 1)
```

2149 Remarkably, having only first-class continuations one can even implement a language with a  
2150 mutable memory cell:

```
1 -- s -> (s, a) ~ ((s, a) -> r) -> (s -> r)
2 -- ~ (a -> s -> r) -> (s -> r) ~ Cont (s -> r) a

3 get :: Cont (s -> r) s -- (s -> s -> r) -> (s -> r)
4 get = Cont \k s -> k s s

5 put :: s -> Cont (s -> r) ()
6 put s' = Cont \k _s -> k () s'

7 runStateC :: (forall r . Cont (s -> r) a) -> s -> (a, s)
8 runStateC comp = runCont comp (,)
```

---

<sup>157</sup><https://hackage.haskell.org/package/kan-extensions-5.2/docs/Control-Monad-Codensity.html>

<sup>158</sup><https://www.serpentine.com/2011/02/25/cps-is-great-cps-is-terrible/>

<sup>159</sup><https://stackoverflow.com/questions/13835817/are-exceptions-in-c-really-slow>

```

9  exampleState :: Int -> ((), Int)
10 exampleState = runStateC do
11     s <- get
12     put (s + 1)

```

2151 The idea is quite simple: capturing the continuation leaves behind a function into which the previous  
2152 continuation call supplies the current state. And indeed, mutable state is simply a fragment of the  
2153 continuation (e.g., the hardware stack) to which we have direct read and write access. Here we have  
2154 placed an application at the beginning of the continuation, to which, by capturing the continuation,  
2155 we jump to obtain the argument and update it to a new one.

2156 **Exercise 51** *Reduce the example above to understand how it works.*

2157 Intuitively, the power of continuations can be explained as follows: maintaining the continuation  
2158 is the primary responsibility of the language’s runtime system; when continuations are handed to  
2159 the user, their code becomes involved in the runtime’s activity and effectively becomes part of it.

## 2160 6.4.2 The mother of all monads

2161 In the previous section (6.4.1) we implemented the capabilities of classical monads in a language  
2162 with continuations. Does this work for an arbitrary monad? It turns out that yes<sup>160161</sup> Filinski [1994].

2163 Recall that a monad is a type constructor  $m$  with which we represent the denotation of a  
2164 computation. The operations `pure` and `bind` allow creating and composing denotations. Thus, we  
2165 have computations represented as first-class data.

2166 If our language provides first-class continuations (let us take `Cont (m r)` as such a language for  
2167 example), we can define two operations: `reflect` — execute a monadic computation in the language  
2168 with continuations and obtain the result  $a$ ; `reify` — obtain a denotation from a computation in  
2169 the language with continuations<sup>162</sup>. In Haskell, this can be expressed as overloading this pair of  
2170 functions for various monads:

```

1  class MonadicReflection m where
2     -- m a -> (a -> m r) -> m r
3     reflect :: m a -> forall r . Cont (m r) a
4     -- (forall r . (a -> m r) -> m r) -> m a
5     reify  :: (forall r . Cont (m r) a) -> m a

```

2171 For example, for the `State` domain the implementation looks as follows:

```

1  newtype State s a = State { runState :: s -> (a, s) }

```

<sup>160</sup><http://www.schoolofhaskell.com/school/to-infinity-and-beyond/pick-of-the-week/the-mother-of-all-monads>

<sup>161</sup><https://blog.poisson.chat/posts/2019-10-26-reasonable-continuations.html>

<sup>162</sup>For a strict language, `reify` should accept a thunk  $(() \rightarrow a) \rightarrow m a$ .

```

2 instance MonadicReflection (State s) where
3   reflect :: State s a -> (forall r . Cont (State s r) a)
4   reflect comp = Cont \k -> State \s ->
5     let (a, s') = runState comp s in
6       runState (k a) s'

7   reify :: (forall r . Cont (State s r) a) -> State s a
8   reify comp = State $ runState (runCont comp (\x -> State (x,)))

9 exampleStateReflection :: Int -> (Int, Int)
10 exampleStateReflection = runState $ reify do
11   x <- reflect get
12   reflect $ put (x + 1)
13   pure x

```

2172 Intuitively, this can be understood as follows: the bind implementation does not actually care  
2173 whether it is given a “small” continuation (assembled via desugaring of `do`-notation) or a “large”  
2174 continuation assembled by the meta-language or the `Cont` monad. And in principle, monads can be  
2175 defined through their embedding into the `Cont` monad<sup>163</sup>.

2176 Thus, if the language supports first-class continuations, one can implement the functionality of  
2177 an arbitrary monad in direct style<sup>164</sup>, without `do`-notation, applicative chains, and other syntactic  
2178 noise<sup>165166</sup>. Later we will see that this approach, unlike monads, natively supports composition (7),  
2179 and also does not impose restrictions on typing (8). This raises the question: do we still need  
2180 monads in programming?

### 2181 6.4.3 Generators and coroutines

2182 Generators and coroutines are usually defined as computations that can be suspended and  
2183 resumed in the same state Moura and Ierusalimsky [2009]. Coroutines generalize generators and  
2184 are typically used as a primitive for asynchronous programming in the form of `async/await`  
2185 or other language constructs Elizarov et al. [2021]. The difference between synchronous and  
2186 asynchronous programming can be understood as follows: in the former case, continuations are  
2187 managed exclusively by the operating system, while in the latter — by the language facilities.

2188 As an example, let us implement generators. As the result of a generator we will use a lazy list,  
2189 an iterator encoded as an `unfold` (see 4.5):

```

1 data Box f = forall s . Box s (s -> f s)
2 data ListF a rec = Nil | Cons a rec
3 type Iterator a = Box (ListF a)

```

<sup>163</sup><https://blog.poisson.chat/posts/2019-10-27-continuation-submonads.html>

<sup>164</sup><https://www.unison-lang.org/docs/fundamentals/abilities/for-monadically-inclined/>

<sup>165</sup><http://www.valuedlessons.com/2008/01/monads-in-python-with-nice-syntax.html>

<sup>166</sup><https://github.com/lampepfl/monadic-reflection/blob/main/TUTORIAL.md>

```

4 box2list :: Iterator a -> [a]
5 box2list (Box s next) = case next s of
6   Nil -> []
7   Cons x s' -> x : box2list (Box s' next)

```

2190 The hidden state of the iterator will be the not-yet-executed continuation. We will obtain access  
2191 to the continuation via the `Cont` monad, where we take `GenState` as the response type. Thus, the  
2192 `yield` operation suspends the computation and saves the continuation in the `Yield` constructor,  
2193 while `makeGen` finishes the given generator by producing the `Stop` constructor:

```

1 data GenState a = Stop | Yield a (() -> GenState a)

2 makeGen :: Cont (GenState a) () -> Iterator a
3 makeGen comp = Box
4   (\() -> runCont comp (\() -> Stop))
5   (\k -> case k () of
6     Stop -> Nil
7     Yield x k' -> Cons x k'
8   )

9 yield :: a -> Cont (GenState a) ()
10 yield x = cont \k -> Yield x k

```

2194 Now we can write code that sequentially produces multiple results:

```

1 exampleGen :: [Int]
2 exampleGen = box2list $ makeGen do
3   yield 1
4   yield 2
5   yield 3

```

2195 **Exercise 52** Reduce `runCont exampleGen (\() -> Stop)` to normal form.

## 2196 6.5 Efficient work with continuations

2197 In practice, implementing generators and coroutines requires suspending and resuming programs,  
2198 in other words, capturing continuations. This needs to be done as efficiently as possible.

### 2199 6.5.1 Contiguous stack

2200 The continuation is represented as a hardware stack. As soon as a continuation needs to be  
2201 captured, the stack is copied to the heap (possibly lazily)<sup>167</sup>. That is, this approach relies entirely  
2202 on support from the language runtime.

<sup>167</sup>(youtube) Ivan Uglyansky - Java Project Loom.

## 2203 6.5.2 Segmented stack

2204 The call stack is represented as a linked list of heap-allocated segments, each corresponding to  
2205 a delimiting operation. Thus, no copying is required; it suffices to swap pointers. However, in this  
2206 case the stack is non-local, which is not great for cache performance.

## 2207 6.5.3 Finite state machine (FSM)

2208 This implementation involves automatic CPS transformation of user code by the compiler. To  
2209 avoid allocating a large number of closures, continuations are defunctionalized and within each  
2210 function are represented by a single mutable object. Thus, the function's entire state is allocated  
2211 on the heap once, and resuming the function body at a particular point is implemented as a state  
2212 machine — using labels and jumps<sup>168</sup>.

2213 This is how, for example, coroutines are implemented in Kotlin<sup>168</sup>, generators in C#<sup>169</sup>...

2214 Even in this form, CPS remains a heavyweight transformation that can slow down code execution  
2215 by orders of magnitude. The issue is, in particular, that variables in this approach are difficult to  
2216 place in registers (functions have many entry and exit points<sup>168</sup>), and they constantly have to be  
2217 written to RAM — performing **spilling**<sup>170</sup>.

---

<sup>168</sup><https://github.com/Kotlin/KEEP/blob/master/proposals/coroutines.md#state-machines>

<sup>169</sup><https://csharpindepth.com/Articles/IteratorBlockImplementation>

<sup>170</sup>[https://en.wikipedia.org/wiki/Register\\_allocation](https://en.wikipedia.org/wiki/Register_allocation)

## 2218 7 Effects and Modular Interpreters

2219 Earlier we recognized managing complexity as the programmer’s main task, and constructing  
2220 embedded languages as the primary tool for solving it (5.1). In this chapter we consider the notion  
2221 of an effect. It is closely related to embedded languages and will give us a better understanding of  
2222 when to construct them, what this gives us, and what one needs to be careful about.

2223 For implementing embedded languages we preferred shallow embedding in the form of tagless  
2224 final (5.4), which maximally reuses the capabilities of the meta-language and allows giving different  
2225 interpretations to the same program. Next we studied the computation process and extracted the  
2226 notion of a continuation (6). It turned out that the tagless final languages we were building around  
2227 monads can be expressed through continuations more conveniently and simply (6.4.1, 6.4.2). In  
2228 this chapter we will understand how this will help us solve the expression problem (5.5) completely.

2229 From a historical-philosophical perspective, the theory of effects is a solid bridge between the  
2230 functional and imperative programming paradigms. The two originated separately — as the Turing  
2231 machine and Church’s  $\lambda$ -calculus — and remained largely isolated schools of thought for decades.  
2232 Only in the last 30 years has this begun to change, and we have come to understand how the  
2233 two worlds complement each other. The theory of effects remains an extremely hot topic to this  
2234 day<sup>171172</sup>.

### 2235 7.1 The notion of an effect

2236 Let us begin the discussion from the opposite end, with the property of purity. A **pure function**  
2237 possesses the following properties:

- 2238 • Its result is always the same given the same set of arguments (it does not depend on anything  
2239 else in any more non-trivial way);
- 2240 • Its only observable result is its return value.

2241 In general, the programming style using pure functions is encouraged, as it possesses many good  
2242 properties. One can conveniently reason about them using equational reasoning; everything needed  
2243 to understand the code is explicitly written in that code; classical type systems work well, providing  
2244 abstraction completeness, quality documentation, and partial specification. . . It is also known that  
2245 everything can be expressed using pure computations, even working with IO Jones [2001].

2246 However, using only pure functions, everything must be done manually. In the case of IO (and  
2247 similarly with state) — passing the resulting world as an argument again and again:

```
1  getList :: Int -> World -> (World, [Int])
2  getList n w | n == 0 = (w, [])
3                | otherwise =
4    let (w', x) = getInt w in
5    let (w'', xs) = getList (n - 1) w' in
```

---

<sup>171</sup>(youtube) The Evolution of Effects — Nickolas Wu.

<sup>172</sup><https://github.com/yallop/effects-bibliography>

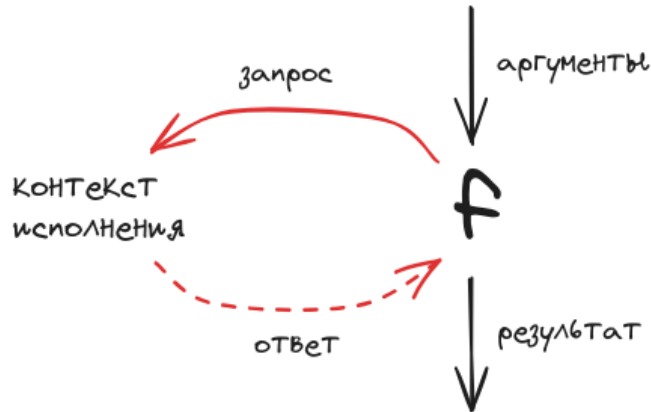


Figure 25: Effects as client-server interaction.

6 (w', x : xs)

2248 Thus, code made of pure functions is filled with inessential ceremony, behind which the business  
 2249 logic and essence are hidden. To focus on the important details, one needs to delegate all this  
 2250 bookkeeping to external code, sweep the uninteresting details under the rug. Then the code above  
 2251 can be rewritten, for example, as follows:

```

1  getList :: Int -> IO [Int]
2  getList n | n == 0 = pure []
3             | otherwise = do
4     x <- getInt
5     xs <- getList (n - 1)
6     return (x : xs)

```

2252 The abstract entity to which we will delegate details inessential to the given fragment of business  
 2253 logic we will call the **execution context**<sup>173</sup>. And an **effect** is the interaction of a function with the  
 2254 execution context, which occurs through calls to **effect operations** (for example, `getInt` from  
 2255 the example above)<sup>174</sup>. Thus, one can view the execution context as a server to which the client  
 2256 function sends requests and receives responses (see Fig. 25). In this model, such a function can  
 2257 violate both properties of pure functions.

2258 In practice, this context is the interpreter of an (embedded) language, and the effect operations  
 2259 are its constructs. If the language is embedded, one speaks of **user-defined effects**. Thus we  
 2260 again return to the task of building modular interpreters (5.5). Moreover, the implementation of  
 2261 a computational context can also delegate the implementation of certain functionality to another  
 2262 execution context, and so on. We obtain the already familiar tower of interpreters (5.1.1).

2263 Let us consider some examples of computational contexts and operations:

<sup>173</sup><https://okmij.org/ftp/Computation/having-effect.html>

<sup>174</sup>Sometimes the term “effect” also refers to the transfer function in dataflow analysis Møller and Schwartzbach [2012]. In essence, it describes the change of context.

- 2264 • Context — memory management subsystem, `modify` — effect operation: the context maintains  
2265 the state of memory cells for us;
- 2266 • Context — exception handler, `throw MyException` — effect operation: the context determines  
2267 how the error will be handled for us (note that here control does not return to the term);
- 2268 • Context — dependency injection settings, requesting functionality — effect operation: the  
2269 context determines the implementation of the functionality we use<sup>175</sup>. As we will see later,  
2270 this and similar simple effects can be implemented simply and efficiently (see 7.2.2).

2271 **Exercise 53** *Give more examples of computational contexts and operations.*

2272 When we talk about a standalone language in which we program (for example, Haskell), any  
2273 action of ours in the program is executed by it. That is, for instance, is addition an effect operation?  
2274 In that case it is reasonable to identify a subset of language constructs that are “interesting” enough  
2275 to be considered as producing an effect.

2276 Which language constructs should be considered “interesting”? Notice that on the one hand  
2277 it is good that effects hide some complexity from us, allowing us to focus on other things. On  
2278 the other hand, this is also bad, because we cease to observe this complexity, yet it permeates  
2279 our code and maintains implicit dependencies between its parts. Thus, effects require additional  
2280 care from the programmer<sup>176</sup>. Accordingly, it is precisely such non-trivial constructs that should  
2281 be considered “interesting.” At the very least, constructs whose use takes a function out of the  
2282 pure category should definitely be considered “interesting.” Also, these can be operations that make  
2283 complex non-local modifications to the control flow (maintained by the interpreter in the form of a  
2284 continuation).

2285 Ultimately, the choice of “interesting” constructs depends on the task and the developer’s  
2286 perspective<sup>173</sup>. Thus, constructs affecting arbitrary observable properties of code, such as termination  
2287 or computational complexity, can motivate considering recursive calls or long-running operations as  
2288 effects.

2289 Later we will learn to track and control the use of effects at the type level using effect systems  
2290 (see later 8).

## 2291 7.2 Effect handlers

2292 Effect handlers are a modern universal method for building modular interpreters of embedded  
2293 languages, directly implementing the client-server metaphor. As is often the case, handlers were  
2294 invented multiple times. In this chapter we will look at the main implementations that will best help  
2295 us understand the concept.

2296 The main idea of effect handlers is fairly simple. A language construct `handle` is introduced  
2297 that allows defining a computational context for a given scope, providing implementations of effect  
2298 operations. A construct `perform` is also introduced that allows invoking an effect operation (sending  
2299 a request to the context). Each operation has a set of parameters, as well as a “return address,”

---

<sup>175</sup>There exists the term **contextual polymorphism** — code in different contexts can have different behavior.

<sup>176</sup>(youtube) Kris Jenkins — Side-Effects Are The Complexity Iceberg.

2300 the continuation of the call site, to which it will return the result. For example, in the experimental  
2301 language Koka<sup>177</sup> a context providing some constant can be implemented as follows (`resume` is the  
2302 name of the call-site continuation, `perform` is inserted implicitly):

```
1 with handler
2   ctl ask() resume(21)
3   ask() + ask()
```

2303 If the nearest context cannot handle the required request, the request is delegated to an outer  
2304 context, and so on until a suitable context is found. The modularity of interpreters defined by  
2305 handlers is based on this idea.

### 2306 7.2.1 Handlers via delimited continuations

2307 As we have already seen, various effects can be implemented using access to the current  
2308 continuation (6.4.1, 6.4.2). Effect handlers extend this idea by using delimited continuations to  
2309 transfer control to various interpreters (handlers).

2310 It is known that the classical operators for manipulating delimited continuations, monadic  
2311 reflection, and effect handlers are expressible in terms of each other Forster et al. [2017].

### 2312 7.2.2 Efficient implementation of handlers

2313 In general, the speed of `perform` is determined by the speed of capturing and restoring delimited  
2314 continuations. However, there exists a class of operations that can be implemented much more  
2315 efficiently.

2316 If we look at the implementation of the `ask` operation, we see that its last action is calling  
2317 the continuation, returning control to the calling code. Such operations are called **tail-resumptive**;  
2318 they strongly resemble ordinary functions, except that their implementations are determined by  
2319 the context (handler). Thus, tail-resumptive operations can be implemented as implicit passing of  
2320 a dictionary of functions from the handler to `perform`, thereby avoiding expensive continuation  
2321 manipulations Xie et al. [2020]<sup>178</sup>.

### 2322 7.2.3 Embedded handlers as explicit client-server communication

2323 To better understand the semantics of handlers, let us implement a language with handlers  
2324 embedded in Haskell. We start with the variant proposed by Oleg Kiselyov, which most straightforwardly  
2325 encodes the idea of client-server communication between the term and the context Kiselyov et al.  
2326 [2013].

2327 Let us begin with the `ask` effect, which requests numbers from the context. We define the data  
2328 type of messages to the context: it is either the final result of the computation, or a request `Ask`  
2329 containing a “return address” — the current continuation:

---

<sup>177</sup><https://koka-lang.github.io/koka/doc/index.html>

<sup>178</sup>Handlers of tail-resumptive operations resemble co-pattern matching and, correspondingly, objects (see 4.6.2).

```
1 data Message res = Val res | Ask (Int -> Message res)
```

2330 We will assemble continuations into a specialized `Monad Cont` with a suitable response type  
2331 (see 6.2.2, 6.4.1):

```
1 newtype Eff res = Eff  
2   { runEff :: forall res' . (res -> Message res') -> Message res' }
```

2332 Then the `ask` effect is implemented simply as “sending” an `Ask` request with the current  
2333 continuation.<sup>179</sup>

```
1 ask :: Eff Int  
2 ask = Eff \k -> Ask k
```

2334 We implement the handler as a “server” that processes requests in a loop until the computation  
2335 sends the final result:

```
1 run :: Eff res -> Message res  
2 run comp = runEff comp Val  
  
3 runReader :: Eff res -> Int -> res  
4 runReader comp env = loop (run comp)  
5   where  
6     loop = \case  
7       Val res -> res  
8       Ask k -> loop (k env)
```

2336 Finally, we can write effectful code:

```
1 exampleReader :: Int -> Int  
2 exampleReader = runReader do  
3   x <- ask  
4   y <- ask  
5   pure (x + y)
```

2337 The `Eff` monad will turn it into a lazy list of messages:

```
1 exampleReader :: Int -> Int  
2 exampleReader = loop $  
3   Ask \x ->  
4   Ask \y ->  
5   Val (x + y)
```

---

<sup>179</sup>Earlier we similarly implemented generators, see 6.4.3.

## 2338 7.2.4 Extensible messages and forwarding

2339 Let us abstract the message type over the “shape” of requests that can participate in them (see  
2340 4.3.2):

```
1 data Message effs res = Val res | Request (effs (Message effs res))
```

2341 The previous message type is obtained by passing the following shape functor:

```
1 newtype Reader env msg = Ask (env -> msg)
```

2342 **Exercise 54** *Verify that `Message (Reader Int) res` is equivalent to the previous message type.*

2343 The coproduct of shape functors is a shape functor (see 5.5.1):

```
1 data (eff |> effs) a = L (eff a) | R (effs a)
```

2344 Now the ask operation admits the existence of other request types effs:

```
1 ask :: Eff (Reader env |> effs) env  
2 ask = Eff \k -> Request $ L $ Ask k
```

2345 The new handler processes only a portion of the requests, forwarding the rest to the handler  
2346 above (composing the continuations in the correct way):

```
1 runReader  
2   :: forall effs env res . Functor effs  
3   => Eff (Reader env |> effs) res  
4   -> env -> Eff effs res  
5 runReader comp env = loop (run comp)  
6   where  
7     loop :: Message (Reader env |> effs) res -> Eff effs res  
8     loop = \case  
9       Val res -> pure res  
10      Request (L (Ask k)) -> loop (k env)  
11      Request (R unknownReq) -> do  
12        response <- Eff \k -> Request (fmap k unknownReq)  
13        loop response
```

2347 Notice that uninterpreted effects remain in the resulting domain. The final domain is obtained  
2348 by applying all the necessary interpreters. This is how the stable denotations problem (5.5) is solved.

## 2349 7.2.5 Free monads

2350 To obtain a full-fledged solution in a simpler and more minimalistic way, let us turn to the most  
2351 classical implementation of handlers via free monads. First, let us discuss free monads themselves.

2352 Consider some algebraic structure, for example, a monoid (a neutral element and an associative  
2353 binary operation). From an arbitrary set  $X$  one can construct a certain monoid  $M(X)$  in the most  
2354 “economical” way — the **free monoid**. This is done as follows: to the set  $X$  one adds expression  
2355 trees with the monoid operations:

```
1 data M x = Element x | Mempty | Mappend (M x) (M x)
2 instance Monoid (M x) where
3   mempty = Mempty
4   mappend l r = Mappend l r
```

2356 One only needs to organize the new set so that there are no trees that are identical from the algebra’s  
2357 perspective (for example, `Mappend Mempty Mempty` and `Mempty`). In the case of a monoid, one can  
2358 choose the list type:

```
1 type M x = [x]
2 instance Monoid (M x) where
3   mempty = []
4   mappend = (++)
```

2359 Free monads are constructed analogously to free monoids. We use the definition of a monad  
2360 as a functor with operations `pure` and `join`. Then from an arbitrary functor  $F$  one can obtain a  
2361 monad  $Free(F)$  (we show this using the equivalent monad definition from Haskell):

```
1 data Free f a = Pure a | Join (f (Free f a))
2 instance Functor f => Monad (Free f) where
3   return = Pure
4   Pure x >>= k = k x
5   Join f >>= k = Join (fmap (>>= k) f)
```

2362 Let us rename the constructors:

```
1 data Term sig var = Var var | Op (sig (Term sig var))
2 (>>=) :: Term sig var -> (var -> Term sig var') -> Term sig var'
3 data MonoidSig subtree = Mempty | MAppend subtree subtree -- (Bool -> subtree)
```

2363 This is nothing other than an encoding of algebraic terms over the signature `sig` with variables  
2364 from the set `var`. And monadic `bind` is substitution.

## 2365 7.2.6 Handlers via free monads

2366 Notice that the following types are isomorphic: `Message`  $\cong$  `Free`  $\cong$  `Term`. Handlers in the classical  
 2367 form arose precisely in the process of studying how to describe effects as algebraic structures Bauer  
 2368 [2018]<sup>180</sup>.

2369 Instead of variables we store the pure result of the computation, and signatures are written in  
 2370 the form  $P \times (A \rightarrow K)$ , where  $P$  is the operation parameter and  $A$  is the operation result, by which  
 2371 the handler selects the needed continuation subtree for resuming the computation:

```
1 data Comp effs res = Res res | Op (effs (Comp effs res))
2 data Reader env comp = Ask () (env -> comp)
3 data State s comp = Get () (s -> comp) | Put s (() -> comp)
```

2372 The `Monad Comp` implementation will compose continuations for us in `do`-notation:

```
1 ask = Op (Ask () (\e -> Res e))
2 example = do
3   x <- ask
4   y <- ask
5   pure (x + y)
6 --  $\alpha\beta$ -equivalent to
7 example =
8   Op (Ask () (\x ->
9     Op (Ask () (\y ->
10      Res (x + y))))))
```

2373 The handler also simply folds the list of operations. Operations it cannot handle are left in the  
 2374 tree. To skip an unknown operation and interpret the subtree (continuation), `fmap` is used:

```
1 runReader
2 :: Functor effs => Comp (Reader env |> effs) res -> env -> Comp effs res
3 runReader comp env = case comp of
4   Res res -> Res res
5   Op (L (Ask () k)) -> runReader (k env) env
6   Op (R other) -> Op (fmap (`runReader` env) other)
```

2375 **Exercise 55** Build an example computation using `Reader` and `State`. What tree results after  
 2376 interpreting one of the effects?

2377 As usual, the fold can be generalized as a catamorphism:

```
1 handle
2 :: Functor effs => (res -> d) -> (effs d -> d) -> Comp effs res -> d
3 handle val alg = \case
4   Res res -> val res
5   Op eff -> alg $ fmap (handle val alg) eff
```

<sup>180</sup>(youtube) What is algebraic about algebraic effects and handlers — Andrej Bauer.

2378 Now in the implementation the continuation (subtree) is already interpreted into the desired domain:

```
1 runReader
2   :: Functor effs => Comp (Reader env |> effs) res -> env -> Comp effs res
3 runReader = handle (\res _env -> pure res) \case
4   L (Ask k) -> \env -> k env env
5   R other -> \env -> Op (fmap ($ env) other)
```

2379 When handlers are implemented as a built-in language feature, a design decision must be made:  
2380 in the provided continuation, is the current effect already interpreted or not? The first variant  
2381 is called **deep handlers**, the second — **shallow handlers**; they are expressible in terms of each  
2382 other Hillerström and Lindley [2018].

2383 Unfortunately, since the various effect variants are ordered both at the term level and at the  
2384 type level, we need an operation that turns a computation with fewer effects into a computation  
2385 with more:

```
1 liftF :: Functor effs => Comp effs res -> Comp (eff |> effs) res
2 liftF = \case
3   Res x -> Res x
4   Op effs -> Op $ R $ liftF <$> effs
5
6 example :: Comp (Reader Int |> State Int) ()
7 example = do
8   env <- ask
9   liftF (put env)
```

2386 The ordering can be eliminated using type classes Swierstra [2008] (3.4.8).

2387 One can notice that besides a large number of allocations, using **Monad Free** can lead to  
2388 quadratic complexity of code due to linear passes in each bind. Various alternative encoding schemes  
2389 exist Ploeg and Kiselyov [2014], Kiselyov and Ishii [2015].

2390 Free monads find other, albeit analogous, applications: trampolining Bjarnarson [2012] and  
2391 pipelines Kiselyov [2012b] [Bragilevsky, chapter 14].

## 2392 7.2.7 Applications of handlers

2393 Let us consider handlers of tail-resumptive operations, which treat continuations in a trivial  
2394 way. They are needed for propagating values and functionality down the stack. As we discussed  
2395 earlier 7.2.2, such handlers are analogous to having dynamic free variables or implicit function  
2396 arguments 3.1.2, and, along with them, records or anonymous classes 4.6.2.

2397 Among non-trivial use cases for continuations, we highlight the following:

- 2398 • The continuation can be left uninvoked; thus, an exception mechanism can be implemented.  
2399 However, in practice continuations may contain resource finalization logic. In that case, some  
2400 special handling is still needed<sup>181</sup>.

---

<sup>181</sup><https://koka-lang.github.io/koka/doc/book.html#sec-resource>

- 2401 • It can be invoked multiple times to emulate nondeterminism. This limits the use of efficient  
2402 mutable continuations 6.5 Leijen [2018]<sup>182</sup>, and also requires non-trivial handling when working  
2403 with resources.
- 2404 • It can be invoked not immediately. This is needed for implementing mutable state, generators,  
2405 and coroutines. For mutable state this is too expensive. Generators are a good application, but  
2406 a built-in generator implementation can be more efficient by saving on continuation allocations  
2407 (an `Iterator` is created directly)<sup>183</sup>. Implementing coroutines on top of handlers gives users  
2408 the ability to write their own schedulers, which is the most valued application of handlers at  
2409 the present time Sivaramakrishnan et al. [2021], Phipps-Costin et al. [2023].

## 2410 7.2.8 Monad transformers

2411 In the classical implementation of handlers via free monads 7.2.6, each handler produces an  
2412 intermediate tree of uninterpreted operations. The desire to deforest 4.4.1 these intermediate  
2413 trees arises naturally. One can go all the way and completely eliminate free monads and signature  
2414 functors Wu and Schrijvers [2015]. In doing so, we effectively obtain the tagless final shallow  
2415 embedding 5.4, another popular solution to the expression (stable denotations) problem — **monad**  
2416 **transformers** Liang et al. [1995], Jones [1995]<sup>184</sup>. In general, these approaches are equivalent in  
2417 expressiveness, but may in practice have different characteristics when embedded in a language Schrijvers  
2418 et al. [2019]. However, unlike transformers, handlers can be made a convenient built-in language  
2419 feature.

## 2420 7.2.9 Algebraicity and higher-order effects

2421 In the early 2000s, the idea emerged of describing effect operations not with monads directly,  
2422 but algebraically, using signatures and equations Plotkin and Power [2002], Bauer [2018]. Effects  
2423 in this formalism are composable by construction, just as the signatures of algebraic theories are  
2424 composable (concatenation of signatures is a signature 7.2.4); thus we have extensible syntax.  
2425 However, a restriction is imposed on operations in the form of the **algebraicity property** — an  
2426 operation commutes with the continuation and “bubbles up” to the top:

$$E[op(v, k)] \equiv op(v, \lambda x. E[k(x)])$$

2427 The corresponding effects are called **algebraic**.

2428 In our implementation (7.2.6), the algebraicity restriction can be seen as follows. Our operation  
2429 signatures are functors in their continuations. This is used by both monadic `bind` for accumulating  
2430 continuations and handlers for interpreting the continuations of unknown operations.

2431 Many useful effects are algebraic, but not all. For example, catching exceptions `catch` is not an  
2432 algebraic operation, since its signature would be written as follows:

<sup>182</sup><https://koka-lang.github.io/koka/doc/book.html#sec-multi-resume>

<sup>183</sup><https://csharpindepth.com/Articles/IteratorBlockImplementation>

<sup>184</sup><https://hackage.haskell.org/package/mtl>

```
1 data Catch e comp = Catch { try :: comp, onExn :: (e -> comp), next :: comp }
```

2433 In this case the only reasonable functor implementation would also work with nested continuations,  
2434 and the entire program, as a result of bind, would end up inside the try block.

2435 In other words, *algebraic effects cannot accept other effectful computations as arguments*. It  
2436 was precisely to model catch that handlers were originally proposed Plotkin and Pretnar [2013]<sup>185</sup>.  
2437 Thus, handlers have two tasks: delimiting the scope of certain functionality and interpreting effects.  
2438 However, as we know, the order of handlers determines the resulting domain and, accordingly, the  
2439 semantics. At the same time, the task of scope delimitation fixes the handler's position, which  
2440 limits expressiveness and makes some domains inaccessible.

2441 Operations that do not satisfy the algebraicity property correspond to **higher-order effects**,  
2442 which can accept other effectful computations as arguments. Embeddings of handlers for such  
2443 effects exist along with an extensive body of research on the topic Wu et al. [2014]<sup>186</sup><sup>187</sup> Yang  
2444 et al. [2022].

2445 The key to higher-order effects lies in the ability to execute the effects of argument computations  
2446 in the context of those handlers that are available at the call site of the higher-order effect van der  
2447 Rest et al. [2022]. This is related to another idea, bidirectional effects, which allows the operation's  
2448 implementation to produce effects at its call site (for example, throw exceptions there), an extremely  
2449 necessary practical capability Zhang et al. [2020]<sup>188</sup>.

---

<sup>185</sup>Despite the fact that algebraic effects are simply a class of "well-behaved" operations, often when they are mentioned, effect handlers are implied.

<sup>186</sup><https://github.com/fused-effects/fused-effects>

<sup>187</sup>(youtube) Building Haskell Programs with Fused Effects — Patrick Thomson

<sup>188</sup><https://effekt-lang.org/docs/concepts/bidirectional>

2450 **8 Effect Systems TODO**

2451 TODO

## 2452 9 Datatype-generic programming TODO

2453 Earlier we said that everything programmers write is interpreters. It turns out that many of them  
2454 can be obtained automatically by reaching a new level of generality. And code that does not have  
2455 to be written — should not be written.

2456 As a rule, such a situation arises when the interpretation is fully determined by the description  
2457 (the shape) of the type. That is, by which data constructors are defined for it, with which fields,  
2458 how those fields are named. . . Examples of such interpretations include serialisation into a particular  
2459 format (e.g. json), structural equality, and so on.

2460 The solution is for the language to support **structural polymorphism** [Maguire, a, chapter 13],  
2461 or **datatype-generic programming**, as it is sometimes called.

2462 One can imagine a language in which datatype-generic programming is just ordinary programming Chapman  
2463 et al. [2010]. To achieve this, one takes a language with dependent types together with a distinguished  
2464 universe of descriptions `Desc`. Concrete types are then obtained by applying an interpretation  
2465 function `Desc -> Type`. Thus, parametric polymorphism is embodied as generalisation over a type,  
2466 while structural polymorphism is generalisation over its description.

2467 TODO

### 2468 9.1 Specialisation

$$\llbracket \llbracket spec \rrbracket (p, x) \rrbracket (y) \equiv \llbracket p \rrbracket (x, y)$$

2469 TODO Liu [2023]

### 2470 9.2 Ad-hoc solutions

2471 TODO

### 2472 9.3 Approaches to datatype-generic programming in Haskell

2473 TODO

#### 2474 9.3.1 Template Haskell

2475 TODO<sup>189</sup>

#### 2476 9.3.2 Deriving strategies

2477 TODO<sup>190</sup>

---

<sup>189</sup><https://markkarpov.com/tutorial/th.html>

<sup>190</sup>[https://ghc.gitlab.haskell.org/ghc/doc/users\\_guide/exts/deriving\\_strategies.html](https://ghc.gitlab.haskell.org/ghc/doc/users_guide/exts/deriving_strategies.html)

2478 **9.3.3 GHC.Generics**

2479     TODO<sup>191</sup>

2480 **9.3.4 Sum Of Products**

2481     TODO

2482 **9.3.5 Uniplate**

2483     TODO

---

<sup>191</sup>[https://ghc.gitlab.haskell.org/ghc/doc/users\\_guide/exts/generics.html](https://ghc.gitlab.haskell.org/ghc/doc/users_guide/exts/generics.html)

## 2484 10 Persistence and Optics TODO

2485 In functional programming it is common to work with data persistently. In this chapter we  
2486 consider when and why this is important. As well as functional optics — a set of abstractions for  
2487 convenient work with data.

### 2488 10.1 Persistent data structures

2489 Data structures are divided into **mutable** and **immutable**. This is a classification solely by  
2490 internal implementation — whether mutable memory cells are used or not.

2491 By usage, one distinguishes **ephemeral** and **persistent** data structures. Changes to an ephemeral  
2492 data structure can be observed through the same reference:

```
1 let xs = MutableList<Int>.empty()  
2 xs.add(42)  
3 print(xs)
```

2493 Whereas modifying a persistent one produces new references each time, while the previous versions  
2494 of the structure remain accessible through the old ones:

```
1 let xs = []  
2 let xs' = 42 : xs  
3 print xs
```

2495 It might seem that ephemeral structures are mutable and persistent ones are immutable. In  
2496 reality these are more or less orthogonal classifications. For instance, one can provide a persistent  
2497 interface to a mutable data structure via copying, and an ephemeral interface to an immutable one  
2498 (using the **State** monad).

2499 In practice, when people talk about persistent data structures, they often mean data structures  
2500 with a persistent interface that are specially optimized for it (not requiring a full copy on every  
2501 operation). Such data structures can be surprisingly optimal and diverse Okasaki [1999].

2502 For example, one can implement an efficient persistent array with logarithmic complexity for all  
2503 operations. In Haskell, such a data structure is **Seq**<sup>192,193</sup>. If small arrays are stored in the nodes,  
2504 which modern processor architectures can efficiently copy, one can substantially reduce the tree  
2505 height and the algorithmic complexity of operations (e.g. `scala.immutable.Vector`).

2506 **Exercise 56** *Implement a persistent implicit-key treap in a non-Haskell language. Which features*  
2507 *of Haskell complicate the use of this structure?*

2508 When should one use ephemeral, and when persistent data structures? Comparing the two, one  
2509 can find that ephemeral structures can be implemented more efficiently in many cases: processor

---

<sup>192</sup><https://hackage.haskell.org/package/containers-0.7/docs/Data-Sequence.html>

<sup>193</sup><http://www.staff.city.ac.uk/~ross/papers/FingerTree.html>

2510 caches work better with local data, fewer allocations and indirections. At the same time, persistent  
2511 structures are bound to be trees, in order to reallocate not the entire structure but only the path  
2512 to the root.

2513 Persistent structures, on the other hand, allow writing more modular and thread-safe code that  
2514 need not account for the possibility of the structure being modified through a reference. Whereas  
2515 working with an ephemeral structure is not pure code but involves producing observable side effects.  
2516 Also, combining persistent results of different calls can be cheaper; for example, concatenation of  
2517 persistent arrays is cheaper than concatenation of ephemeral ones (logarithmic complexity versus  
2518 linear).

2519 Thus, within a limited, easily surveyable scope it is better to use ephemeral structures due  
2520 to their efficiency (for example, to initially populate a collection with elements). However, across  
2521 abstraction boundaries it is better to pass only persistent structures (or rely on an effect system,  
2522 see 8). That is, each data structure should support two phases of its life. For example, this is how  
2523 it is done in Scala, where many persistent collections have a [Builder](#) version.

## 2524 10.2 The idea of optics

2525 **Functional optics** allows building functional references that focus on certain properties of data  
2526 structures. Subsequently, using eliminating functions, one can read a property or persistently update  
2527 it given a functional reference and an object. For example, we can focus on the first component of  
2528 a pair using the reference `_1` and read the first component:

```
1 ghci> view _1 (42, 43)
2 42
```

2529 Or we can focus on the property of membership of a certain element in a set `member`, set the  
2530 property to `True`, and obtain a new set with the desired element:

```
1 ghci> set (member 42) (Set.fromList [1, 2]) True
2 Set.fromList [1, 2, 42]
```

2531 Functional references compose (here, using `%`). So, if the set is in the first component of a pair,  
2532 we can still add an element to it:

```
1 ghci> set (_1 % member 42) (Set.fromList [1, 2], 3) True
2 (Set.fromList [1, 2, 42], 3)
```

2533 It is important to note that functional references, unlike ordinary pointers, are on the one hand  
2534 more abstract and do not necessarily simply refer to a memory region, and on the other hand are  
2535 detached from a particular object, thus more closely resembling pointers to methods in C++.

2536 Thus, optics helps bring the convenience of programming with persistent data structures closer  
2537 to that of programming with mutable ones. This is important because good programming practices  
2538 should be at least as accessible as less optimal ones.

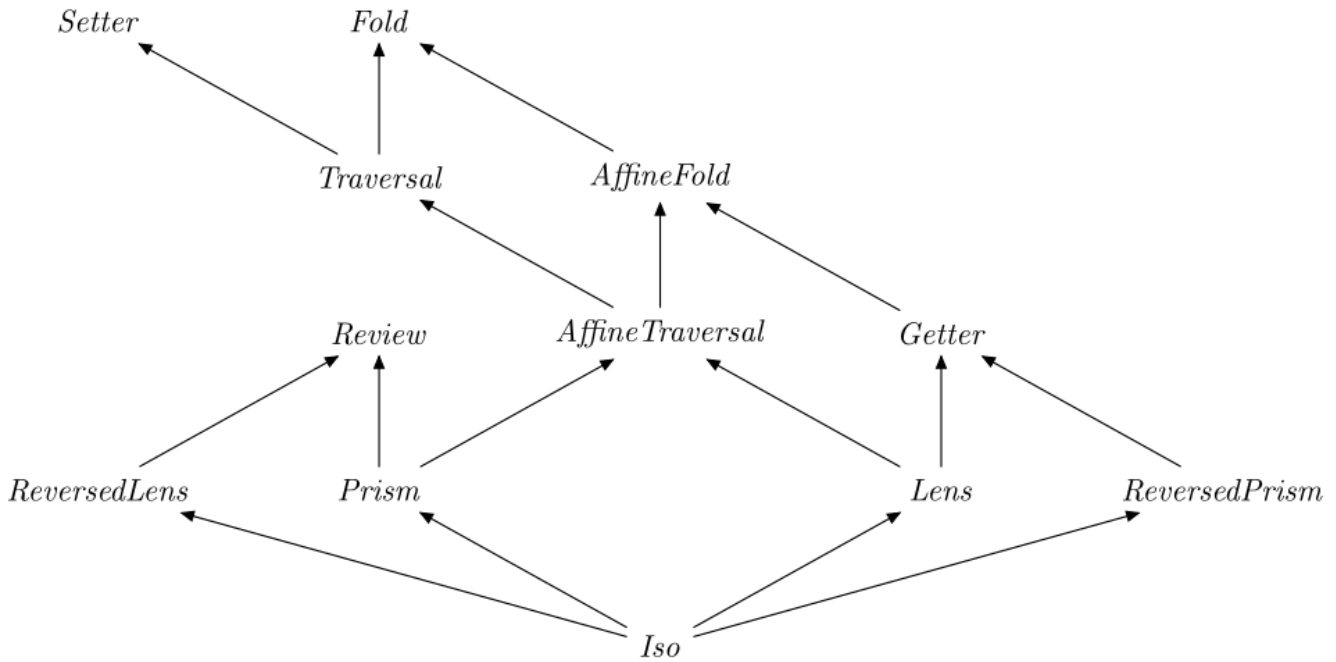


Figure 26: Hierarchy of optical devices of the optics library.

2539 Functional references can provide various usage interfaces, which are called **optical devices**.  
 2540 Devices form a hierarchy (Fig. 26): when composing two devices, the result is the nearest common  
 2541 ancestor with functionality equal to the intersection of the functionalities of the originals.

2542 Both references above are the optical device **lens**, which historically appeared first. Initially, lenses  
 2543 were proposed as a solution to the view-update problem in databases Bohannon et al. [2006], Foster  
 2544 et al. [2008]. Namely — how should the actual database be changed when the view is changed. Or  
 2545 how to reconstruct the whole after extracting and modifying a part. Lenses here serve as a means  
 2546 of bidirectional programming — they allow simultaneously describing both the view and the way  
 2547 to update. In parallel<sup>194</sup>, lenses are mentioned in a series of blog posts describing an attempt to  
 2548 work more conveniently with mutable state when writing games<sup>195196</sup> (lenses in games continue to  
 2549 delight<sup>197</sup>).

2550 Numerous optics libraries exist for both Haskell<sup>198199200</sup> and other languages: Kotlin<sup>201</sup>, Scala<sup>202</sup>,  
 2551 Swift<sup>203</sup> (language support).

<sup>194</sup><https://github.com/ekmett/lens/wiki/History-of-Lenses>

<sup>195</sup>(post) Making Haskell nicer for game programming

<sup>196</sup>(post) Haskell State Accessors (second attempt: Composability)

<sup>197</sup><http://www.timphilipwilliams.com/posts/2019-07-25-minecraft.html>

<sup>198</sup><https://hackage.haskell.org/package/optics-0.1/docs/Optics.html>

<sup>199</sup><http://lens.github.io/>

<sup>200</sup><https://github.com/marcosh/existential-optics/tree/main>

<sup>201</sup><https://arrow-kt.io/>

<sup>202</sup><https://github.com/optics-dev/Monocle>

<sup>203</sup><https://github.com/swiftlang/swift-evolution/blob/main/proposals/0161-key-paths.md>

## 2552 10.3 Using optics

2553 As an example, let us consider the use of several optical devices from the `optics` library, which  
2554 has excellent documentation<sup>204</sup>.

2555 Examples:

```
1 data Pet = Cat String | Dog Int deriving Show
2 newtype UserName = UserName { _getUserName :: String } deriving Show
3 data User = User { _userName :: UserName, _userCats :: [Pet] } deriving Show

4 exampleUser :: User
5 exampleUser = User (UserName "Bob") [Cat "Kitty", Dog 42, Dog 4]

6 exampleName :: User -> User
7 exampleName = over (userName % getUserName) (++ "!")

8 exampleIx :: User -> User
9 exampleIx = set (userName % getUserName % ix 2) 'k'

10 exampleFold :: User -> Int
11 exampleFold = getSum $ foldMapOf (userName % getUserName % folded) (const 1)

12 examplePrint :: User -> IO User
13 examplePrint = traverseOf (userCats % traversed % _Dog) (\x -> x <$ print x)
```

2556 TODO

## 2557 10.4 Data optics

### 2558 10.4.1 Lenses

2559 The simplest lens is formed by a pair of functions — viewing and setting a property:

```
1 data SimpleLens' s a = SimpleLens'
2   { view' :: s -> a
3     , set'  :: s -> a -> s
4   }
```

2560 Natural laws are imposed on these functions:

```
1 view l (set l s x) ≡ x
2 set l s (view l s) ≡ s
3 set l (view l s x) y ≡ set l s y
```

---

<sup>204</sup><https://hackage.haskell.org/package/optics-0.4.2.1/docs/Optics.html>

2561 For example, one can easily modify a user's age:

```
1 newtype Age = Age { _getAge :: Int }
2 data User = User { _userName :: String, _userAge :: Age }

3 userAge :: SimpleLens' User Age
4 userAge = SimpleLens' { view' = _userAge, set' = \s x -> s { _userAge = x } }

5 getAge :: SimpleLens' Age Int
6 getAge = SimpleLens' { view' = _getAge, set' = \s x -> s { _getAge = x } }

7 ghci> set (userAge % getAge) user 1
```

2562 One can generalize lenses to polymorphic lenses, which allow recreating the structure with new  
2563 type parameters:

```
1 data SimpleLens s t a b = SimpleLens
2   { view :: s -> a
3     , set :: s -> b -> t
4   }

5 _1 :: SimpleLens (a, c) (b, c) a b
6 _1 = SimpleLens { view = \ (x, _) -> x, set = \ (x, c) y -> (y, c) }
```

2564 One can notice that a lens in this representation is a costate comonad coalgebra<sup>205</sup>. Indeed, the  
2565 pair of functions can be replaced with a function returning a pair:

```
1 data Store a s = Store (a, a -> s)
2 data DataLens s a = DataLens (s -> Store a s)
```

2566 If we define a comonad instance for `Store a s` and write out the laws of compatibility with the  
2567 coalgebra, we obtain exactly the lens laws.

2568 TODO

## 2569 10.4.2 Prisms

2570 Prisms allow obtaining a property that may be absent, and setting it when present. Thus, a  
2571 specific field of a sum type may fail to be extracted if an unexpected constructor is passed.

2572 Simple prisms can be defined analogously to lenses, accounting for possible absence of the  
2573 property:

```
1 data SimplePrism' s a = SimplePrism'
2   { preview' :: s -> Maybe a
3     , review' :: a -> s
4   }
```

---

<sup>205</sup>(youtube) Category Theory II 9.1: Lenses.

2574 A polymorphic prism must provide the structure with an updated type in case the preview fails:

```
1 data SimplePrism s t a b = SimplePrism
2   { preview :: s -> Either t a
3     , review :: b -> t
4   }
```

2575 One can define a prism for working with the contents of the `Left` constructor:

```
1 _Left :: SimplePrism (Either a c) (Either b c) a b
2 _Left = SimplePrism
3   { preview = \case Left a -> Right a; Right c -> Left (Right c)
4     , review = Left
5   }
```

2576 Composition of prisms is defined in a natural way.

### 2577 10.4.3 Composition of lenses and prisms

2578 Let us introduce a type class with which we will learn to compose various combinations of lenses  
2579 and prisms. The resulting device is uniquely determined by the operands of the composition.

```
1 class Composable o1 o2 o3 | o1 o2 -> o3 where
2   compose :: o1 s t a b -> o2 a b c d -> o3 s t c d
```

2580 Now we can naturally define composition for various devices. One can notice that this approach  
2581 requires a quadratic number of instances in the number of optical devices.

```
1 instance Composable SimpleLens SimpleLens SimpleLens where
2   compose :: SimpleLens s t a b -> SimpleLens a b c d -> SimpleLens s t c d
3   compose l2 l1 = SimpleLens
4     { view = view l1 . view l2
5       , set = \s x -> set l2 s (set l1 (view l2 s) x)
6     }
```

2582 Composition of lenses and prisms forms a different device — affine traversals, which are something  
2583 in between. The presence of a prism in the composition does not allow guaranteed viewing of the  
2584 property, while a lens does not allow creating a new value without the original.

```
1 data AffineTraversal s t a b = AffineTraversal
2   { apreview :: s -> Either t a
3     , aset :: s -> b -> t
4   }
5 instance Composable SimplePrism SimpleLens SimplePrism where
```

```
6   compose :: SimplePrism s t a b -> SimpleLens a b c d -> SimplePrism s t c d
7   compose p l = SimplePrism
8     { preview = fmap (view l) . preview p
9     , pset = \s d -> case preview p s of
10        Left t -> t
11        Right a -> pset p s (set l a d)
12     }
```

## 2585 **10.5 The path to profunctor optics**

2586 TODO

## 2587 **10.6 Generating optics**

2588 TODO

2589

*To be continued...*

## References

- 2590
- 2591 Christopher Strachey. Fundamental concepts in programming languages.  
2592 *Higher-order and symbolic computation*, 13:11–49, 2000. URL [https://facweb.cdm.depaul.edu/smitsch/courses/csc447fa23/assets/articles/](https://facweb.cdm.depaul.edu/smitsch/courses/csc447fa23/assets/articles/strachey-fundamental-concepts-in-programming-languages.pdf)  
2593 [strachey-fundamental-concepts-in-programming-languages.pdf](https://facweb.cdm.depaul.edu/smitsch/courses/csc447fa23/assets/articles/strachey-fundamental-concepts-in-programming-languages.pdf).  
2594
- 2595 Luca Cardelli and Peter Wegner. On understanding types, data abstraction, and polymorphism.  
2596 *ACM Computing Surveys (CSUR)*, 17(4):471–523, 1985. URL [https://doi.org/10.1145/](https://doi.org/10.1145/6041.6042)  
2597 [6041.6042](https://doi.org/10.1145/6041.6042).
- 2598 Benjamin C Pierce. *Types and programming languages*. MIT press, 2002.
- 2599 John Launchbury and Simon L Peyton Jones. State in haskell. *Lisp and symbolic computation*, 8(4):  
2600 293–341, 1995. URL [https://www.microsoft.com/en-us/research/wp-content/uploads/](https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/state-lasc.pdf)  
2601 [2016/07/state-lasc.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/state-lasc.pdf).
- 2602 Sandy Maguire. *Thinking with Types, Type-Level Programming in Haskell*. a. URL <https://leanpub.com/thinking-with-types/>  
2603 <https://leanpub.com/thinking-with-types/>?
- 2604 Simon Peyton Jones. Type inference as constraint solving: how ghc's type inference engine actually  
2605 works. Keynote talk at Zurihac 2019, 2019.
- 2606 Benjamin C Pierce and David N Turner. Local type inference. *Acm transactions on programming*  
2607 *languages and systems (toplas)*, 22(1):1–44, 2000. URL [https://doi.org/10.1145/345099.](https://doi.org/10.1145/345099.345100)  
2608 [345100](https://doi.org/10.1145/345099.345100).
- 2609 David Raymond Christiansen. Bidirectional typing rules: A tutorial. 2013. URL <https://davidchristiansen.dk/tutorials/bidirectional.pdf>.  
2610 <https://davidchristiansen.dk/tutorials/bidirectional.pdf>.
- 2611 Jana Dunfield and Neelakantan R Krishnaswami. Sound and complete bidirectional typechecking  
2612 for higher-rank polymorphism with existentials and indexed types. *Proceedings of the ACM on*  
2613 *Programming Languages*, 3(POPL):1–28, 2019. URL <https://doi.org/10.1145/3290322>.
- 2614 Alejandro Serrano, Jurriaan Hage, Simon Peyton Jones, and Dimitrios Vytiniotis. A  
2615 quick look at impredicativity. *Proceedings of the ACM on Programming Languages*, 4  
2616 (ICFP):1–29, 2020. URL [https://www.microsoft.com/en-us/research/publication/](https://www.microsoft.com/en-us/research/publication/a-quick-look-at-impredicativity/)  
2617 [a-quick-look-at-impredicativity/](https://www.microsoft.com/en-us/research/publication/a-quick-look-at-impredicativity/).
- 2618 Frank Emrich, Sam Lindley, Jan Stolarek, James Cheney, and Jonathan Coates. Freezeml: Complete  
2619 and easy type inference for first-class polymorphism. In *Proceedings of the 41st ACM SIGPLAN*  
2620 *Conference on Programming Language Design and Implementation*, pages 423–437, 2020. URL  
2621 <https://link.springer.com/article/10.1208/s12249-010-9382-3>.

2622 Brent A Yorgey, Stephanie Weirich, Julien Cretin, Simon Peyton Jones, Dimitrios Vytiniotis, and  
2623 José Pedro Magalhães. Giving haskell a promotion. In *Proceedings of the 8th ACM SIGPLAN*  
2624 *Workshop on Types in Language Design and Implementation*, pages 53–66, 2012. URL <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/p53-yorgey.pdf>.  
2625

2626 Stephanie Weirich, Justin Hsu, and Richard A Eisenberg. System fc with explicit kind equality. *ACM*  
2627 *SIGPLAN Notices*, 48(9):275–286, 2013. URL <https://doi.org/10.1145/2544174.2500599>.

2628 Vitaly Bragilevsky. *Haskell in Depth*. Manning. URL <https://www.manning.com/books/haskell-in-depth>.  
2629

2630 Bruno Blanchet. Escape analysis for object-oriented languages: application to java. *Acm Sigplan*  
2631 *Notices*, 34(10):20–34, 1999. URL <https://doi.org/10.1145/320385.320387>.

2632 Richard A Eisenberg and Simon Peyton Jones. Levity polymorphism. *ACM SIGPLAN Notices*, 52  
2633 (6):525–539, 2017. URL <https://doi.org/10.1145/3140587.3062357>.

2634 Paul Hudak, John Hughes, Simon Peyton Jones, and Philip Wadler. A history  
2635 of haskell: being lazy with class. In *Proceedings of the third ACM SIGPLAN*  
2636 *conference on History of programming languages*, pages 12–1, 2007. URL [https://www.researchgate.net/profile/Simon-Peyton-Jones/publication/221501761\\_](https://www.researchgate.net/profile/Simon-Peyton-Jones/publication/221501761_A_history_of_Haskell_Being_lazy_with_class/links/0c960517e31f50f743000000/A-history-of-Haskell-Being-lazy-with-class.pdf)  
2637 [A\\_history\\_of\\_Haskell\\_Being\\_lazy\\_with\\_class/links/0c960517e31f50f743000000/](https://www.researchgate.net/profile/Simon-Peyton-Jones/publication/221501761_A_history_of_Haskell_Being_lazy_with_class/links/0c960517e31f50f743000000/A-history-of-Haskell-Being-lazy-with-class.pdf)  
2638 [A-history-of-Haskell-Being-lazy-with-class.pdf](https://www.researchgate.net/profile/Simon-Peyton-Jones/publication/221501761_A_history_of_Haskell_Being_lazy_with_class/links/0c960517e31f50f743000000/A-history-of-Haskell-Being-lazy-with-class.pdf).  
2639

2640 Philip Wadler and Stephen Blott. How to make ad-hoc polymorphism less ad hoc. In *Proceedings*  
2641 *of the 16th ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, pages  
2642 60–76, 1989. URL <https://doi.org/10.1145/75277.75283>.

2643 Cordelia V Hall, Kevin Hammond, Simon L Peyton Jones, and Philip L Wadler. Type classes in  
2644 haskell. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 18(2):109–  
2645 138, 1996. URL <https://doi.org/10.1145/227699.227700>.

2646 Simon Peyton Jones, Mark Jones, and Erik Meijer. Type classes: an exploration of the design  
2647 space. In *In Haskell Workshop*, 1997. URL [https://courses.cs.washington.edu/courses/](https://courses.cs.washington.edu/courses/cse590p/06sp/multi.pdf)  
2648 [cse590p/06sp/multi.pdf](https://courses.cs.washington.edu/courses/cse590p/06sp/multi.pdf).

2649 Jeffrey R Lewis, John Launchbury, Erik Meijer, and Mark B Shields. Implicit parameters: Dynamic  
2650 scoping with static types. In *Proceedings of the 27th ACM SIGPLAN-SIGACT symposium on*  
2651 *Principles of programming languages*, pages 108–118, 2000. URL [https://doi.org/10.1145/](https://doi.org/10.1145/325694.325708)  
2652 [325694.325708](https://doi.org/10.1145/325694.325708).

2653 Simon Peyton Jones. Type inference as constraint solving: how ghc’s type inference engine  
2654 actually works, 2019. URL [https://www.microsoft.com/en-us/research/publication/](https://www.microsoft.com/en-us/research/publication/type-inference-as-constraint-solving-how-ghcs-type-inference-engine-actually-works/)  
2655 [type-inference-as-constraint-solving-how-ghcs-type-inference-engine-actually-works/](https://www.microsoft.com/en-us/research/publication/type-inference-as-constraint-solving-how-ghcs-type-inference-engine-actually-works/).

- 2656 Martin Sulzmann, Gregory J Duck, Simon Peyton-Jones, and Peter J Stuckey. Understanding  
2657 functional dependencies via constraint handling rules. *Journal of functional programming*, 17(1):  
2658 83–129, 2007a. URL <https://doi.org/10.1017/S0956796806006137>.
- 2659 Oleg Kiselyov and Chung-chieh Shan. Functional pearl: implicit configurations—or, type classes  
2660 reflect the values of types. In *Proceedings of the 2004 ACM SIGPLAN workshop on*  
2661 *Haskell*, pages 33–44, 2004. URL [https://d1wqtxts1xzle7.cloudfront.net/43582096/  
2662 Functional\\_pearl\\_implicit\\_configurations20160310-32037-1bu6179-libre.pdf](https://d1wqtxts1xzle7.cloudfront.net/43582096/Functional_pearl_implicit_configurations20160310-32037-1bu6179-libre.pdf).
- 2663 Filip Křikava, Heather Miller, and Jan Vitek. Scala implicits are everywhere: A large-scale study of  
2664 the use of scala implicits in the wild. *Proceedings of the ACM on Programming Languages*, 3  
2665 (OOPSLA):1–28, 2019. URL <https://doi.org/10.1145/3360589>.
- 2666 Bruno CdS Oliveira, Adriaan Moors, and Martin Odersky. Type classes as objects and implicits. *ACM*  
2667 *Sigplan Notices*, 45(10):341–360, 2010. URL [https://citeseerx.ist.psu.edu/document?  
2668 repid=rep1&type=pdf&doi=d30d65ca9ce7891352024a5c71ebe0ae8c41f7ac](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=d30d65ca9ce7891352024a5c71ebe0ae8c41f7ac).
- 2669 Dominique Devriese and Frank Piessens. On the bright side of type classes: instance arguments in  
2670 agda. *ACM SIGPLAN Notices*, 46(9):143–155, 2011. URL [https://archive.alvb.in/msc/  
2671 thesis/reading/typeclasses-agda\\_Devriese.pdf](https://archive.alvb.in/msc/thesis/reading/typeclasses-agda_Devriese.pdf).
- 2672 Conor McBride. Faking it simulating dependent types in haskell. *Journal of functional programming*,  
2673 12(4-5):375–392, 2002. URL <https://doi.org/10.1017/S0956796802004355>.
- 2674 Simon Peyton Jones, Andrew Tolmach, and Tony Hoare. Playing by the rules: rewriting as a practical  
2675 optimisation technique in ghc. In *Haskell workshop*, volume 1, pages 203–233, 2001. URL [https:  
2676 //www.microsoft.com/en-us/research/wp-content/uploads/2001/09/rules.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2001/09/rules.pdf).
- 2677 Sandy Maguire. *Algebra-Driven Design, Elegant Solutions from Simple Building Blocks*. b. URL  
2678 <https://leanpub.com/algebra-driven-design/>.
- 2679 Li-yao Xia. *Defunctionalization*. URL [https://poisson.chat/aquarium/  
2680 defunctionalization.pdf](https://poisson.chat/aquarium/defunctionalization.pdf).
- 2681 John C Reynolds. Definitional interpreters for higher-order programming languages. In *Proceedings*  
2682 *of the ACM annual conference-Volume 2*, pages 717–740, 1972. URL [https://doi.org/10.  
2683 1145/800194.805852](https://doi.org/10.1145/800194.805852).
- 2684 John C Reynolds. Definitional interpreters revisited. *Higher-Order and Symbolic Computation*, 11:  
2685 355–361, 1998. URL <https://doi.org/10.1023/A:1010075320153>.
- 2686 Jeremy Yallop and Leo White. Lightweight higher-kinded polymorphism. In *International Symposium*  
2687 *on Functional and Logic Programming*, pages 119–135. Springer, 2014. URL [https://www.cl.  
2688 cam.ac.uk/~jdy22/papers/lightweight-higher-kinded-polymorphism.pdf](https://www.cl.cam.ac.uk/~jdy22/papers/lightweight-higher-kinded-polymorphism.pdf).

- 2689 Manuel MT Chakravarty, Gabriele Keller, and Simon Peyton Jones. Associated type synonyms.  
2690 *ACM SIGPLAN Notices*, 40(9):241–253, 2005a. URL [https://www.microsoft.com/en-us/  
2691 research/wp-content/uploads/2005/01/at-syns.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2005/01/at-syns.pdf).
- 2692 Manuel MT Chakravarty, Gabriele Keller, Simon Peyton Jones, and Simon Marlow. Associated  
2693 types with class. In *Proceedings of the 32nd ACM SIGPLAN-SIGACT symposium on Principles  
2694 of programming languages*, pages 1–13, 2005b. URL [https://www.microsoft.com/en-us/  
2695 research/wp-content/uploads/2005/01/assoc.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2005/01/assoc.pdf).
- 2696 Tom Schrijvers, Simon Peyton Jones, Manuel Chakravarty, and Martin Sulzmann. Type checking  
2697 with open type functions. In *Proceedings of the 13th ACM SIGPLAN international conference  
2698 on Functional programming*, pages 51–62, 2008. URL [https://www.microsoft.com/en-us/  
2699 research/wp-content/uploads/2008/01/icfp2008.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2008/01/icfp2008.pdf).
- 2700 Richard A Eisenberg, Dimitrios Vytiniotis, Simon Peyton Jones, and Stephanie Weirich.  
2701 Closed type families with overlapping equations. *ACM SIGPLAN Notices*, 49(1):671–683,  
2702 2014. URL [https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/  
2703 pop1137-eisenberg.pdf](https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/pop1137-eisenberg.pdf).
- 2704 Oleg Kiselyov, Simon Peyton Jones, and Chung-chieh Shan. Fun with type functions. *Reflections  
2705 on the Work of CAR Hoare*, pages 301–331, 2010. URL [https://www.microsoft.  
2706 com/en-us/research/wp-content/uploads/2016/07/typefun.pdf?from=https:  
2707 //research.microsoft.com/~simonpj/papers/assoc-types/fun-with-type-funs/  
2708 typefun.pdf&type=exact](https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/typefun.pdf?from=https://research.microsoft.com/~simonpj/papers/assoc-types/fun-with-type-funs/typefun.pdf&type=exact).
- 2709 Mark P Jones. Type classes with functional dependencies. In *European Symposium  
2710 on Programming*, pages 230–244. Springer, 2000. URL [https://doi.org/10.1007/  
2711 3-540-46425-5\\_15](https://doi.org/10.1007/3-540-46425-5_15).
- 2712 Nada Amin, Tiark Rompf, and Martin Odersky. Foundations of path-dependent types. *Acm Sigplan  
2713 Notices*, 49(10):233–249, 2014. URL [http://lampwww.epfl.ch/~amin/dot/fpdt\\_post.pdf](http://lampwww.epfl.ch/~amin/dot/fpdt_post.pdf).
- 2714 Jan Stolarek, Simon Peyton Jones, and Richard A Eisenberg. Injective type families for haskell.  
2715 *ACM SIGPLAN Notices*, 50(12):118–128, 2015. URL [https://repository.brynmawr.edu/  
2716 cgi/viewcontent.cgi?article=1070&context=compsci\\_pubs](https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1070&context=compsci_pubs).
- 2717 Richard A Eisenberg and Jan Stolarek. Promoting functions to type families in haskell. *ACM  
2718 SIGPLAN Notices*, 49(12):95–106, 2014. URL [https://repository.brynmawr.edu/cgi/  
2719 viewcontent.cgi?article=1000&context=compsci\\_pubs](https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1000&context=compsci_pubs).
- 2720 Csongor Kiss, Tony Field, Susan Eisenbach, and Simon Peyton Jones. Higher-order type-level  
2721 programming in haskell. *Proceedings of the ACM on Programming Languages*, 3(ICFP):1–26,  
2722 2019. URL <https://dl.acm.org/doi/pdf/10.1145/3341706>.

- 2723 Dominic Orchard and Tom Schrijvers. Haskell type constraints unleashed. In *International*  
2724 *Symposium on Functional and Logic Programming*, pages 56–71. Springer, 2010. URL <https://kar.kent.ac.uk/57498/1/constraint-families.pdf>.  
2725
- 2726 Edsko de Vries and Andres Löb. True sums of products. In *Proceedings of the 10th ACM SIGPLAN*  
2727 *workshop on Generic programming*, pages 83–94, 2014. URL <https://doi.org/10.1145/2633628.2633634>.  
2728
- 2729 Gert-Jan Bottu, Georgios Karachalias, Tom Schrijvers, Bruno C d S Oliveira, and Philip Wadler.  
2730 Quantified class constraints. *ACM SIGPLAN Notices*, 52(10):148–161, 2017. URL <https://www.pure.ed.ac.uk/ws/portalfiles/portal/42495988/quantcc.pdf>.  
2731
- 2732 Martin Sulzmann, Manuel MT Chakravarty, Simon Peyton Jones, and Kevin Donnelly. System f  
2733 with type equality coercions. In *Proceedings of the 2007 ACM SIGPLAN international workshop*  
2734 *on Types in languages design and implementation*, pages 53–66, 2007b. URL <https://doi.org/10.1145/1190315.1190324>.  
2735
- 2736 Dimitrios Vytiniotis, Simon Peyton Jones, Tom Schrijvers, and Martin Sulzmann. Outsidein (x)  
2737 modular type inference with local assumptions. *Journal of functional programming*, 21(4-5):  
2738 333–412, 2011. URL <https://doi.org/10.1017/S0956796811000098>.
- 2739 Joachim Breitner, Richard A Eisenberg, Simon Peyton Jones, and Stephanie Weirich. Safe zero-  
2740 cost coercions for haskell. In *Proceedings of the 19th ACM SIGPLAN international conference*  
2741 *on Functional programming*, pages 189–202, 2014. URL [https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1013&context=compsci\\_pubs](https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1013&context=compsci_pubs).  
2742
- 2743 Simon Peyton Jones, Stephanie Weirich, Richard A Eisenberg, and Dimitrios Vytiniotis. A  
2744 reflection on types. In *A List of Successes That Can Change the World: Essays Dedicated*  
2745 *to Philip Wadler on the Occasion of His 60th Birthday*, pages 292–317. Springer, 2016.  
2746 URL [https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1002&context=compsci\\_pubs](https://repository.brynmawr.edu/cgi/viewcontent.cgi?article=1002&context=compsci_pubs).  
2747
- 2748 Simon Marlow, Simon Peyton Jones, Andrew Moran, and John Reppy. Asynchronous exceptions  
2749 in haskell. In *Proceedings of the ACM SIGPLAN 2001 conference on Programming language*  
2750 *design and implementation*, pages 274–285, 2001. URL <https://classes.cs.uchicago.edu/archive/2007/spring/32102-1/papers/p274-marlow.pdf>.  
2751
- 2752 Simon Marlow. An extensible dynamically-typed hierarchy of exceptions. In *Proceedings*  
2753 *of the 2006 ACM SIGPLAN Workshop on Haskell*, pages 96–106, 2006. URL  
2754 [https://archive.alvb.in/msc/03\\_infoafp/papers/2012-12-18\\_HoorCollege\\_](https://archive.alvb.in/msc/03_infoafp/papers/2012-12-18_HoorCollege_MarlowExtensibleExceptions_dk.pdf)  
2755 [MarlowExtensibleExceptions\\_dk.pdf](https://archive.alvb.in/msc/03_infoafp/papers/2012-12-18_HoorCollege_MarlowExtensibleExceptions_dk.pdf).
- 2756 Simon Peyton Jones. Tackling the awkward squad: monadic input/output, concurrency, exceptions,  
2757 and foreign-language calls in haskell. *NATO SCIENCE SERIES SUB SERIES III COMPUTER*

2758 *AND SYSTEMS SCIENCES*, 180:47–96, 2001. URL [https://citeseerx.ist.psu.edu/](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=2e6c9d76f9cb690dc18019fc894ba9572a8c2812)  
2759 [document?repid=rep1&type=pdf&doi=2e6c9d76f9cb690dc18019fc894ba9572a8c2812](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=2e6c9d76f9cb690dc18019fc894ba9572a8c2812).

2760 Ralf Hinze, Jennifer Hackett, and Daniel WH James. Functional pearl: F for functor. 2012. URL  
2761 [www.cs.ox.ac.uk/people/daniel.james/functor/functor.pdf](http://www.cs.ox.ac.uk/people/daniel.james/functor/functor.pdf).

2762 Ralf Hinze and Daniel WH James. Reason isomorphically! In *Proceedings of the 6th ACM*  
2763 *SIGPLAN workshop on Generic programming*, pages 85–96, 2010. URL [http://www.cs.ox.](http://www.cs.ox.ac.uk/people/ralf.hinze/publications/WGP10.pdf)  
2764 [ac.uk/people/ralf.hinze/publications/WGP10.pdf](http://www.cs.ox.ac.uk/people/ralf.hinze/publications/WGP10.pdf).

2765 Gérard Huet. The zipper. *Journal of functional programming*, 7(5):549–554, 1997. URL [https:](https://doi.org/10.1017/S0956796897002864)  
2766 [//doi.org/10.1017/S0956796897002864](https://doi.org/10.1017/S0956796897002864).

2767 Conor McBride. The derivative of a regular type is its type of one-hole contexts. *Unpublished*  
2768 *manuscript*, pages 74–88, 2001. URL [https://citeseerx.ist.psu.edu/document?repid=](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=7de4f6fddb11254d1fd5f8adfd67b6e0c9439eaa)  
2769 [rep1&type=pdf&doi=7de4f6fddb11254d1fd5f8adfd67b6e0c9439eaa](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=7de4f6fddb11254d1fd5f8adfd67b6e0c9439eaa).

2770 Michael Abbott, Thorsten Altenkirch, Neil Ghani, and Conor McBride. Derivatives of containers.  
2771 In *International Conference on Typed Lambda Calculi and Applications*, pages 16–30. Springer,  
2772 2003. URL <http://www.strictlypositive.org/derivcont.pdf>.

2773 Erik Meijer, Maarten Fokkinga, and Ross Paterson. Functional programming with bananas,  
2774 lenses, envelopes and barbed wire. In *Conference on functional programming languages and*  
2775 *computer architecture*, pages 124–144. Springer, 1991. URL [https://ris.utwente.nl/ws/](https://ris.utwente.nl/ws/portalfiles/portal/6142049/meijer91functional.pdf)  
2776 [portalfiles/portal/6142049/meijer91functional.pdf](https://ris.utwente.nl/ws/portalfiles/portal/6142049/meijer91functional.pdf).

2777 Erik Meijer and Graham Hutton. Bananas in space: Extending fold and unfold to exponential types.  
2778 In *Proceedings of the seventh international conference on Functional programming languages*  
2779 *and computer architecture*, pages 324–333, 1995. URL [https://dl.acm.org/doi/pdf/10.](https://dl.acm.org/doi/pdf/10.1145/224164.224225)  
2780 [1145/224164.224225](https://dl.acm.org/doi/pdf/10.1145/224164.224225).

2781 Olivier Danvy and Kevin Millikin. Refunctionalization at work. *Science of Computer Programming*,  
2782 74(8):534–549, 2009. ISSN 0167-6423. doi: [https://doi.org/10.1016/j.scico.2007.10.](https://doi.org/10.1016/j.scico.2007.10.007)  
2783 [007](https://doi.org/10.1016/j.scico.2007.10.007). URL <https://www.sciencedirect.com/science/article/pii/S0167642309000227>.  
2784 Special Issue on Mathematics of Program Construction (MPC 2006).

2785 Philip Wadler. Deforestation: Transforming programs to eliminate trees. In *European Symposium on*  
2786 *Programming*, pages 344–358. Springer, 1988. URL [https://link.springer.com/content/](https://link.springer.com/content/pdf/10.1007/3-540-19027-9_23.pdf)  
2787 [pdf/10.1007/3-540-19027-9\\_23.pdf](https://link.springer.com/content/pdf/10.1007/3-540-19027-9_23.pdf).

2788 СА Романенко. ВЫЯВЛЕНИЕ И ДОКАЗАТЕЛЬСТВО СВОЙСТВ ФУНКЦИОНАЛЬ-  
2789 НЫХ ПРОГРАММ МЕТОДАМИ СУПЕРКОМПИЛЯЦИИ. URL [https://www.keldysh.ru/](https://www.keldysh.ru/council/1/klyuchnikov-diss.pdf)  
2790 [council/1/klyuchnikov-diss.pdf](https://www.keldysh.ru/council/1/klyuchnikov-diss.pdf).

- 2791 Andrew Gill, John Launchbury, and Simon L Peyton Jones. A short cut to deforestation. In  
2792 *Proceedings of the conference on Functional programming languages and computer architecture*,  
2793 pages 223–232, 1993. URL <https://dl.acm.org/doi/pdf/10.1145/165180.165214>.
- 2794 Jeremy Gibbons. Unfolding abstract datatypes. In *International Conference on*  
2795 *Mathematics of Program Construction*, pages 110–133. Springer, 2008. URL  
2796 [https://ora.ox.ac.uk/objects/uuid:b51d2af6-25de-41ae-b3e7-9ac997de87b3/files/](https://ora.ox.ac.uk/objects/uuid:b51d2af6-25de-41ae-b3e7-9ac997de87b3/files/m7c19077c7bfba562a6cc9ceb3db62641)  
2797 [m7c19077c7bfba562a6cc9ceb3db62641](https://ora.ox.ac.uk/objects/uuid:b51d2af6-25de-41ae-b3e7-9ac997de87b3/files/m7c19077c7bfba562a6cc9ceb3db62641).
- 2798 William R Cook. On understanding data abstraction, revisited. In *Proceedings of the 24th ACM*  
2799 *SIGPLAN conference on Object oriented programming systems languages and applications*, pages  
2800 557–572, 2009. URL [https://www.phaazon.net/media/uploads/on\\_understanding\\_data\\_](https://www.phaazon.net/media/uploads/on_understanding_data_abstraction.pdf)  
2801 [abstraction.pdf](https://www.phaazon.net/media/uploads/on_understanding_data_abstraction.pdf).
- 2802 Duncan Coutts, Roman Leshchinskiy, and Don Stewart. Stream fusion: From lists  
2803 to streams to nothing at all. *ACM SIGPLAN Notices*, 42(9):315–326, 2007.  
2804 URL [https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=a4fad0d182605fcd155bebe3d620b7ffa0456968)  
2805 [a4fad0d182605fcd155bebe3d620b7ffa0456968](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=a4fad0d182605fcd155bebe3d620b7ffa0456968).
- 2806 Oleg Kiselyov, Aggelos Biboudis, Nick Palladinos, and Yannis Smaragdakis. Stream fusion,  
2807 to completeness. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of*  
2808 *Programming Languages*, pages 285–299, 2017. URL <https://arxiv.org/pdf/1612.06668>.
- 2809 Paul Downen, Zachary Sullivan, Zena M Ariola, and Simon Peyton Jones. Codata in action.  
2810 In *European Symposium on Programming*, pages 119–146. Springer International Publishing  
2811 Cham, 2019. URL [https://library.oapen.org/bitstream/handle/20.500.12657/23330/](https://library.oapen.org/bitstream/handle/20.500.12657/23330/1006825.pdf?sequence=1#page=132)  
2812 [1006825.pdf?sequence=1#page=132](https://library.oapen.org/bitstream/handle/20.500.12657/23330/1006825.pdf?sequence=1#page=132).
- 2813 Jon Louis Bentley. Little languages. *Commun. ACM*, 29(8):711–721, 1986. URL [https://doi.](https://doi.org/10.1145/6424.315691)  
2814 [org/10.1145/6424.315691](https://doi.org/10.1145/6424.315691).
- 2815 Jeremy Gibbons and Nicolas Wu. Folding domain-specific languages: deep and shallow embeddings  
2816 (functional pearl). In *Proceedings of the 19th ACM SIGPLAN international conference on*  
2817 *Functional programming*, pages 339–347, 2014. URL [https://ora.ox.ac.uk/objects/uuid:](https://ora.ox.ac.uk/objects/uuid:0547e439-48d5-4749-a5eb-811e65cfaaee/files/mf31b32c6dcbd81dd6a8ef43166a8cc91)  
2818 [0547e439-48d5-4749-a5eb-811e65cfaaee/files/mf31b32c6dcbd81dd6a8ef43166a8cc91](https://ora.ox.ac.uk/objects/uuid:0547e439-48d5-4749-a5eb-811e65cfaaee/files/mf31b32c6dcbd81dd6a8ef43166a8cc91).
- 2819 Jeremy Gibbons. Functional programming for domain-specific languages. In *Central*  
2820 *European Functional Programming School*, pages 1–28. Springer, 2013. URL  
2821 [https://ora.ox.ac.uk/objects/uuid:489909da-8818-4d6e-a4fd-0107b589ee76/files/](https://ora.ox.ac.uk/objects/uuid:489909da-8818-4d6e-a4fd-0107b589ee76/files/m00be42b36cabbc6d11c8748fac8c7c71)  
2822 [m00be42b36cabbc6d11c8748fac8c7c71](https://ora.ox.ac.uk/objects/uuid:489909da-8818-4d6e-a4fd-0107b589ee76/files/m00be42b36cabbc6d11c8748fac8c7c71).
- 2823 Benjamin C Pierce, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cătălin Hrițcu, Vilhelm  
2824 Sjöberg, and Brent Yorgey. Software foundations. *Webpage: http://www.cis.upenn.*  
2825 *edu/bcpierce/sf/current/index.html*, 16, 2010. URL [https://idris-hackers.github.io/](https://idris-hackers.github.io/software-foundations/pdf/sf-idris-2018.pdf)  
2826 [software-foundations/pdf/sf-idris-2018.pdf](https://idris-hackers.github.io/software-foundations/pdf/sf-idris-2018.pdf).

2827 Graham Hutton. Fold and unfold for program semantics. *ACM SIGPLAN Notices*, 34(1):280–288,  
2828 1998. URL <https://dl.acm.org/doi/abs/10.1145/291251.289457>.

2829 Robert Nystrom. *Crafting Interpreters*. URL <https://craftinginterpreters.com/contents.html>.

2831 Dom Hutchinson. Language engineering - notes. Course COMS22201 notes. URL <https://github.com/dajhutchinson/Language-Engineering>.

2833 Simon Marlow. Parallel and concurrent programming in haskell. In *Central*  
2834 *European Functional Programming School*, pages 339–401. Springer, 2011.  
2835 URL <https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=bd127ea1952996864c1542e1453973a78973ad5c>.

2837 Arthur Charguéraud. The locally nameless representation. *Journal of automated reasoning*, 49:  
2838 363–408, 2012. URL <https://www.chargueraud.org/research/2009/ln/main.pdf>.

2839 Stephen Chong. *CS153: Compilers Lecture 12: Closures and Environments*. URL <https://groups.seas.harvard.edu/courses/cs153/2018fa/lectures/Lec12-Functions.pdf>.

2841 Oleg Kiselyov. Typed tagless final interpreters. In *Generic and indexed programming: International*  
2842 *spring school, sSGIP 2010, oxford, uK, march 22-26, 2010, revised lectures*, pages 130–174.  
2843 Springer, 2012a. URL <https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=f0b6a04331821a294bad4d08975637b4f363ec5a#page=136>.

2845 Frank Pfenning and Conal Elliott. Higher-order abstract syntax. *ACM sigplan notices*, 23(7):  
2846 199–208, 1988. URL <https://doi.org/10.1145/960116.54010>.

2847 Jeff Epstein, Andrew P Black, and Simon Peyton-Jones. Towards haskell in the cloud.  
2848 In *Proceedings of the 4th ACM symposium on Haskell*, pages 118–129, 2011. URL  
2849 <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/07/remote.pdf?from=https://research.microsoft.com/en-us/um/people/simonpj/papers/parallel/remote.pdf&type=exact>.

2852 Jacques Carette, Oleg Kiselyov, and Chung-chieh Shan. Finally tagless, partially  
2853 evaluated: Tagless staged interpreters for simpler typed languages. In *Programming*  
2854 *Languages and Systems: 5th Asian Symposium, APLAS 2007, Singapore, November*  
2855 *29-December 1, 2007. Proceedings 5*, pages 222–238. Springer, 2007. URL <https://www.cambridge.org/core/journals/journal-of-functional-programming/article/finally-tagless-partially-evaluated-tagless-staged-interpreters-for-simpler-typed-lang/7B2DC44A2127EBBA71ADE63809D9425F>.

2859 Eugenio Moggi. *Computational lambda-calculus and monads*. University of Edinburgh, Department  
2860 of Computer Science, Laboratory for ..., 1988. URL <http://www.lfcs.inf.ed.ac.uk/reports/88/ECS-LFCS-88-66/ECS-LFCS-88-66.pdf>.

2861

- 2862 Philip Wadler. Comprehending monads. In *Proceedings of the 1990 ACM Conference on LISP and*  
2863 *Functional Programming*, pages 61–78, 1990. URL [https://dl.acm.org/doi/pdf/10.1145/](https://dl.acm.org/doi/pdf/10.1145/91556.91592)  
2864 [91556.91592](https://dl.acm.org/doi/pdf/10.1145/91556.91592).
- 2865 Philip Wadler. The essence of functional programming. In *Proceedings of the 19th ACM SIGPLAN-*  
2866 *SIGACT symposium on Principles of programming languages*, pages 1–14, 1992. URL <https://dl.acm.org/doi/pdf/10.1145/143165.143169>.
- 2868 Conor McBride and Ross Paterson. Applicative programming with effects. *Journal of functional*  
2869 *programming*, 18(1):1–13, 2008. URL <https://doi.org/10.1017%2FS0956796807006326>.
- 2870 John C Reynolds. The discoveries of continuations. *Lisp and symbolic computation*, 6:233–247,  
2871 1993. URL [https://ics.uci.edu/~jajones/INF102-S18/readings/08\\_histcont.pdf](https://ics.uci.edu/~jajones/INF102-S18/readings/08_histcont.pdf).
- 2872 Peter J Landin. Histories of discoveries of continuations: Belles-lettres with equivocal  
2873 tenses. In *Proceedings of the Second ACM SIGPLAN Workshop on Continuations*  
2874 *(CW'97), Technical report BRICS NS-96-13, University of Aarhus*, page 1. Citeseer,  
2875 1997. URL [https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=14c609276166517ed1afb56e6b84a41ba157030b)  
2876 [14c609276166517ed1afb56e6b84a41ba157030b](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=14c609276166517ed1afb56e6b84a41ba157030b).
- 2877 Steven E Ganz, Daniel P Friedman, and Mitchell Wand. Trampoline style. In *Proceedings of the*  
2878 *fourth ACM SIGPLAN international conference on Functional programming*, pages 18–27, 1999.  
2879 URL <https://doi.org/10.1145/317636.317779>.
- 2880 RO Bjarnarson. Stackless scala with free monads. *Scala Days*, 2012. URL [https://days2012.](https://days2012.scala-lang.org/sites/days2012/files/bjarnason_trampoline.pdf)  
2881 [scala-lang.org/sites/days2012/files/bjarnason\\_trampoline.pdf](https://days2012.scala-lang.org/sites/days2012/files/bjarnason_trampoline.pdf).
- 2882 Conor McBride. Clowns to the left of me, jokers to the right (pearl) dissecting data  
2883 structures. In *Proceedings of the 35th annual ACM SIGPLAN-SIGACT symposium on Principles*  
2884 *of programming languages*, pages 287–295, 2008. URL [https://web.archive.org/web/](https://web.archive.org/web/20170705104305id_/http://www.cis.upenn.edu/~cis39903/static/clowns.pdf)  
2885 [20170705104305id\\_/http://www.cis.upenn.edu/~cis39903/static/clowns.pdf](https://web.archive.org/web/20170705104305id_/http://www.cis.upenn.edu/~cis39903/static/clowns.pdf).
- 2886 Daniel Hillerström. Foundations for programming and implementing effect handlers. 2022. URL  
2887 <http://dx.doi.org/10.7488/era/2122>.
- 2888 Serkan Muhcu, Philipp Schuster, Michel Steuwer, and Jonathan Immanuel Brachthäuser. Multiple  
2889 resumptions and local mutable state, directly. *Proceedings of the ACM on Programming*  
2890 *Languages*, 9(ICFP):704–733, 2025. URL <https://doi.org/10.1145/3747529>.
- 2891 Olivier Danvy. Defunctionalized interpreters for programming languages. *ACM Sigplan Notices*, 43  
2892 (9):131–142, 2008. URL <http://jfla.inria.fr/2014/danvy-ICFP08.pdf>.
- 2893 Jeremy Gibbons. Continuation-passing style, defunctionalization, accumulations, and associativity.  
2894 *arXiv preprint arXiv:2111.10413*, 2021. URL <https://doi.org/10.48550/arXiv.2111.10413>.

- 2895 Olivier Danvy. An analytical approach to program as data objects, 2006. URL [https://www.cs.tufts.edu/comp/150FP/archive/olivier-danvy/d\\_scientia.pdf](https://www.cs.tufts.edu/comp/150FP/archive/olivier-danvy/d_scientia.pdf).
- 2896
- 2897 Atze van der Ploeg and Oleg Kiselyov. Reflection without remorse: revealing a hidden sequence to  
2898 speed up monadic reflection. In *Proceedings of the 2014 ACM SIGPLAN Symposium on Haskell*,  
2899 pages 133–144, 2014. URL <https://doi.org/10.1145/2633357.2633360>.
- 2900 Pierre-Louis Curien and Hugo Herbelin. The duality of computation. *ACM sigplan notices*, 35(9):  
2901 233–243, 2000. URL <https://doi.org/10.1145/357766.351262>.
- 2902 R Kent Dyvbig, Simon Peyton Jones, and Amr Sabry. A monadic framework for delimited  
2903 continuations. *Journal of functional programming*, 17(6):687–730, 2007. URL <https://doi.org/10.1017/S0956796807006259>.
- 2904
- 2905 Andrzej Filinski. Representing monads. In *Proceedings of the 21st ACM SIGPLAN-SIGACT*  
2906 *symposium on Principles of programming languages*, pages 446–457, 1994. URL <https://doi.org/10.1145/174675.178047>.
- 2907
- 2908 Ana Lúcia De Moura and Roberto Ierusalimsky. Revisiting coroutines. *ACM Transactions on*  
2909 *Programming Languages and Systems (TOPLAS)*, 31(2):1–31, 2009. URL <https://doi.org/10.1145/1462166.1462167>.
- 2910
- 2911 Roman Elizarov, Mikhail Belyaev, Marat Akhin, and Ilmir Usmanov. Kotlin coroutines: design and  
2912 implementation. In *Proceedings of the 2021 ACM SIGPLAN International Symposium on New*  
2913 *Ideas, New Paradigms, and Reflections on Programming and Software*, pages 68–84, 2021. URL  
2914 [https://www.researchgate.net/profile/Mikhail-Belyaev-3/publication/355361443\\_](https://www.researchgate.net/profile/Mikhail-Belyaev-3/publication/355361443_Kotlin_coroutines_design_and_implementation/links/624c63c34f88c3119ce12a49/Kotlin-coroutines-design-and-implementation.pdf)  
2915 [Kotlin\\_coroutines\\_design\\_and\\_implementation/links/624c63c34f88c3119ce12a49/](https://www.researchgate.net/profile/Mikhail-Belyaev-3/publication/355361443_Kotlin_coroutines_design_and_implementation/links/624c63c34f88c3119ce12a49/Kotlin-coroutines-design-and-implementation.pdf)  
2916 [Kotlin-coroutines-design-and-implementation.pdf](https://www.researchgate.net/profile/Mikhail-Belyaev-3/publication/355361443_Kotlin_coroutines_design_and_implementation/links/624c63c34f88c3119ce12a49/Kotlin-coroutines-design-and-implementation.pdf).
- 2917 Anders Møller and Michael I Schwartzbach. Static program analysis. *Notes. Feb*, 2012. URL  
2918 <https://users-cs.au.dk/amoeller/spa/spa.pdf>.
- 2919 Yannick Forster, Ohad Kammar, Sam Lindley, and Matija Pretnar. On the expressive power of user-  
2920 defined effects: Effect handlers, monadic reflection, delimited control. *Proceedings of the ACM*  
2921 *on Programming Languages*, 1(ICFP):1–29, 2017. URL <https://doi.org/10.1145/3110257>.
- 2922 Ningning Xie, Jonathan Immanuel Brachthäuser, Daniel Hillerström, Philipp Schuster, and Daan  
2923 Leijen. Effect handlers, evidently. *Proceedings of the ACM on Programming Languages*, 4  
2924 (ICFP):1–29, 2020. URL <https://doi.org/10.1145/3408981>.
- 2925 Oleg Kiselyov, Amr Sabry, and Cameron Swords. Extensible effects: an alternative to monad  
2926 transformers. *ACM SIGPLAN Notices*, 48(12):59–70, 2013. URL <https://doi.org/10.1145/2578854.2503791>.
- 2927
- 2928 Andrej Bauer. What is algebraic about algebraic effects and handlers? *arXiv preprint*  
2929 *arXiv:1807.05923*, 2018. URL <https://doi.org/10.48550/arXiv.1807.05923>.

- 2930 Daniel Hillerström and Sam Lindley. Shallow effect handlers. In *Asian Symposium on Programming*  
2931 *Languages and Systems*, pages 415–435. Springer, 2018. URL [https://www.research.ed.ac.uk/files/76099718/shallow\\_effect\\_handlers.pdf](https://www.research.ed.ac.uk/files/76099718/shallow_effect_handlers.pdf).
- 2933 Wouter Swierstra. Data types à la carte. *Journal of functional programming*, 18(4):423–  
2934 436, 2008. URL <https://www.cs.tufts.edu/~nr/cs257/archive/wouter-swierstra/DataTypesALaCarte.pdf>.
- 2936 Oleg Kiselyov and Hiromi Ishii. Freer monads, more extensible effects. *ACM SIGPLAN Notices*, 50  
2937 (12):94–105, 2015. URL <https://doi.org/10.1145/2887747.2804319>.
- 2938 Oleg Kiselyov. Iteratees. In *FLOPS'12: Proceedings of the 11th international conference*  
2939 *on Functional and Logic Programming*, 2012b. URL <https://okmij.org/ftp/Haskell/Iteratee/describe.pdf>.
- 2941 Daan Leijen. Algebraic effect handlers with resources and deep finalization. Technical report,  
2942 Tech. Rep. MSR-TR-2018-10, Microsoft Research (April 2018), 2018. URL <https://www.microsoft.com/en-us/research/wp-content/uploads/2018/04/resource-v1.pdf>.
- 2944 KC Sivaramakrishnan, Stephen Dolan, Leo White, Tom Kelly, Sadiq Jaffer, and Anil Madhavapeddy.  
2945 Retrofitting effect handlers onto ocaml. In *Proceedings of the 42nd ACM SIGPLAN International*  
2946 *Conference on Programming Language Design and Implementation*, pages 206–221, 2021. URL  
2947 <https://doi.org/10.1145/3453483.3454039>.
- 2948 Luna Phipps-Costin, Andreas Rossberg, Arjun Guha, Daan Leijen, Daniel Hillerström,  
2949 KC Sivaramakrishnan, Matija Pretnar, and Sam Lindley. Continuing webassembly with effect  
2950 handlers. *Proceedings of the ACM on Programming Languages*, 7(OOPSLA2):460–485, 2023.  
2951 URL <https://doi.org/10.1145/3622814>.
- 2952 Nicolas Wu and Tom Schrijvers. Fusion for free: Efficient algebraic effect handlers. In *International*  
2953 *Conference on Mathematics of Program Construction*, pages 302–322. Springer, 2015. URL  
2954 <https://lirias.kuleuven.be/retrieve/322544>.
- 2955 Sheng Liang, Paul Hudak, and Mark Jones. Monad transformers and modular interpreters. In  
2956 *Proceedings of the 22nd ACM SIGPLAN-SIGACT symposium on Principles of programming*  
2957 *languages*, pages 333–343, 1995. URL <https://dl.acm.org/doi/pdf/10.1145/199448.199528>.
- 2959 Mark P Jones. Functional programming with overloading and higher-order polymorphism. In  
2960 *International School on Advanced Functional Programming*, pages 97–136. Springer, 1995. URL  
2961 <http://web.cecs.pdx.edu/mpj/pubs/springschool95.pdf>.
- 2962 Tom Schrijvers, Maciej Piróg, Nicolas Wu, and Mauro Jaskelioff. Monad transformers and modular  
2963 algebraic effects: what binds them together. In *Proceedings of the 12th ACM SIGPLAN*  
2964 *International Symposium on Haskell*, pages 98–113, 2019. URL <https://doi.org/10.1145/3331545.3342595>.

- 2966 Gordon Plotkin and John Power. Notions of computation determine monads. In *International*  
2967 *Conference on Foundations of Software Science and Computation Structures*, pages 342–  
2968 356. Springer, 2002. URL [https://era.ed.ac.uk/bitstream/handle/1842/196/Comp\\_Eff\\_](https://era.ed.ac.uk/bitstream/handle/1842/196/Comp_Eff_Monads.pdf?sequence=1)  
2969 [Monads.pdf?sequence=1](https://era.ed.ac.uk/bitstream/handle/1842/196/Comp_Eff_Monads.pdf?sequence=1).
- 2970 Gordon D Plotkin and Matija Pretnar. Handling algebraic effects. *Logical methods in computer*  
2971 *science*, 9, 2013. URL [https://doi.org/10.2168/LMCS-9\(4:23\)2013](https://doi.org/10.2168/LMCS-9(4:23)2013).
- 2972 Nicolas Wu, Tom Schrijvers, and Ralf Hinze. Effect handlers in scope. In *Proceedings of the 2014*  
2973 *ACM SIGPLAN Symposium on Haskell*, pages 1–12, 2014. URL [https://doi.org/10.1145/](https://doi.org/10.1145/2633357.2633358)  
2974 [2633357.2633358](https://doi.org/10.1145/2633357.2633358).
- 2975 Zhixuan Yang, Marco Paviotti, Nicolas Wu, Birthe van den Berg, and Tom Schrijvers. Structured  
2976 handling of scoped effects. In *European Symposium on Programming*, pages 462–491. Springer  
2977 International Publishing Cham, 2022. URL [https://library.oapen.org/bitstream/handle/](https://library.oapen.org/bitstream/handle/20.500.12657/54028/1/978-3-030-99336-8.pdf#page=476)  
2978 [20.500.12657/54028/1/978-3-030-99336-8.pdf#page=476](https://library.oapen.org/bitstream/handle/20.500.12657/54028/1/978-3-030-99336-8.pdf#page=476).
- 2979 Cas van der Rest, Jaro Reinders, and Casper Bach Poulsen. Handling higher-order effects. *arXiv*  
2980 *preprint arXiv:2203.03288*, 2022. URL <https://arxiv.org/pdf/2203.03288>.
- 2981 Yizhou Zhang, Guido Salvaneschi, and Andrew C Myers. Handling bidirectional control flow.  
2982 *Proceedings of the ACM on Programming Languages*, 4(OOPSLA):1–30, 2020. URL [https://](https://doi.org/10.1145/3428207)  
2983 [doi.org/10.1145/3428207](https://doi.org/10.1145/3428207).
- 2984 James Chapman, Pierre-Évariste Dagand, Conor McBride, and Peter Morris. The gentle art of  
2985 levitation. *ACM Sigplan Notices*, 45(9):3–14, 2010. URL [https://citeseerx.ist.psu.edu/](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=55ab751e0af3808c741585df0210ca1785c235d6)  
2986 [document?repid=rep1&type=pdf&doi=55ab751e0af3808c741585df0210ca1785c235d6](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=55ab751e0af3808c741585df0210ca1785c235d6).
- 2987 Yanhong A Liu. Incremental computation: What is the essence? *arXiv preprint arXiv:2312.07946*,  
2988 2023. URL <https://arxiv.org/pdf/2312.07946>.
- 2989 Chris Okasaki. *Purely functional data structures*. Cambridge University Press, 1999. URL [https://](https://www.cs.tufts.edu/~nr/cs257/archive/chris-okasaki/dissertation.pdf)  
2990 [www.cs.tufts.edu/~nr/cs257/archive/chris-okasaki/dissertation.pdf](https://www.cs.tufts.edu/~nr/cs257/archive/chris-okasaki/dissertation.pdf).
- 2991 Aaron Bohannon, Benjamin C Pierce, and Jeffrey A Vaughan. Relational lenses: a  
2992 language for updatable views. In *Proceedings of the twenty-fifth ACM SIGMOD-*  
2993 *SIGACT-SIGART symposium on Principles of database systems*, pages 338–347,  
2994 2006. URL [https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=8ee979ec0a1d8288a07a813c7bde91e4a2ae3187)  
2995 [8ee979ec0a1d8288a07a813c7bde91e4a2ae3187](https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=8ee979ec0a1d8288a07a813c7bde91e4a2ae3187).
- 2996 J Nathan Foster, Alexandre Pilkiewicz, and Benjamin C Pierce. Quotient lenses. *ACM Sigplan*  
2997 *Notices*, 43(9):383–396, 2008. URL [https://repository.upenn.edu/server/api/core/](https://repository.upenn.edu/server/api/core/bitstreams/31dcfba8-8b4f-4143-bd2c-9443dc19c499/content)  
2998 [bitstreams/31dcfba8-8b4f-4143-bd2c-9443dc19c499/content](https://repository.upenn.edu/server/api/core/bitstreams/31dcfba8-8b4f-4143-bd2c-9443dc19c499/content).